

“Dependence and Parallelism”

To remember

- Statement order must not matter.
- Statements must not have dependences.
- Some dependences can be removed.
- Some dependences may not be obvious.

Summary

- No dependence (can run in parallel)
S1: $X = K + 3$;
S2: $Y = Z * 5$;
- True dependence (cannot run in parallel)
S1: $X = 3$;
S2: $Y = X * 5$;
- Anti dependence (cannot run in parallel)
S1: $Y = X * 4$;
S2: $X = 3$;
- Output dependence (cannot run in parallel)
S1: $X = Y * 4$;
S2: $X = 3$;
- Loop-carried dependence

Parallel algorithm Design

Programming Models


- The programming model
 - o determines the basic concepts of the parallel implementation and
 - o abstracts from the hardware as well as from the programming language or API.
- The names used for programming models differ in the literature.

Programming Models (2)

- **Sequential Model:** The sequential program is automatically parallelized.
 - o Advantage: Familiar programming model
 - o Disadvantage: Limitations in compiler analysis
- **Message Passing Model:** The application consists of a set of processes with separate address spaces. The processes exchange messages by explicit send/receive operations.
 - o Advantage: Full control of performance aspects
 - o Disadvantage: Complex programming

Assembly programming for parallel architectures!


MP Programming Model



Programming Models (3)

- **Shared Memory Model:** Application consists of parallel threads, accessing shared data structures protected by synchronization operations.
 - o **Thread-Based SM Model:** Explicit programming of cooperating threads.
 - Advantage: Portability
 - Disadvantage: Complex programming
 - o **Directive-Based SM Model:** Parallel loops with implicit synchronization.
 - Advantage: Easier to program
 - Disadvantage: Difficult to control performance
 - o **Remote Memory Access Model:** Parallel threads accessing mainly private data structures. Shared data structures are accessed via special operations.
 - Advantage: Can be easily combined with message passing
 - Disadvantage: Complex programming

Shared Memory Programming Model



Programming Models (4)

- **Data Parallel Programming Model:** Synchronized execution of parallel operations on large distributed data structures.
 - o Advantage: High-level programming
 - o Disadvantage: Performance of current implementations

Together...

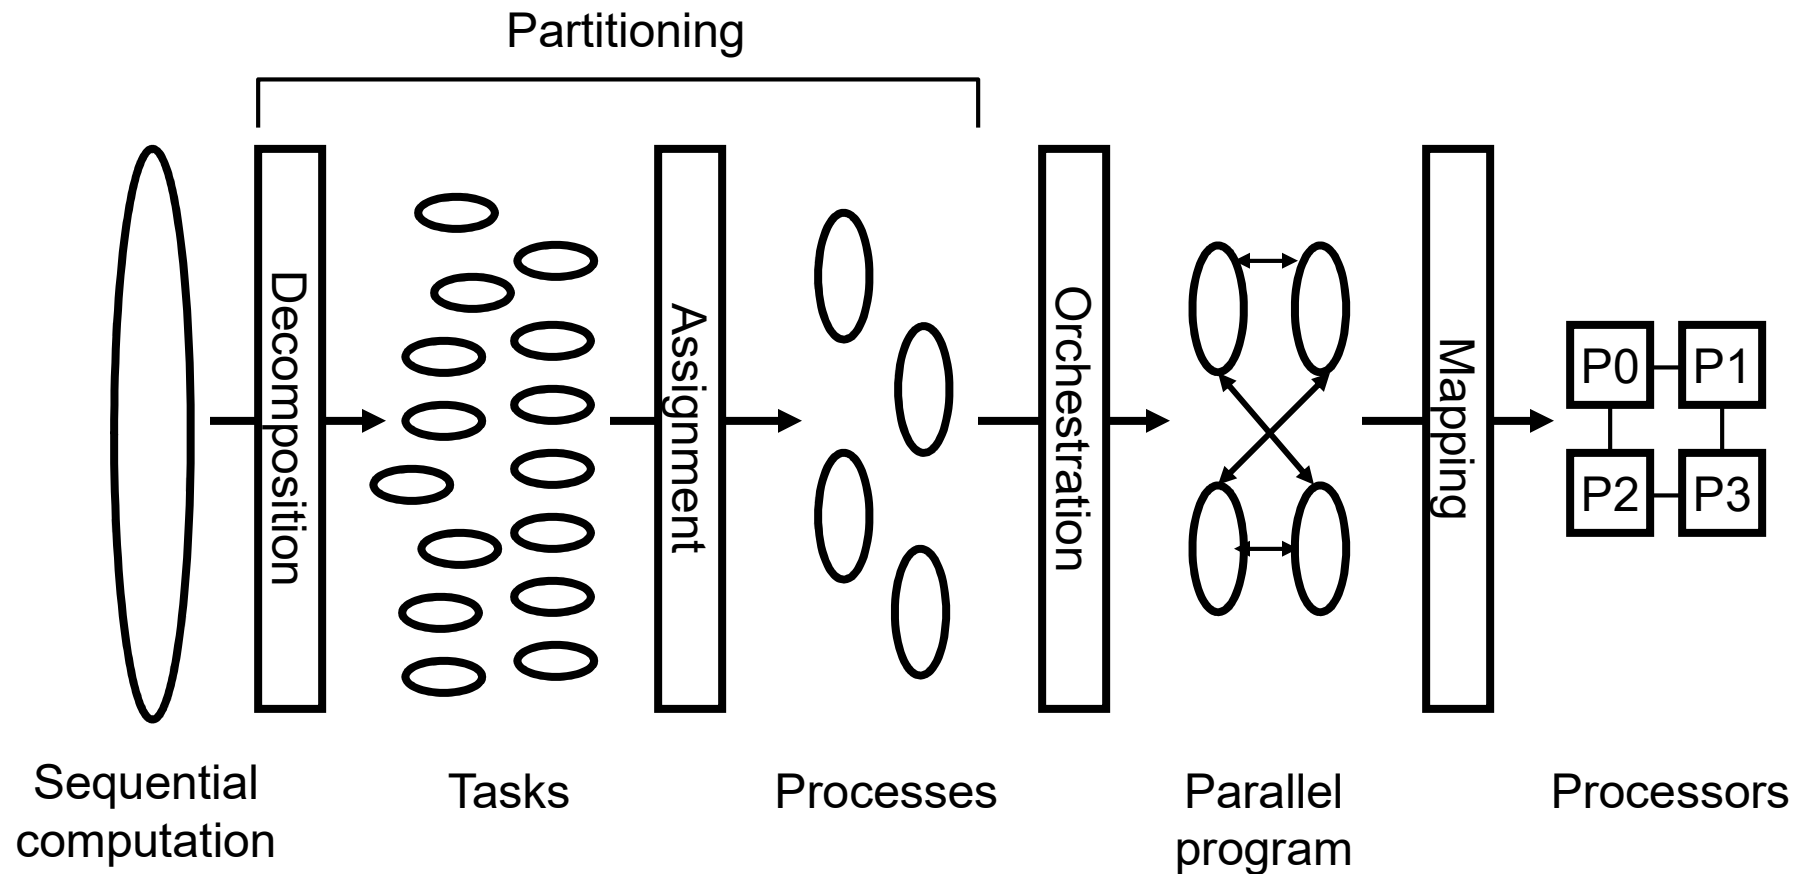
- Parallel programming models
 - Shared memory (Pthreads and openMP)
 - Message passing (MPI)
 - Massive data parallelism (GPU programming)
 - Hybrid
- Distributed programming models
 - Cloud programming, MapReduce, Spark

Concepts of Parallel Programming

➤ Concepts:

- o **Task**: arbitrary piece of work performed by a single process
- o **Thread**: is an abstract entity as part of a process that performs tasks. Defines a unit for scheduling.
- o **Process**: is active entity with resources that performs tasks.
- o **Processor**: is a physical resource executing processes


Phases in the Parallelization Process



Decomposition


- Dividing computation and data into pieces
- Domain decomposition
 - o Divide data into pieces
 - o Determine how to associate computations with the data
- Functional decomposition
 - o Divide computation into pieces
 - o Determine how to associate data with the computations

Example Domain Decompositions

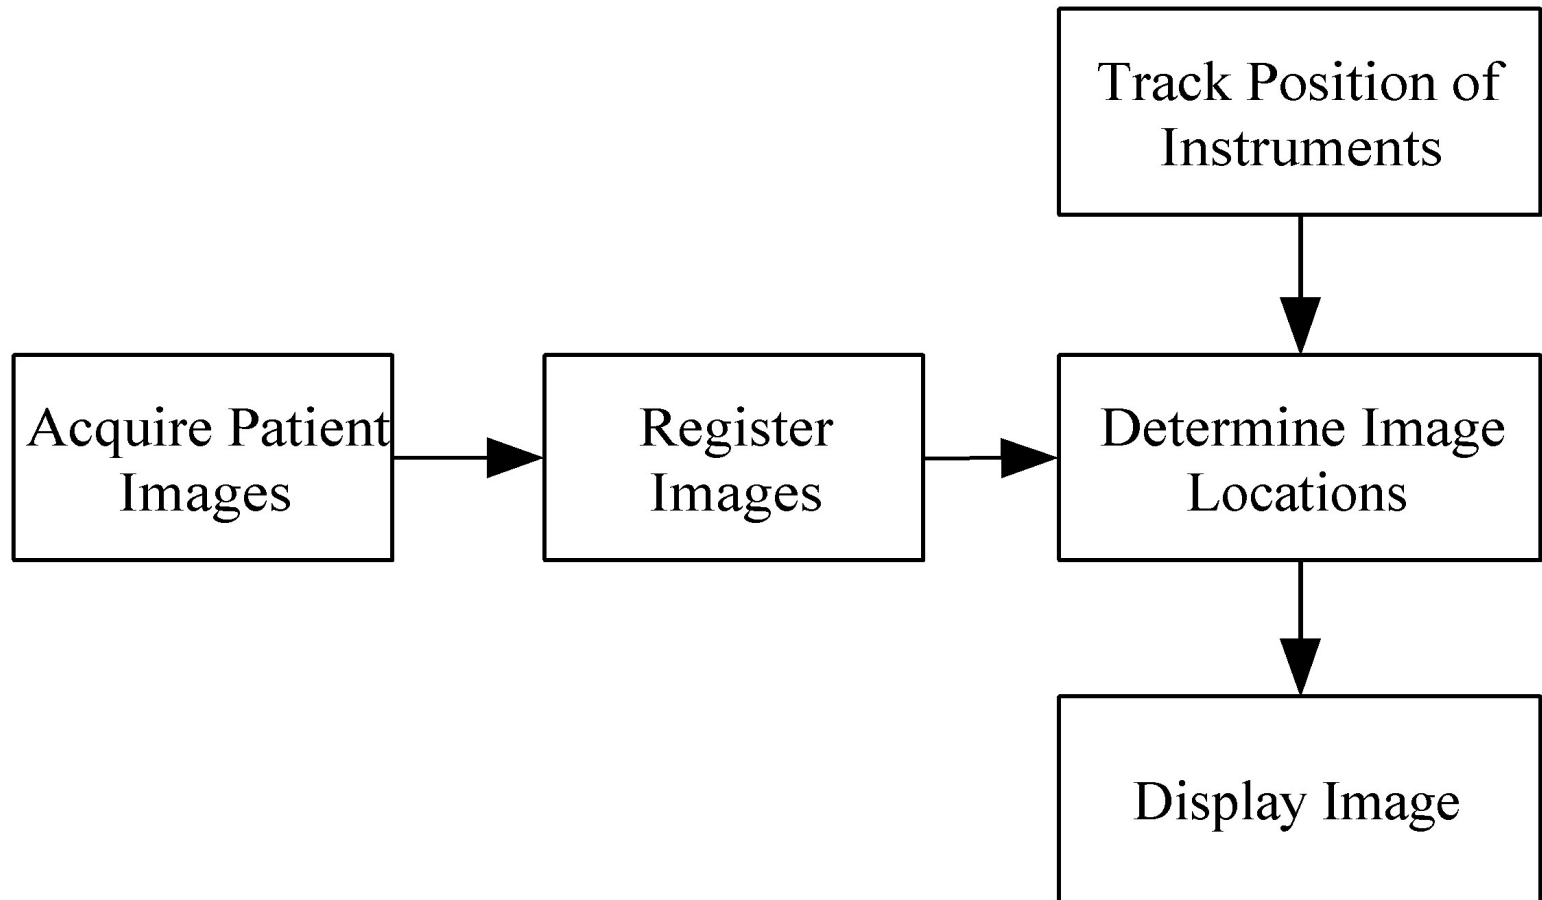


Domain Decomposition

- The **data** associated with a problem is decomposed!
 - o If possible, divide data into small pieces of equal size
- Each parallel task then works on a portion of the data
 - o Generates number of tasks, each with some data and set of operations on data



Example Functional Decomposition



Functional Decomposition

- Breaking the computation into a collection of tasks
 - o Tasks may be of variable length
 - o Tasks require data to be executed
 - o Tasks may become available dynamically


Data vs Functional Parallelism

➤ Data parallelism

- o The same operations are executed in parallel for the elements of large data structures, e.g. arrays.
- o Tasks are the operations on each individual element or on subsets of the elements.
- o Whether tasks are of same length or variable length depends on the application. Quite some applications have tasks of same length.


Example: Data Parallelism

```
for (i=0;i<n;i++)  
  a[i]=b[i]+c[i]
```



Example: Data Parallelism, Variable Length

```
for (i=0;i<n;i++)  
  for (j=0;j<=i;j++)  
    a[i]=a[i]+b[i][j]
```



Data vs Functional Parallelism

➤ Functional parallelism

- o Entirely different calculations can be performed concurrently on either the same or different data.
- o The tasks are usually specified via different functions or different code regions.
- o The degree of available functional parallelism is usually modest.
- o Tasks are of different length in most cases.

Example: Function Parallelism


```
f (a) ;  
g (b) ;
```

```
i=i+10 ;  
j=2*k+z**2 ;
```

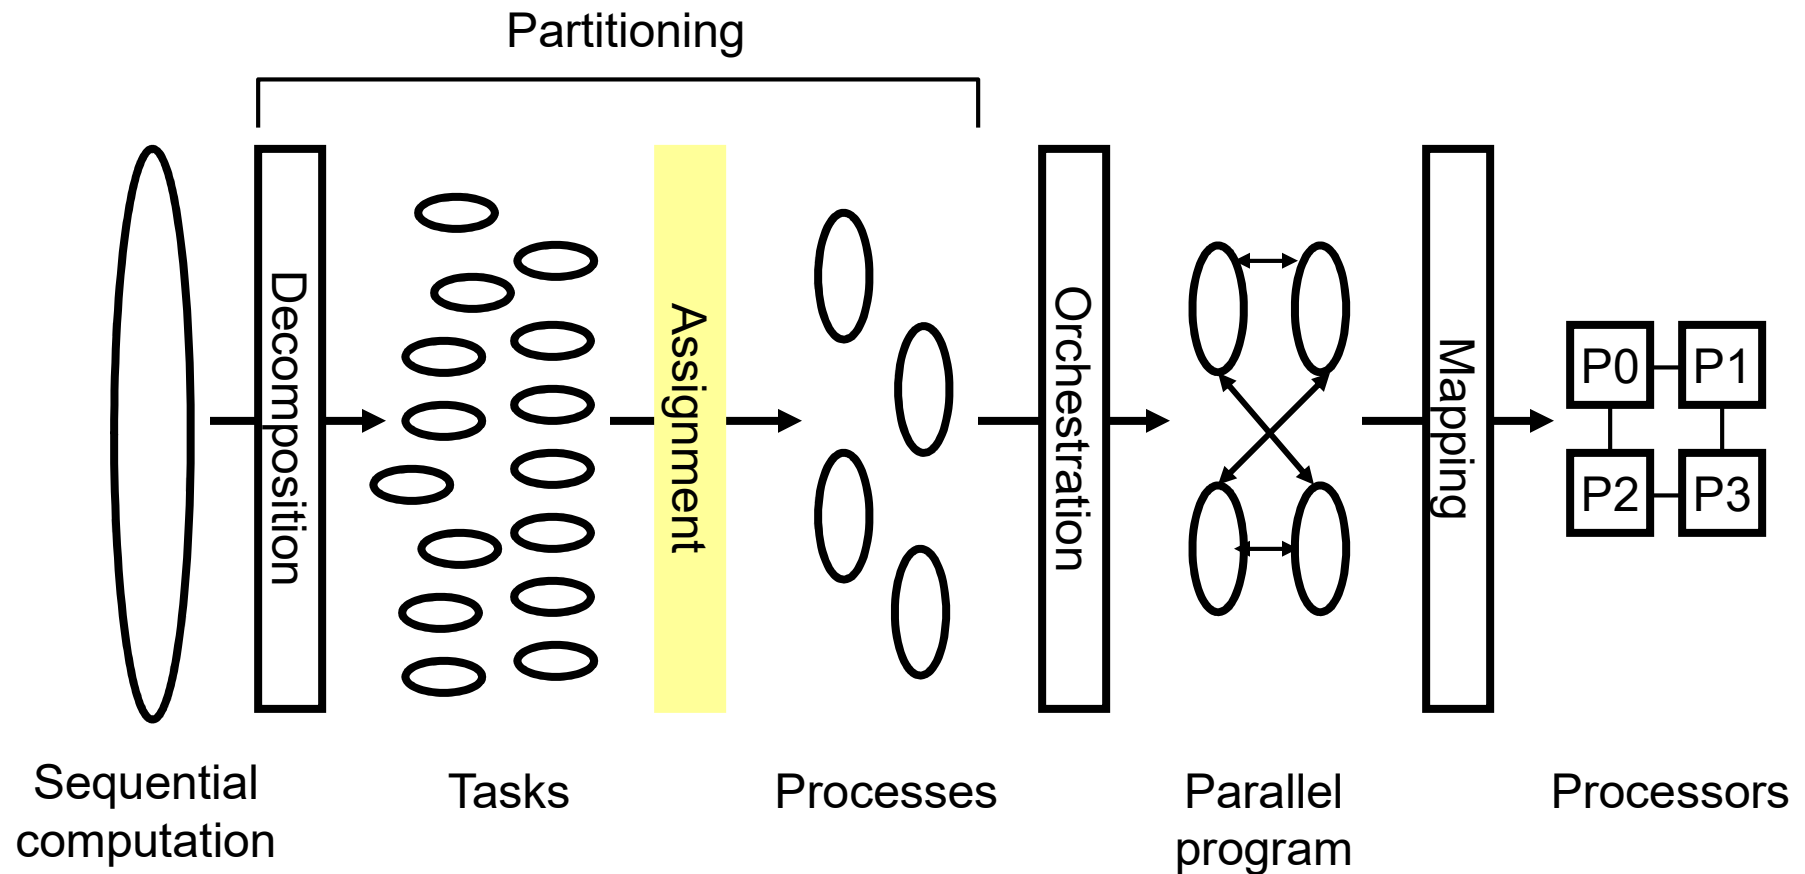
- The functions or statements can be executed in parallel.
- These are different operations → functional parallelism

Functional Decomposition Example

- A simulation of the earth's climate comprises atmosphere, ocean, hydrology, etc.



Phases in the Parallelization Process




Assignment

- Assignment means specifying the mechanism by which tasks will be distributed among processes.
- Goals:
 - o Balance workload
 - o Reduce interprocess communication
 - o Reduce assignment overhead
- Assignment time
 - o **Static**: fixed assignment during compilation or program creation
 - o **Dynamic**: adaptive assignment during execution


Example: Balance Workload

```
for (i=0;i<n;i++)  
  a[i]=b[i]+c[i]
```



Example: Balance Workload (2)

```
for (i=0;i<n;i++)
    for (j=0;j<i;j++)
        a[i]=a[i]+b[i][j]
```



➤ Static Load Balancing

- o 2 different assignments

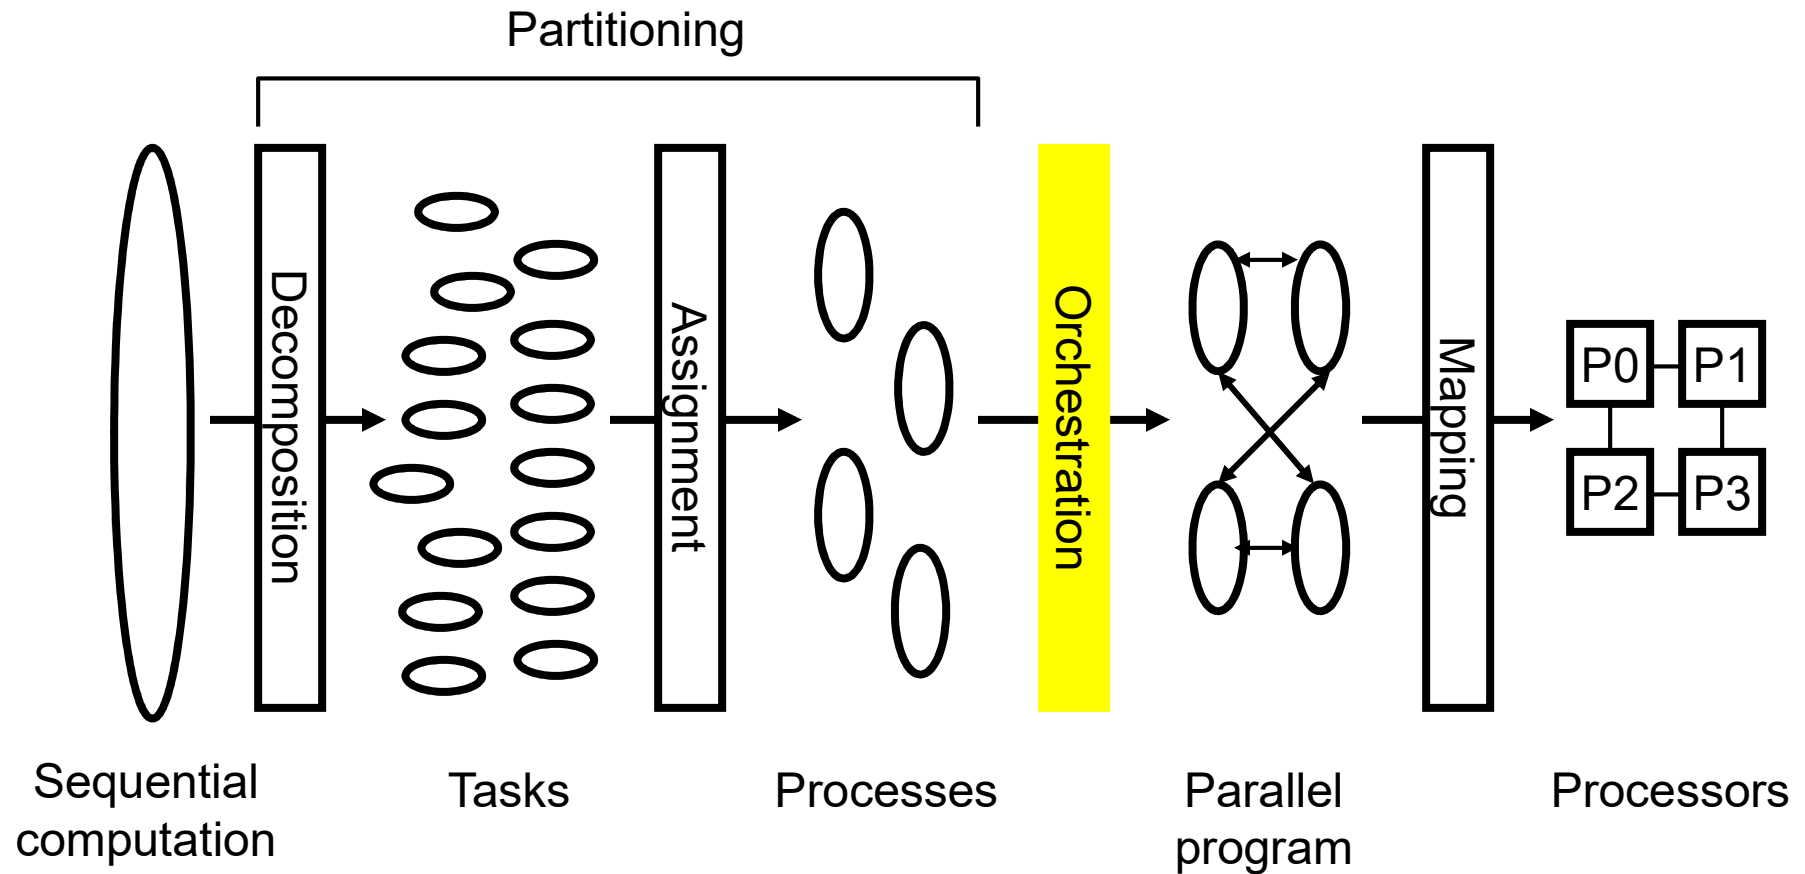
➤ Dynamic Load Balancing

- o postpone assignment until execution time

Partitioning Checklist

- Does your partition define at least an order of magnitude more tasks than there are processors?
- Does your partition avoid redundant computation and storage requirements?
- Are tasks of equal size?
- Does no. of tasks scale with problem size?
- Have you identified alternate partitions?

Phases in the Parallelization Process



Orchestration

- Implementation in a given programming model
- Means for
 - o Naming and accessing data
 - o Exchanging data
 - o Synchronization
- Questions
 - o How to organize data structures?
 - o How to schedule assigned tasks to improve locality?
 - o Whether to communicate in large or small messages?
- Performance goal
 - o Reduction of communication and synchronization overhead
 - o Reduction of parallelization overhead
 - o Reduction of idle time
 - o Ideal load balancing
- Sometimes in some cases goals can be conflicting
 - o reduction of communication versus reduction of parallelization overhead

Communication (MPI)

- Tasks often execute concurrently, but not independently
 - o The computation to be performed in one task requires data from another
 - o Information flow is specified in communication phase
- We can conceptualize a need for comm between two tasks as a channel linking tasks
- Comm is easy for functional decomposition, not easy for domain decomposition
- Latency vs Bandwidth
 - o Sending many small messages can cause latency to dominate communication overhead.
 - o Package small messages into a larger message (aka increasing the effective comm bandwidth)

Categories of Communication

- Local versus global
 - o Local comm: ?
 - o Global comm: ?
- Structured versus unstructured
 - o Structured comm: ?
 - o Unstructured comm: ?
- Static versus dynamic
 - o Static: ?
 - o Dynamic: ?
- Synchronous versus asynchronous
 - o Synchronous: ?
 - o Asynchronous: ?

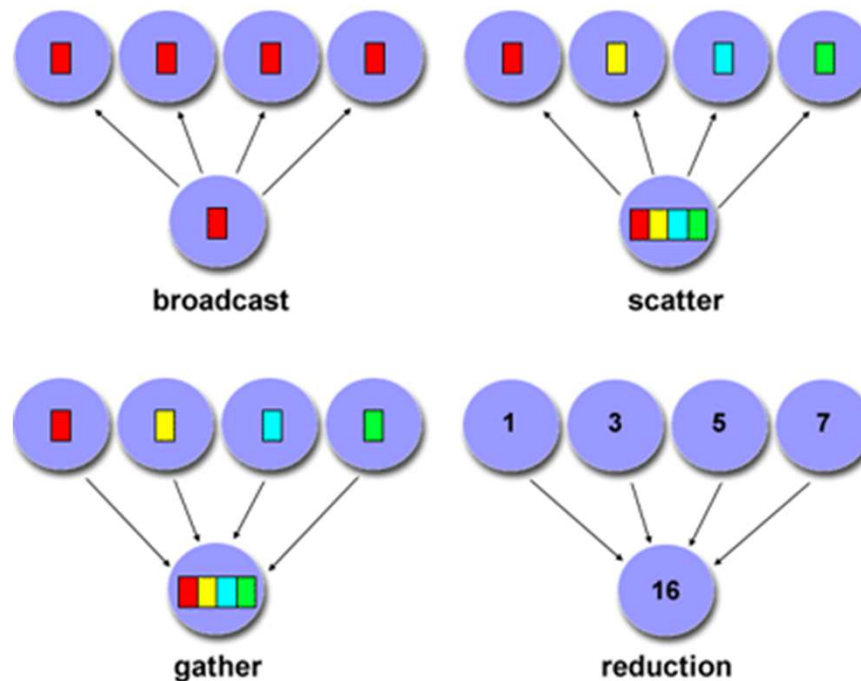
Scope of Communications

➤ Point-to-point

- o involves two tasks with one acting as the sender and the other acting as the receiver

➤ Collective


- o Involves data sharing among more than two tasks



Local Communication


- Consider comm requirements of Jacobi finite difference scheme

$$x_{i,j}^{(t+1)} = \frac{4x_{i,j}^{(t)} + x_{i-1,j}^{(t)} + x_{i+1,j}^{(t)} + x_{i,j-1}^{(t)} + x_{i,j+1}^{(t)}}{8}$$



Global Communication

- Consider problem of performing a parallel reduction of N numbers







Communication Design Checklist

- Do all tasks perform same amount of communication operations?
- Does each task communicate with small number of neighbors?
- Are communication operations able to proceed concurrently?
- Is computation for different tasks able to proceed concurrently?

Agglomeration (MPI)


- In first two phases, we partitioned the computation to be performed into a set of tasks and introduced comm to provide data required by these tasks
- In agglomerate stage, we consider if it is useful to combine tasks to provide smaller number of tasks
- Determine if worthwhile to replicate data or computation




Increasing Granularity

- In partitioning phase, efforts focused on exposing parallelism
- Large number of parallel tasks does not result in efficient parallel algorithm
 - o For comm, send same data in less number messages
 - o Less task creation costs and task scheduling costs
- **Fine-grain vs coarse-grain** parallelism
 - o Dependent on the algorithm and the hardware environment in which it runs

(a)







(b)




Replicating Computations

- Consider problem of replicating the sum of N numbers in N processors
- Use a sum reduction followed by a broadcast takes $2(N-1)$ in ring and $2\log N$ for a tree



Replicating Computations

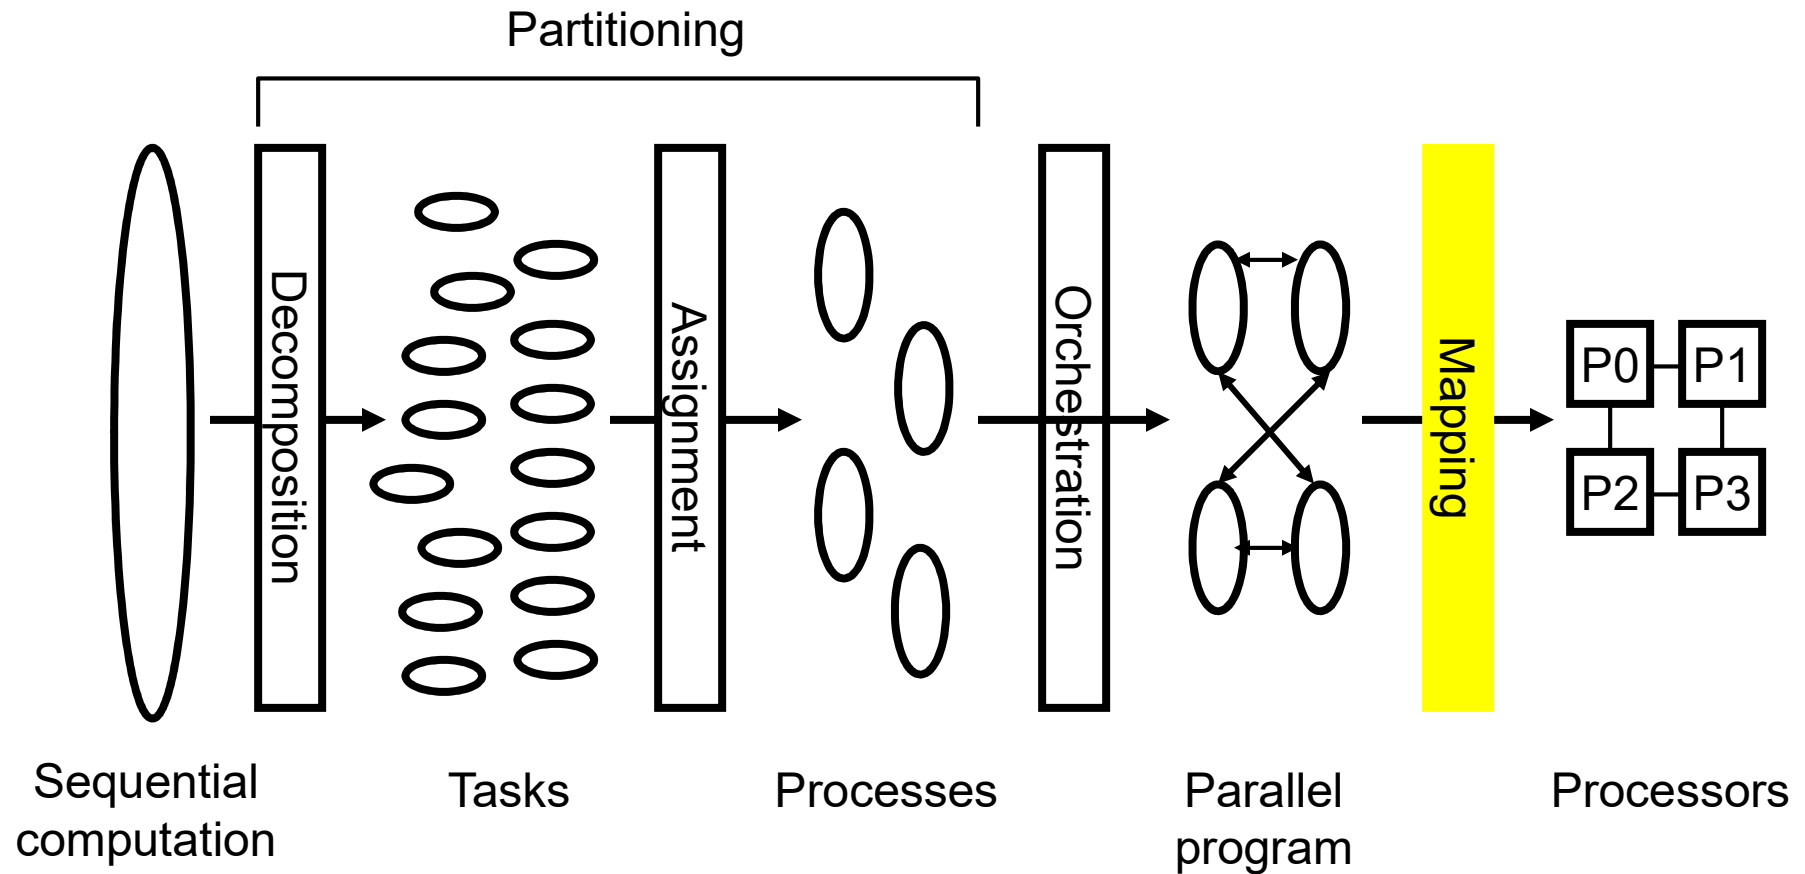
- Can accomplish in $\log(N)$ steps in butterfly algorithm



Agglomeration Checklist

- Has agglomeration reduced comm costs by increasing locality?
- If agglomeration uses replicated computations, do benefits outweigh costs?
- If agglomeration replicates data, is scalability affected?
- Has agglomeration yielded tasks of similar computation and comm costs?
- Does number of tasks still scale with problem size?


Phases in the Parallelization Process




Mapping

- Mapping processes to processors
- Done by the program and/or operating system
- Shared memory system: mapping done by operating system
- Distributed memory system: mapping done by user
- Conflicting goals of mapping
 - o Maximize processor utilization
 - o Minimize interprocessor communication

Mapping Example




(a)



(b)

Mapping Problem

- Mapping problem to minimize execution time is **NP-complete** (definition?)
 - o Hence resort to heuristics
- On modern computer systems, it is also a multilevel problem



Performance Goals

Step	Architecture-Dependent?	Major Performance Goal
Decomposition	Mostly No	➤ Expose enough concurrency but not too much
Assignment	Mostly no	➤ Balance workload ➤ Reduce communication volume
Orchestration	Yes	➤ Reduce unnecessary communication via data locality ➤ Reduce communication and synchronization cost as seen by the processor ➤ Reduce serialization at shared resources
Mapping	Yes	➤ Put related processes on the same processor if necessary ➤ Exploit locality in network topology

Application Structure

- Frequently used patterns for parallel applications:
 - o Single Program Multiple Data – SPMD (Domain Decomposition)
 - o Embarrassingly Parallel
 - o Master / Slave
 - o Work Pool
 - o Divide and Conquer
 - o Pipeline
 - o Competition