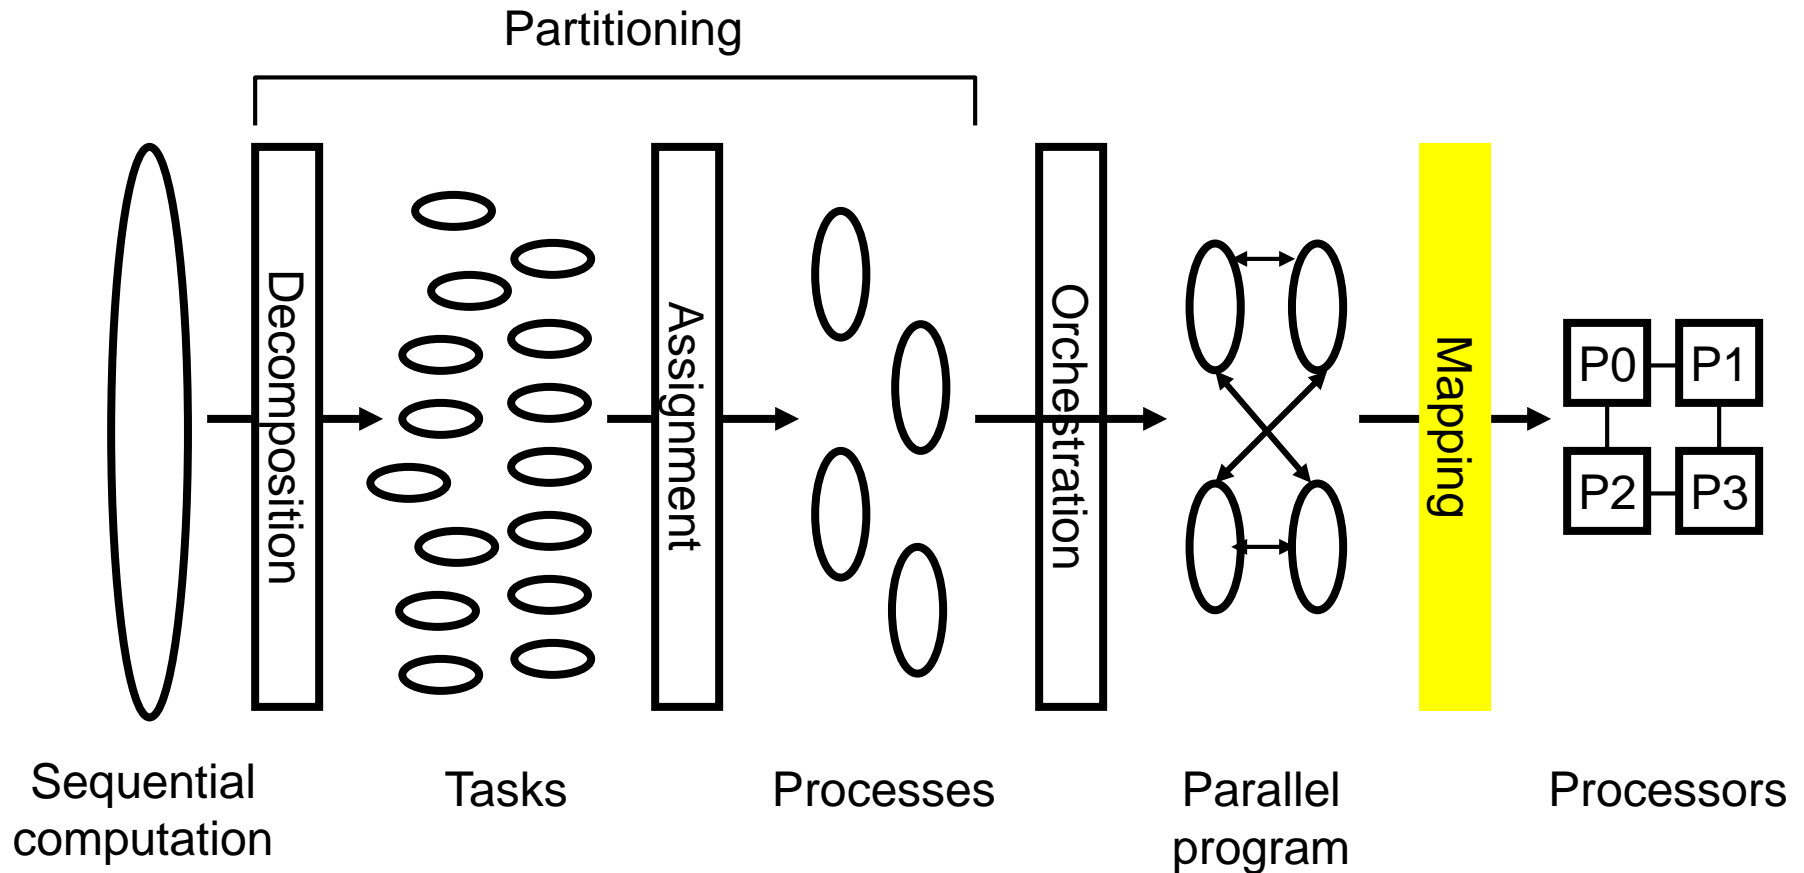



Phases in the Parallelization Process




Mapping

- Mapping processes to processors
- Done by the program and/or operating system
- Shared memory system: mapping done by operating system
- Distributed memory system: mapping done by user
- Conflicting goals of mapping
 - Maximize processor utilization
 - Minimize interprocessor communication

Mapping Example




(a)



(b)

Mapping Problem

- Mapping problem to minimize execution time is **NP-complete** (definition?)
 - o Hence resort to heuristics
- On modern computer systems, it is also a multilevel problem



Performance Goals

Step	Architecture-Dependent?	Major Performance Goal
Decomposition	Mostly No	➤ Expose enough concurrency but not too much
Assignment	Mostly no	➤ Balance workload ➤ Reduce communication volume
Orchestration	Yes	➤ Reduce unnecessary communication via data locality ➤ Reduce communication and synchronization cost as seen by the processor ➤ Reduce serialization at shared resources
Mapping	Yes	➤ Put related processes on the same processor if necessary ➤ Exploit locality in network topology

Parallel Algorithm Design

Case Study: Tridiagonal Solvers

Xian-He Sun

Department of Computer Science
Illinois Institute of Technology
sun@iit.edu

Outline

- Problem Description
- Parallel Algorithms
 - The Partition Method
 - The PPT Algorithm
 - The PDD Algorithm
 - The LU Pipelining Algorithm
 - The PTH Method and PPD Algorithm
- Implementations

Problem Description

- Tridiagonal linear system

$$Ax = d$$

$$A = \begin{pmatrix} b_0 & c_0 & & & \\ a_1 & b & c_1 & & \\ & \ddots & \ddots & \ddots & \\ & & a_{n-2} & b_{n-2} & c_{n-2} \\ & & & a_{n-1} & b_{n-1} \end{pmatrix} = \tilde{A} + \Delta A$$

Sequential Solver

Problem

$$a_k x_{k-1} + b_k x_k + c_k x_{k+1} = d_k$$
$$(k=2, \dots, N)$$

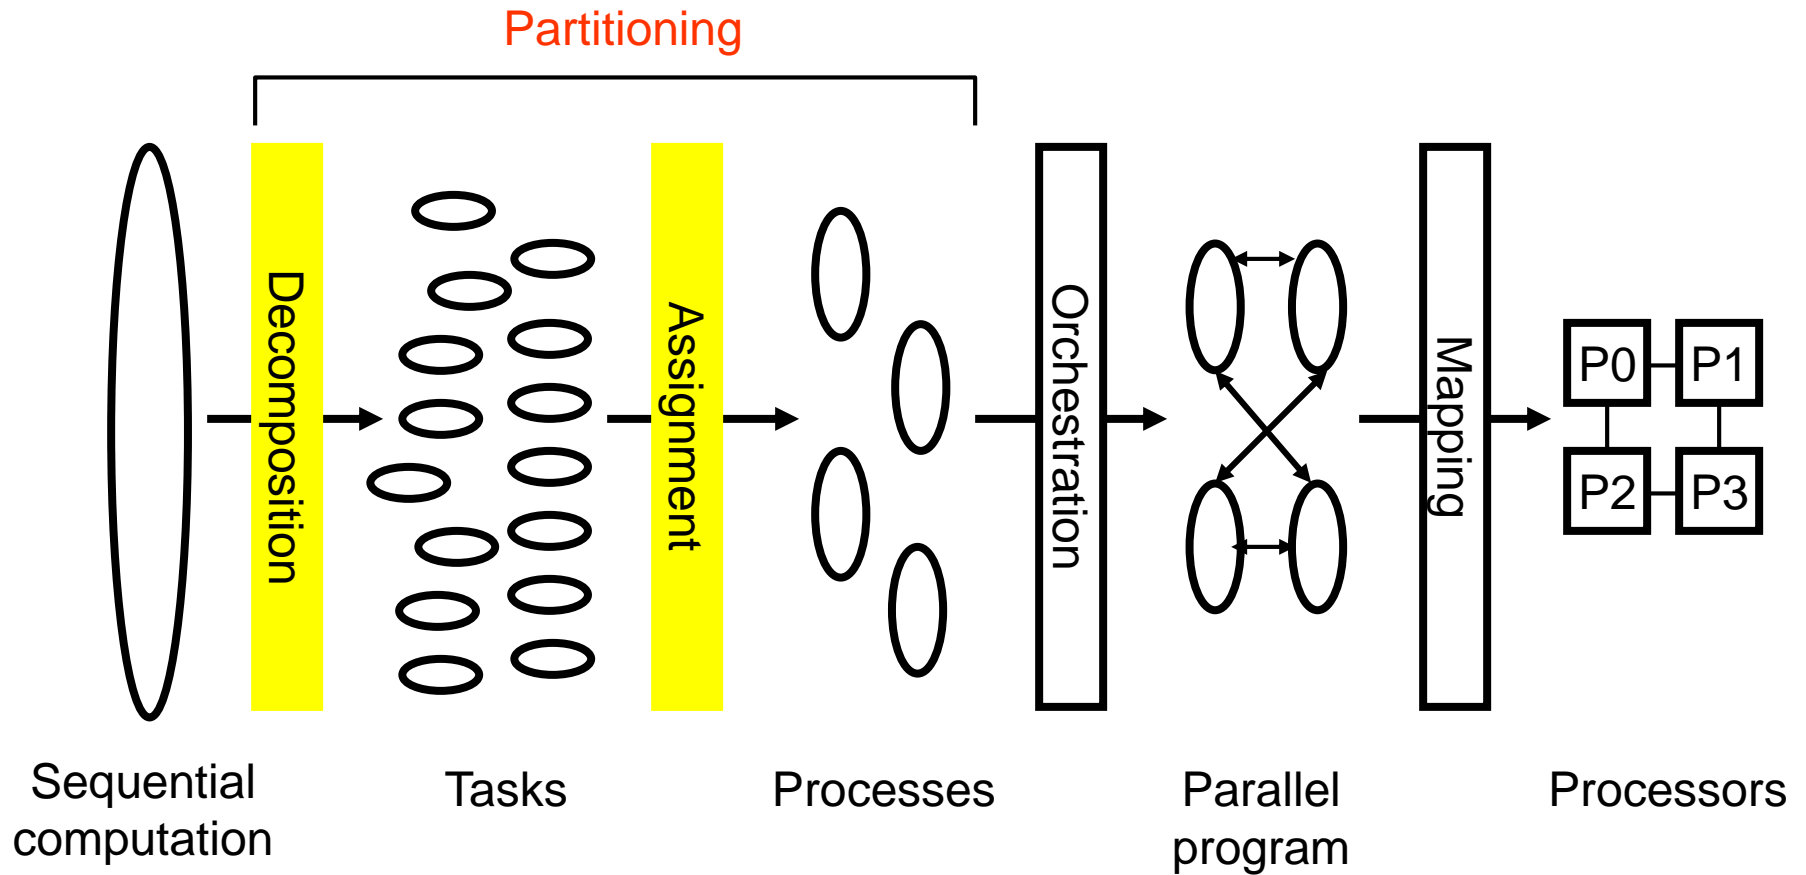
Forward step

$$\beta_1 = b_1 \qquad \beta_k = b_k - a_k c_{k-1} / \beta_{k-1}$$
$$\alpha_1 = d_1 / \beta_1 \qquad \alpha_k = (-a_k \alpha_{k-1} + d_k) / \beta_k$$
$$(k=2, \dots, N)$$

Backward Step

$$x_n = \alpha_n \qquad x_k = (\alpha_k - x_{k+1} c_k) / \beta_k$$
$$(k=N-1, \dots, 1)$$

Partition



The Matrix Modification Formula

$$x = \mathbf{A}^{-1} \mathbf{d} = (\tilde{\mathbf{A}} + \Delta \mathbf{A})^{-1} \mathbf{d} = (\tilde{\mathbf{A}} + \mathbf{V} \mathbf{E}^T)^{-1} \mathbf{d}$$

$$x = \tilde{\mathbf{A}}^{-1} \mathbf{d} - \tilde{\mathbf{A}}^{-1} \mathbf{V} (\mathbf{I} + \mathbf{E}^T \tilde{\mathbf{A}}^{-1} \mathbf{V})^{-1} \mathbf{E}^T \tilde{\mathbf{A}}^{-1} \mathbf{d}$$

$$\tilde{\mathbf{A}} = \begin{bmatrix} A_0 & & & \\ & \text{---} & & \\ & & A_1 & \\ & & & \text{---} \\ & & & & \ddots & \\ & & & & & A_{p-1} \end{bmatrix}$$

$$\Delta \mathbf{A} = \begin{bmatrix} & c_{m-1} & & \\ a_m & & & \\ & c_{2m-1} & & \\ & & a_{2m} & \\ & & & \ddots & \\ & & & & c_{m(p-1)-1} \\ & & & & & a_{m(p-1)} \end{bmatrix}$$

The Partition of Tridiagonal Systems

$$A = \tilde{A} + VE^T$$

$$\Delta A = [a_m e_m, c_{m-1} e_{m-1}, a_{2m} e_{2m}, c_{2m-1} e_{2m-1}, \dots, c_{(p-1)m-1} e_{(p-1)m-1}] \cdot \begin{bmatrix} e_{m-1}^T \\ e_m^T \\ \cdot \\ \cdot \\ \cdot \\ e_{(p-1)m-1}^T \\ e_{(p-1)m}^T \end{bmatrix} = VE^T$$

e_i are column vector with i th element being one and all the other entries being zero.

The Solving process

1. Solve the subsystems in parallel
2. Solve the reduced system
3. Modification

$$\tilde{\mathbf{A}}^{-1} \mathbf{A} \mathbf{x} = \tilde{\mathbf{A}}^{-1} \mathbf{d}$$

$$\begin{bmatrix} A_0^{-1} & & & \\ & A_1^{-1} & & \\ & & & \\ & & & A_{p-1}^{-1} \end{bmatrix} \begin{bmatrix} A_0 & * & & \\ * & A_1 & & \\ & * & * & \\ & & & * & A_{p-1} \end{bmatrix} \begin{pmatrix} x_0 \\ x_1 \\ \cdot \\ \cdot \\ \cdot \\ x_{n-1} \end{pmatrix} = \begin{bmatrix} A_0^{-1} & & & \\ & A_1^{-1} & & \\ & & & \\ & & & A_{p-1}^{-1} \end{bmatrix} \begin{pmatrix} d_0 \\ d_1 \\ \cdot \\ \cdot \\ \cdot \\ d_{n-1} \end{pmatrix}$$

The Solving Procedure

$$\tilde{A}x = d$$

$$\tilde{A}Y = V$$

$$h = E^T \tilde{x}$$

$$Z = I + E^T Y$$

$$Zy = h$$

$$\Delta x = Yy$$

$$x = \tilde{x} - \Delta x$$

The Reduced System $(Zy=h)$

$$\begin{bmatrix} \mathbf{I} & \text{red} \\ \text{red} & \mathbf{I} \end{bmatrix} \begin{pmatrix} \cdot \\ \cdot \\ \cdot \\ \cdot \end{pmatrix} = \begin{pmatrix} \cdot \\ \cdot \\ \cdot \\ \cdot \end{pmatrix}$$

Needs global communication

The Parallel Partition LU (PPT) Algorithm

Step 1. Allocate $A_i, d^{(i)}$ and elements $a_{im}, c_{(i+1)m-1}$ to the i th node, where $0 \leq i \leq p-1$.

Step 2. Use the LU decomposition method to solve

$$A_i[\tilde{x}^{(i)}, v^{(i)}, w^{(i)}] = [d^{(i)}, a_{im}e_0, c_{(i+1)m-1}e_{m-1}]$$

Step 3. Send $\tilde{x}_0^{(i)}, \tilde{x}_{m-1}^{(i)}, v_0^{(i)}, v_{m-1}^{(i)}, w_0^{(i)}, w_{m-1}^{(i)}$ from the i th node to the other nodes $0 \leq i \leq p-1$.

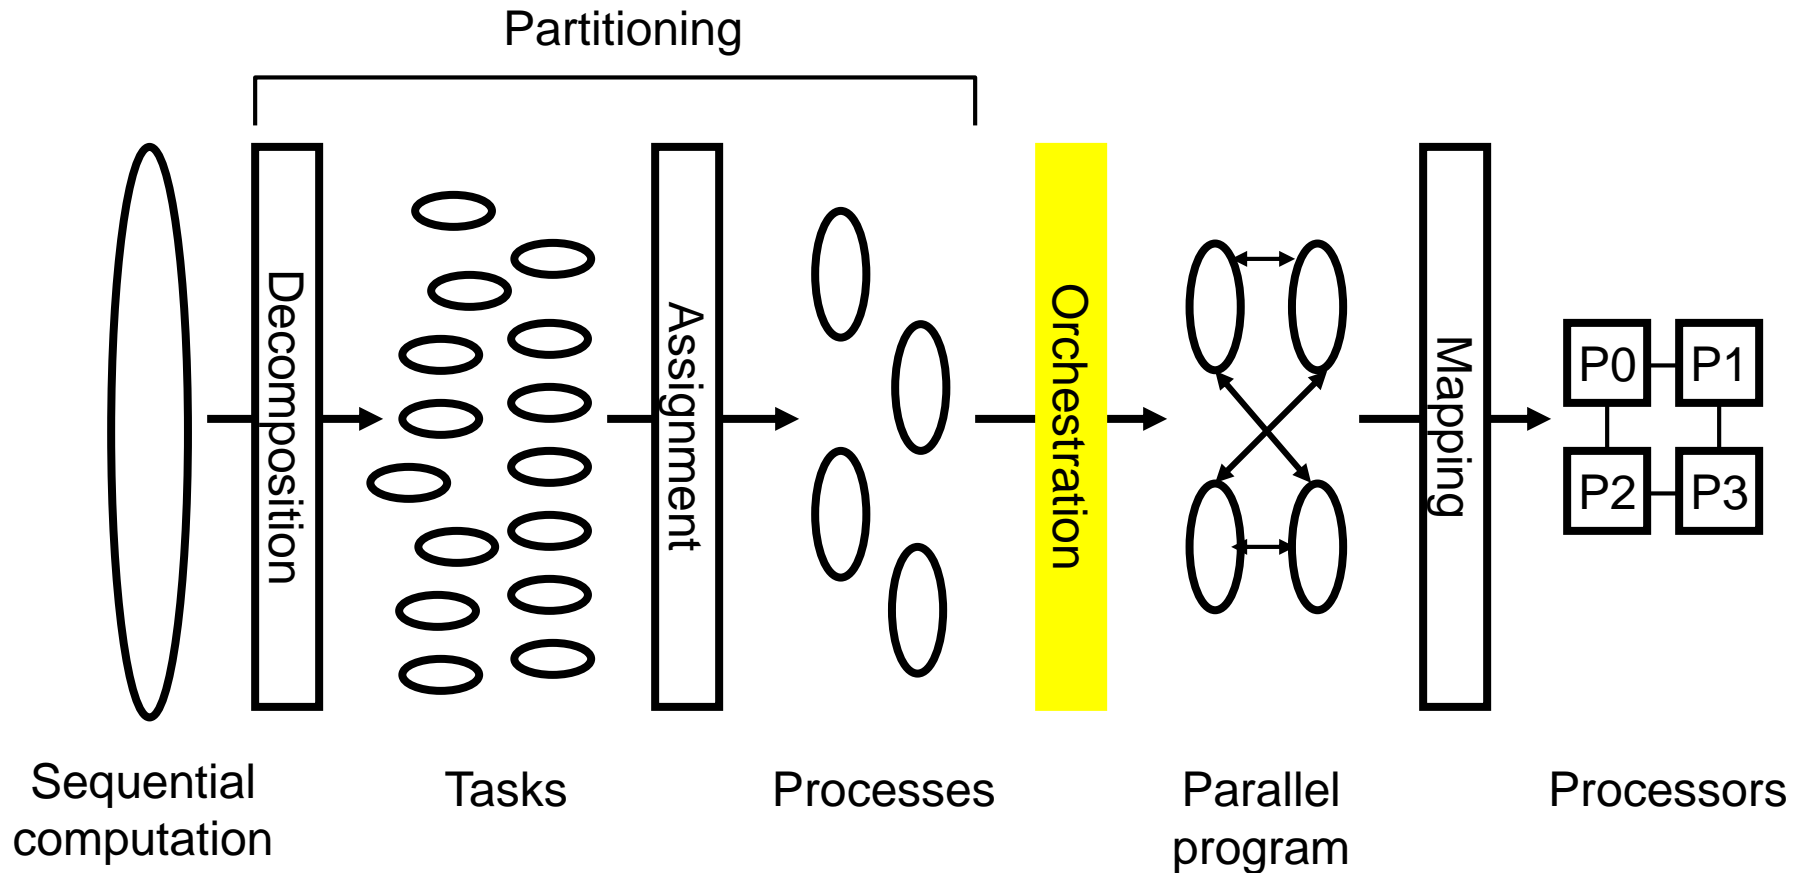
Step 4. Use the LU method to solve $Zy = h$ on all nodes

Step 5. Compute in parallel on p processors

$$\Delta x^{(i)} = [v^{(i)}, w^{(i)}] \begin{bmatrix} y_{2i-1} \\ y_{2i} \end{bmatrix}$$

$$x^{(i)} = \tilde{x}^{(i)} - \Delta x^{(i)}$$

Orchestration




Orchestration


Orchestration is implied in the PPT algorithm

- Intuitively, the reduced system should be solved on one node
 - A tree-reduction communication to get the data
 - Solve
 - A reversed tree-reduction communication to set the results
 - $2 \log(p)$ communication, one solving
- In PPT algorithm (step 3)
 - One total data exchange
 - All nodes solve the reduced system concurrently
 - $1 \log(p)$ communication, one solving

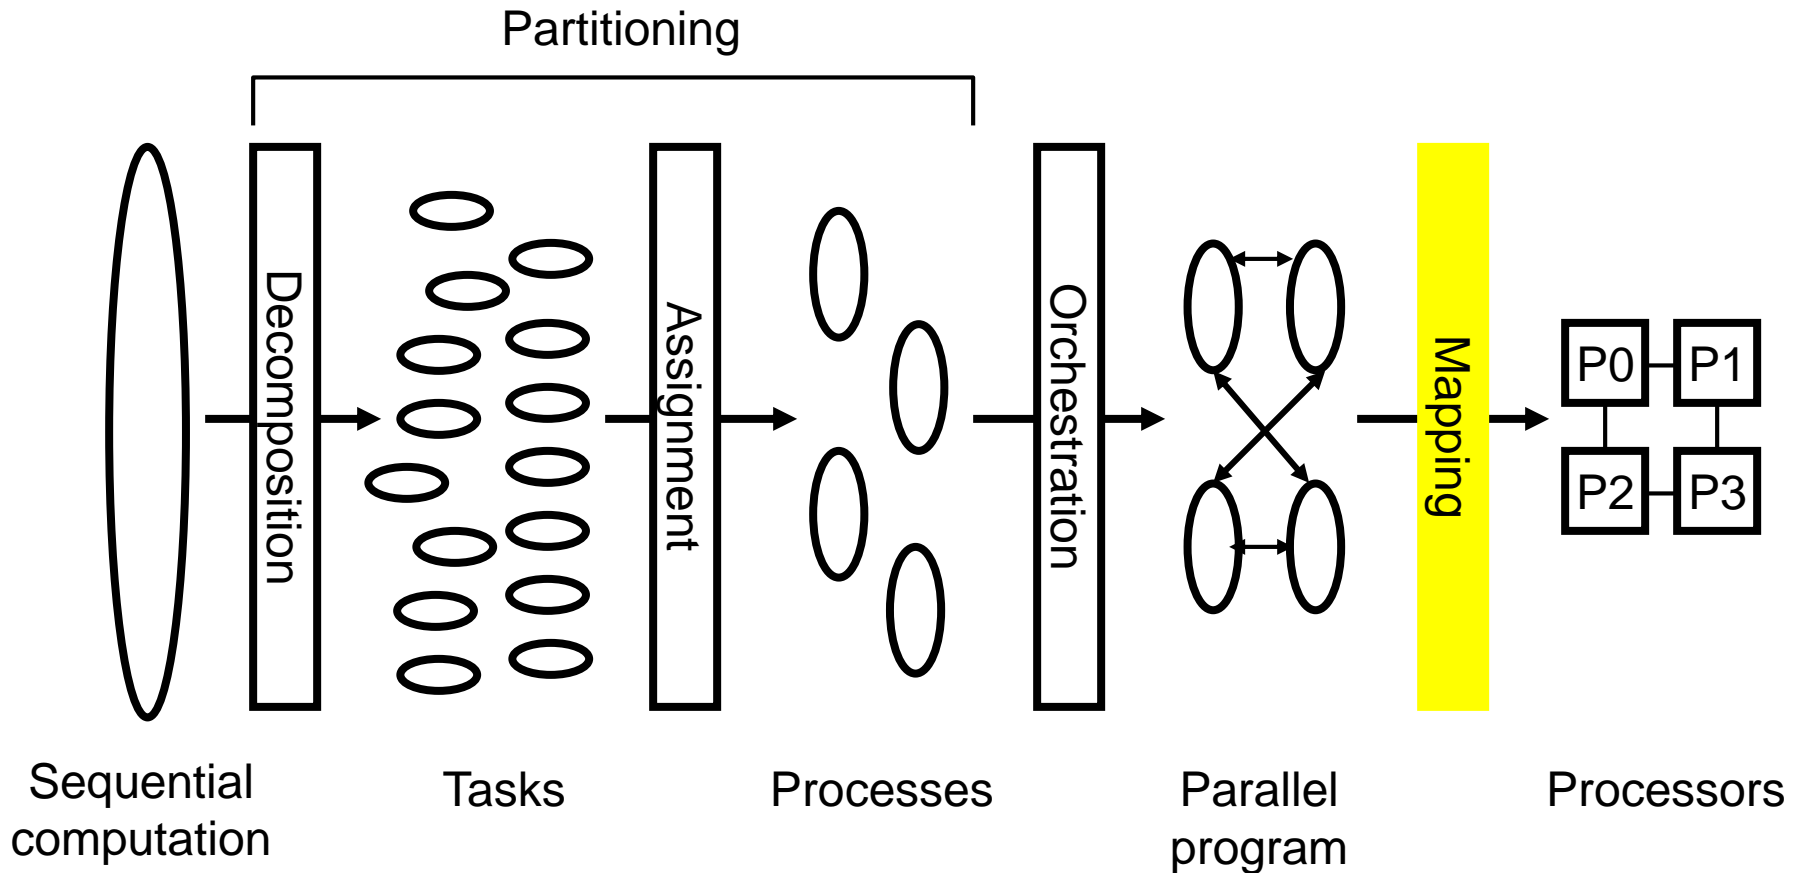
Tree Reduction (Data gathering/scattering)



All-to-All Total Data Exchange



Mapping



Mapping

- Try to reduce the communication
 - Reduce time
 - Reduce message size
 - Reduce cost: distance, contention, congestion, etc
- In total data exchange
 - Try to make every comm. a direct comm.
 - Can be achieved in hypercube architecture

The PPT Algorithm

- Advantage
 - Perfect parallel
- Disadvantage
 - Increased computation (vs. sequential alg.)
 - Global communication
 - Sequential bottleneck


Problem Description

- Parallel codes have been developed during last decade
- The performances of many codes suffer in a scalable computing environment
- Need to identify and overcome the scalability bottlenecks

Diagonal Dominant Systems

$$\begin{pmatrix} 1 & \frac{2}{9} & & & \\ \frac{2}{7} & 1 & \frac{1}{5} & & \\ \cdot & \cdot & \cdot & & \\ & \cdot & \cdot & \frac{1}{6} & \\ & & & 1 & \frac{3}{7} \\ & & & \frac{3}{8} & 1 \end{pmatrix} \begin{pmatrix} x_0 \\ x_1 \\ \cdot \\ \cdot \\ \cdot \\ x_{n-1} \end{pmatrix} = \begin{pmatrix} d_0 \\ \cdot \\ \cdot \\ \cdot \\ \cdot \\ d_{n-1} \end{pmatrix}$$

- The Reduced System of Diagonal Dominant Systems




- Decay Bound for Inverses of Band Matrices

$$\left| A^{-1}(i, j) \right| \leq Cq^{|i-j|/m}, \quad 0 < q < 1$$

$$q = q(r) = \frac{\sqrt{r} - 1}{\sqrt{r} + 1} \quad r = \frac{\lambda_{\max}}{\lambda_{\min}}$$

The Reduced communication



Generally needs global communication, Decay for diagonal dominant systems

The figure shows two tensor network diagrams, labeled Z and \tilde{Z} , which are shown to be equivalent (\approx).

Diagram Z : It consists of a vertical identity I on the left. In the center, there are two horizontal identity I tensors. The top horizontal I has a vertical line passing through it, and the bottom horizontal I has a vertical line passing through it. The vertical lines are connected by horizontal lines, forming a grid. The bottom horizontal line is labeled with a vertical ellipsis and a horizontal ellipsis, indicating it continues.

Diagram \tilde{Z} : It is similar to Z , but the top and bottom horizontal lines are labeled with a vertical ellipsis and a horizontal ellipsis, indicating they continue. The vertical identity I is on the left. The horizontal identity I tensors are in the middle and bottom. The top horizontal line has a vertical line passing through it, and the bottom horizontal line has a vertical line passing through it. The vertical lines are connected by horizontal lines, forming a grid. The bottom horizontal line is labeled with a vertical ellipsis and a horizontal ellipsis, indicating it continues.

\tilde{Z}

The Parallel Diagonal Dominant (PDD) Algorithm

Step 1. Allocate $A_i, d^{(i)}$ and elements $a_{im}, c_{(i+1)m-1}$ to the i th node, where $0 \leq i \leq p-1$.

Step 2. Use the LU decomposition method to solve

$$A_i[\tilde{x}^{(i)}, v^{(i)}, w^{(i)}] = [d^{(i)}, a_{im}e_0, c_{(i+1)m-1}e_{m-1}]$$

Step 3. Send $\tilde{x}_0^{(i)}, v_0^{(i)}$ to the $(i-1)$ th node.

Step 4. Solve

$$\begin{pmatrix} w_{m-1}^{(i)} & 1 \\ 1 & v_0^{(i+1)} \end{pmatrix} \begin{pmatrix} y_{2i} \\ y_{2i+1} \end{pmatrix} = \begin{pmatrix} \tilde{x}_{m-1}^{(i)} \\ \tilde{x}_0^{(i+1)} \end{pmatrix}$$


in parallel and send y_{2i+1} to $(i+1)$ th node

Step 5. Compute in parallel on p processors


$$\Delta x^{(i)} = [v^{(i)}, w^{(i)}] \begin{bmatrix} y_{2i-1} \\ y_{2i} \end{bmatrix}$$

$$x^{(i)} = \tilde{x}^{(i)} - \Delta x^{(i)}$$


Orchestration



Computing/Communication of PDD



Non-periodic



Periodic

Orchestration

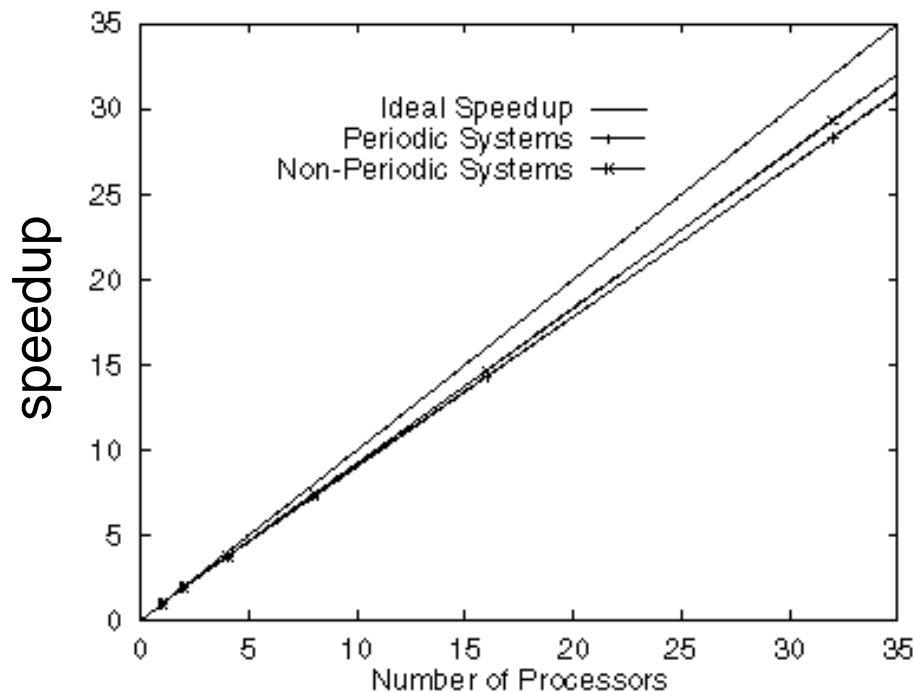
- Orchestration is implied in the algorithm design
- Only two one-to-one neighboring communication

Mapping

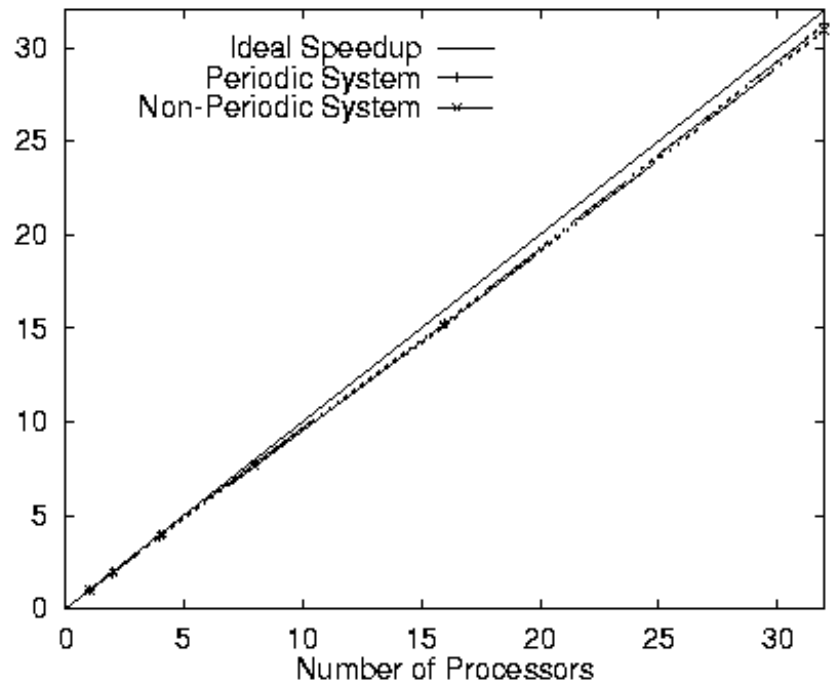
- Communication has reduced
 - Take the special mathematical property
 - Formal analysis can be performed based on the mathematical partition formula
- Two neighboring communication
 - Can be achieved on array communication network

The PDD Algorithm

- Advantage
 - Perfect parallel
 - Constant, minimum communication
- Disadvantage
 - Increased computation (vs. sequential alg.)
 - Applicability
 - Diagonal dominant
 - Subsystems are reasonably large



Scaled Speedup of the PDD Algorithm on Paragon. *1024 System of order 1600, periodic & non-periodic*



Scaled Speedup of the Reduced PDD Algorithm on SP2. *1024 System of Order 1600, periodic & non-periodic*

Problem Description

- For tridiagonal systems we may need new algorithms

Problem Description

- Tridiagonal linear systems


$$AX = D$$

$$A = \begin{pmatrix} b_0 & c_0 & & & \\ a_1 & b_1 & c_1 & & \\ & \ddots & \ddots & \ddots & \\ & & a_{n-2} & b_{n-2} & c_{n-2} \\ & & a_{n-1} & b_{n-1} & \end{pmatrix}$$

$$X = \begin{pmatrix} x_{00} & x_{01} & \cdot & \cdot & x_{0,n-1} \\ x_{10} & x_{11} & x_{12} & \cdot & \cdot \\ \cdot & \ddots & \ddots & \ddots & \cdot \\ \cdot & \cdot & x_{n-2,n-3} & x_{n-2,n-2} & x_{n-2,n-1} \\ x_{n-1,0} & \cdot & \cdot & x_{n-1,n-2} & x_{n-1,n-1} \end{pmatrix}$$

$$D = \begin{pmatrix} d_{00} & d_{01} & \cdot & \cdot & d_{0,n-1} \\ d_{10} & d_{11} & d_{12} & \cdot & \cdot \\ \cdot & \ddots & \ddots & \ddots & \cdot \\ \cdot & \cdot & d_{n-2,n-3} & d_{n-2,n-2} & d_{n-2,n-1} \\ d_{n-1,0} & \cdot & \cdot & d_{n-1,n-2} & d_{n-1,n-1} \end{pmatrix}$$

The Pipelined Method




- Exploit temporal parallelism of multiple systems
- Passing the results from solving a subset to the next before continuing
- Communication is high, $3p$
- Pipelining delay, p
- Optimal computing

The Parallel Two-Level Hybrid Method


- PDD is scalable but has limited applicability
- The pipelined method is mathematically efficient but not scalable
- Combine these two algorithms, outer PDD, inner pipelining
- Can combine with other algorithms too

The Parallel Two-Level Hybrid Method



- Use an accurate parallel tridiagonal solver to solve the m super-subsystems concurrently, each with k processors
- Modify PDD algorithm and consider communications only between the m super-subsystems.

The Partition Pipeline diagonal Dominant (PPD) algorithm



•Evaluation of Algorithms

System	Algorithm	Computation	Communication
Multiple systems	Best Sequential	$8n - 7$	0
	Pipelining	$\frac{(n_1 - 1 + p)(8n - 7)}{p}$	$3(n_1 - 1 + p)(\alpha + 4\beta)$
	PDD	$(17\frac{n}{p} - 14) * n_1$	$(2\alpha + 12 * n_1 * \beta)$
	PPD	$(n_1 - 1 + k) \frac{13n}{p} + 4n_1(\frac{n}{p} + 1)$	$3(2\alpha + 12\beta) + [2 + \log(k)](\alpha + 12n_1\beta)$

Practical Motivation

- NLOM (NRL Layered Ocean Model) is a well-used naval parallel ocean simulation code (see http://www7320.nrlssc.navy.mil/global_nlom/index.html).
- Fine tuned with the best algorithms available at the time
- Efficiency goes down when the number of processors increases.
- Poisson solver is the identified scalability bottleneck

Project Objectives

- Incorporate the best scalable solver, the PDD algorithm, into NLOM
- Increase the scalability of NLOM
- Accumulate experience for a general toolkits solution for other naval simulation codes

Experimental Testing


- Fast Poisson solvers (FACR) (Hockney, 1965)
- One of the most successful rapid elliptic solvers

$$f_q \xrightarrow{FFT} \bar{f}_q^k \xrightarrow[\text{Tridiagonal System}]{\text{Diagonal Dominant}} \bar{\varphi}_q^k \xrightarrow{FFT} \varphi_q$$

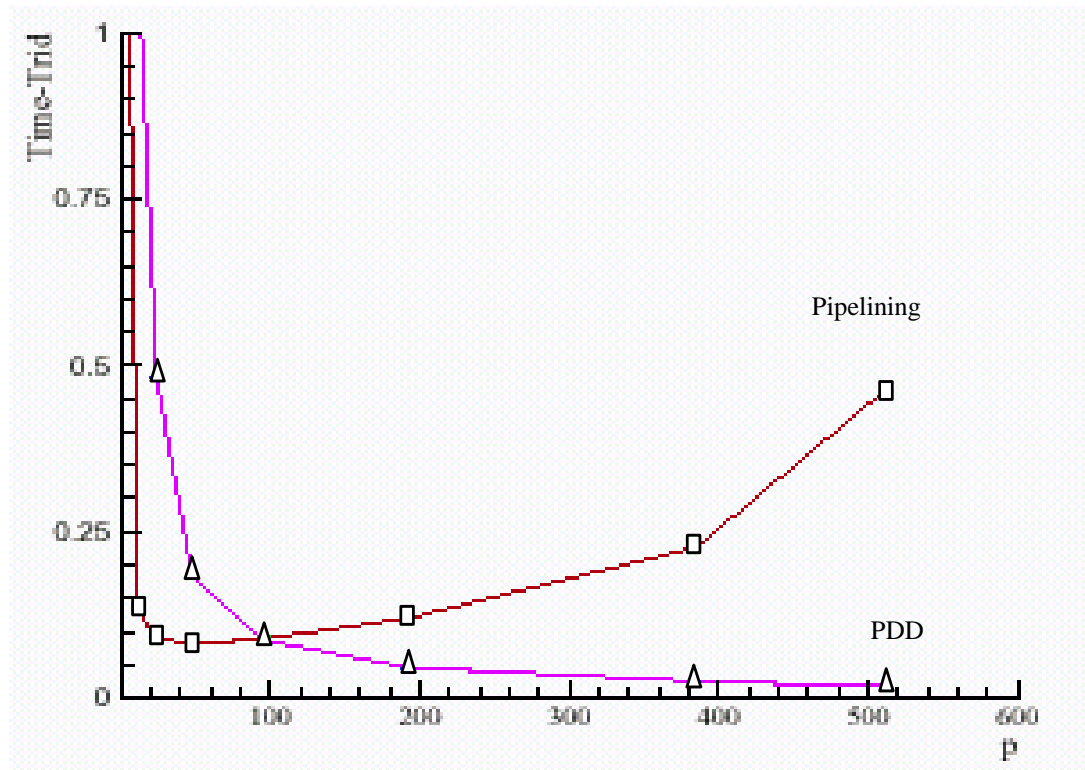
- Large number of systems, each node has a piece of each system
- NLOM implementation, highly optimized pipelining
- Burn At Both Ends (BABE), trade computation with comm. (p, 2p)

NLOM Implementation

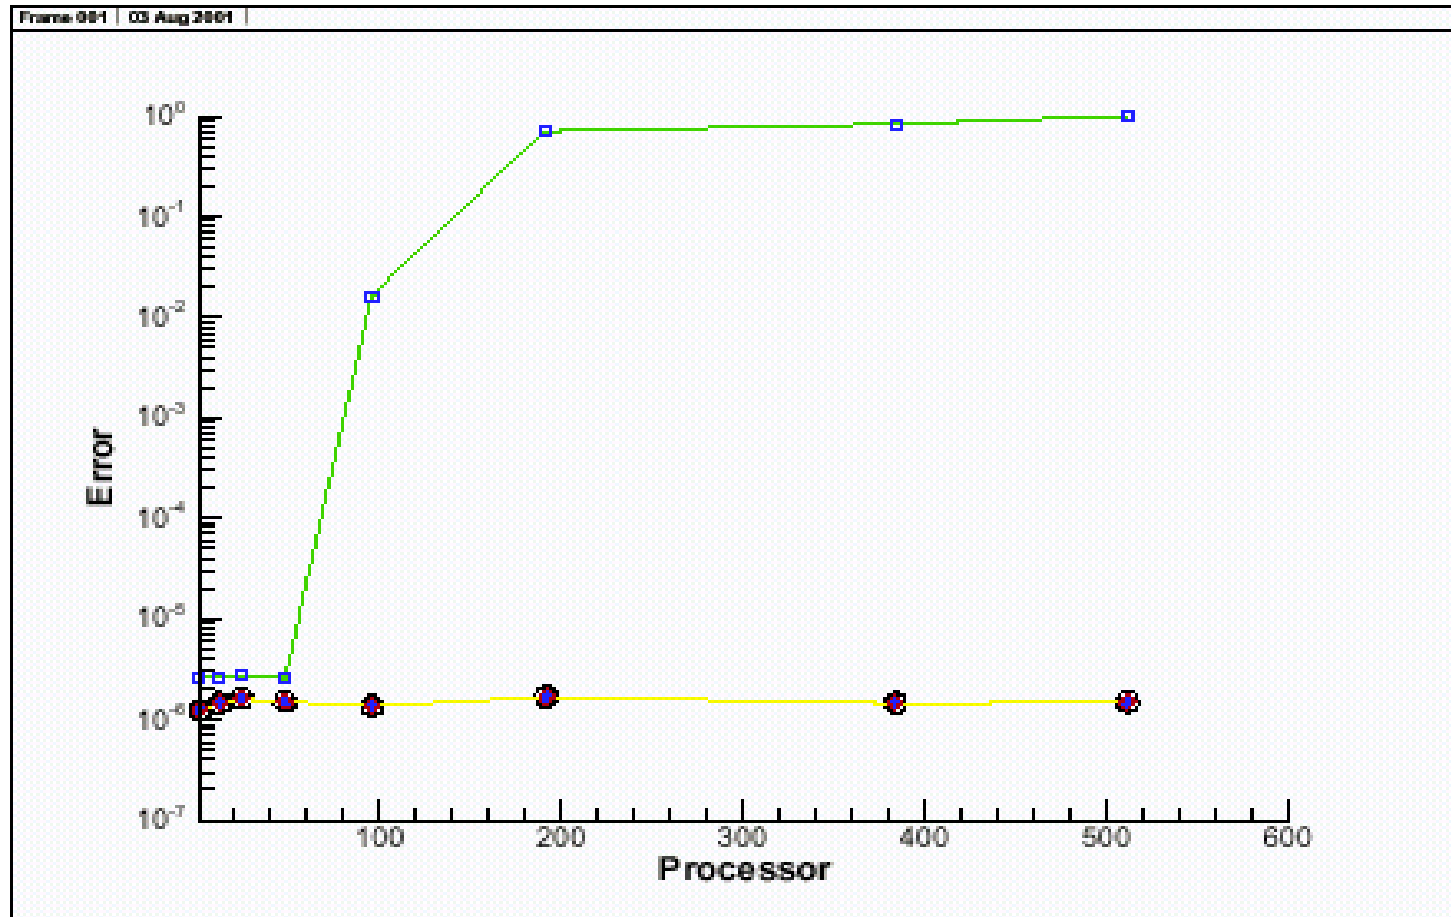
- NLOM has a special data structure and partition
 - Large number of systems, each node has a piece of each system
- Pipelined method, highly optimized
- Burn At Both Ends (BABE), pipelining at both sides, trade computation with comm. (p , $2p$)



Tridiagonal solver runtime: Pipelining (square) and PDD (delta)




Accuracy: Circle - BABE Square - PDD Diamond - PPD



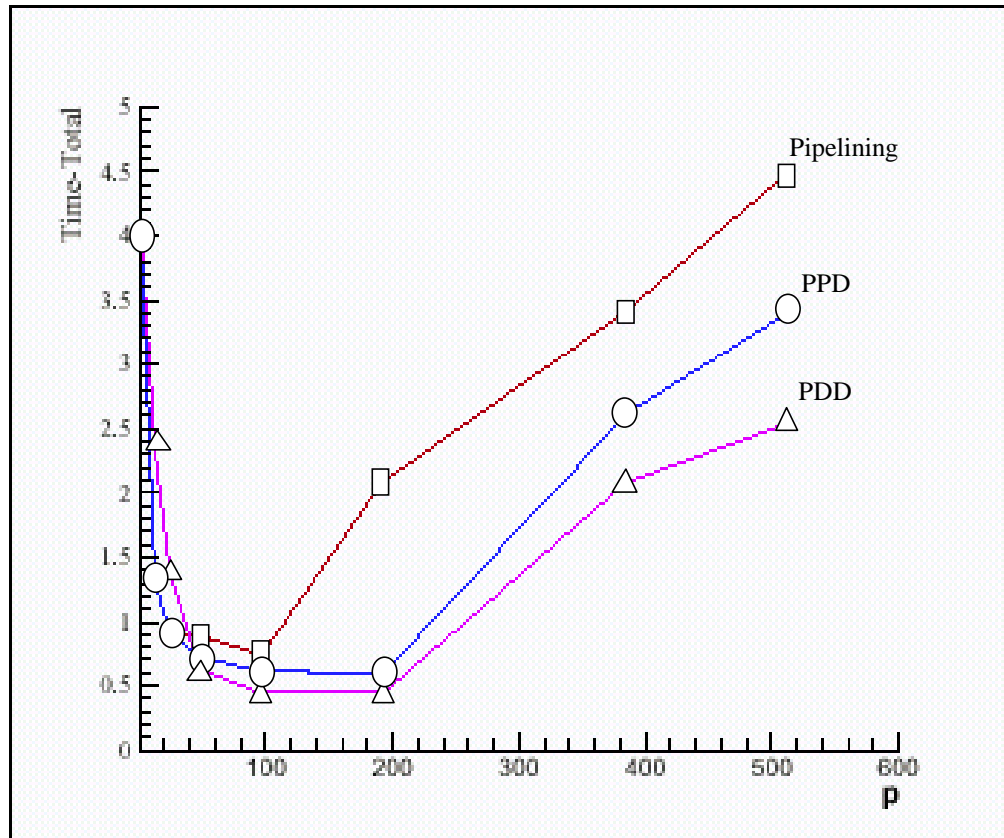
NLOM Application

- Pipelined method is not scalable
- PDD is scalable but loses accuracy, due to the subsystems are very small
- Need the two-level combined method

Trid. Solver Time: Pipelining (square), PDD (delta), PPD (circle)



Total runtime: Pipelining (square), PDD (delta), PPD (circle)



Parallel Two-Level Hybrid (PTH) Method

- Use an accurate parallel tridiagonal solver to solve the m super-subsystems concurrently, each with k processors, where $p = l \cdot k$, and solving three unknowns as given in the *Step 2* of PDD algorithm.
- Modify the solutions of *Step 1* with *Steps 3-5* of PDD algorithm, or of PPT algorithm if PPT is chosen as the outer solver.

The PTH method and related algorithms

Abbreviat ion	Full Name	Explanation
PPT	Parallel ParTition LU Algorithm	A parallel solver based on rank-one modification
PDD	Parallel Diagonal Dominant Alg	A variant of PPT for diagonal dominant system
PTH	Parallel Two-level Hybrid Method	A novel two-level approach
PPD	Partition Pipelined diagonal Dominant Algorithm	A PTH uses PDD and pipelining as outer/inner solver

Perform Evaluation

•Evaluation of Algorithms

Comparison of computation and communication (non periodic)			
System	Algorithm	Computation	Communication
Single system	Best Sequential	$8n - 7$	0
	PPT	$17\frac{n}{p} + 16p - 23$	$(2\alpha + 8p\beta)(\sqrt{p} - 1)$
	PDD	$17\frac{n}{p} - 4$	$2\alpha + 12\beta$
	Reduced PDD	$11\frac{n}{p} + 6j - 4$	$2\alpha + 12\beta$
Multiple system	Best Sequential	$(5n - 3).n1$	0
	PPT	$(9\frac{n}{p} + 10p - 11).n1$	$(2\alpha + 8p.n1.\beta)(\beta - 1)$
	PDD	$(9\frac{n}{p} + 1).n1$	$(2\alpha + 8n1.\beta)$
	Reduced PDD	$(5\frac{n}{p} + 4j + 1).n1$	$(2\alpha + 8n1.\beta)$

Algorithm Analysis:

1. LU-Pipelining

$$(n_1 - 1 + p)[(8n - 7)\frac{\tau_{comp}}{p} + 3(\alpha + 4\beta)]$$

2. The PDD Algorithm

$$(17\frac{n}{p} - 14) * n_1 \tau_{comp} + (2\alpha + 12 * n_1 * \beta)$$

3. The PPD Algorithm

$$(n_1 - 1 + p_1)[\frac{13n}{p} \tau_{comp} + 3(2\alpha + 12\beta)] +$$

$$4n_1(\frac{n}{p} + 1)\tau_{comp} + [2 + \log(p_1)](\alpha + 12n_1\beta)$$

Where

n - the order of each system

n_1 - the number of systems

p - the number of processors

p_1 - the number of processors used
for LU-pipelining

τ_{comp} - the computing speed

α - the communication start time

β - the transmission time

Parameters on IBM Blue Horizon at SDSC

$$\tau_{comp} = 0.01696 \mu s$$

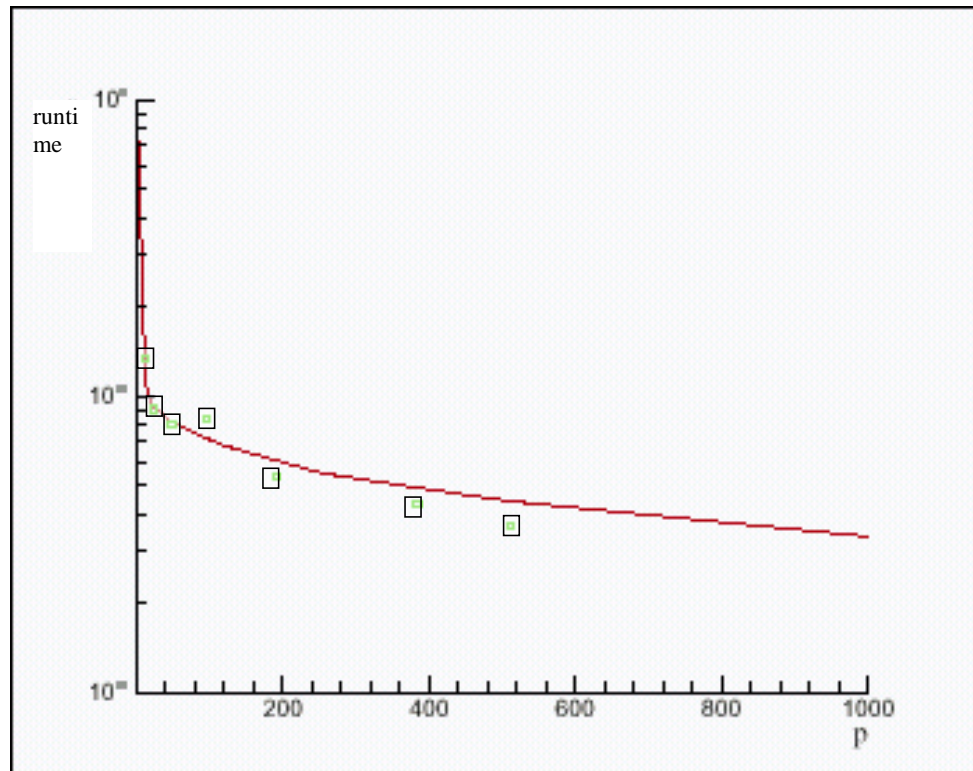
$$\alpha = 31.5 \mu s$$

$$\beta = 9.52 \times 10^{-3} \mu s / byte$$


Computation
and Comm.
Count
(multiple
right sides)

Algorithm	Computation	Communication
Pipelined Algorithm	$(m-1+p)\frac{8n-7}{p}m_1$	$(m-1+p)(3\alpha+12m_1\beta)$
PPT non-pivoting	$\left(17\frac{n}{p}+16p-45\right)n_1$	$((\log(p))\alpha+16(p-1)n_1\beta)$
PDD	$\left(17\frac{n}{p}-14\right)n_1$	$(2\alpha+12n_1\beta)$
PPD PDD/Pipeline	$(m-1+k)\frac{13n}{p}m_1+\left(\frac{4n}{p}+7\right)n_1$	$(m-1+k)(3\alpha+24m_1\beta)+$ $(2+\log(k))\alpha+(8\log(k)+12)n_1\beta$
PPT/ Pipeline	$(m-1+k)\frac{13n}{p}m_1+\left(\frac{4n}{p}+16\frac{p}{k}-23\right)n_1$	$(m-1+k)(3\alpha+24m_1\beta)+$ $(\log(p))\alpha+\left[16\left(\frac{p}{k}-1\right)+8\log(k)\right]n_1\beta$
PDD/PPT	$\left(30\frac{n}{p}+21k-41\right)n_1$	$(2+2\log(k))\alpha+$ $(16(k-1)+8\log(k)+20)n_1\beta$

PPD: The predicted (line) and numerical (square) runtime



Pipelining: The predicted (line) and numerical (square) runtime



Significance

- Advances in massively parallelism, grid computing, and hierarchical data access make performance sensitive to system and problem size
- Scalability is becoming increasingly important
- Poisson solver is a kernel solver used in many naval applications.
- The PPD algorithm provides a scalable solution for Poisson solver
- We also have proposed the general PTH method

Reference

- X.-H. Sun, H. Zhang, and L. Ni, "[Efficient Tridiagonal Solvers on Multicomputers](#)," *IEEE Trans. on Computers*, Vol. 41, No. 3, pp.286-296, March 1992.
- X.-H. Sun, "[Application and Accuracy of the Parallel Diagonal Dominant Algorithm](#)" *Parallel Computing*, August, 1995.
- X.-H. Sun and W. Zhang, "[A Parallel Two-Level Hybrid Method for Tridiagonal Systems, and its Application to Fast Poisson Solvers](#)," *IEEE Trans. on Parallel and Distributed Systems*, Vol. 15, No. 2, pp: 97-106, 2004.
- X.-H. Sun, and S. Moitra, "Performance Comparison of a Set of Periodic and Non-Periodic Tridiagonal Solvers on SP2 and Paragon Parallel Computers," *Concurrency: Practice and Experience*, pp.1-21, Vol.8(10), 1997.
- X.H. Sun, and D. Joslin, "A Simple Parallel Prefix Algorithm for Almost Toeplitz Tridiagonal Systems," *High Speed Computing*, Vol.7, No.4, pp. 547-576, Dec. 1995.
- Y. Zhuang, and X.H. Sun, "[A High Order Fast Direct Solver for Singular Poisson Equations](#)," *Journal of Computational Physics*, Vol. 171, pp. 79-94 (2001).
- Y. Zhuang, and X.H. Sun, "[A High Order ADI Method For Separable Generalized Helmholtz Equations](#)," *International Journal on Advances in Engineering Software*, Vol. 31, pp. 585-592, August 2000.

All the references can be found at

<http://www.cs.iit.edu/~scs/research/scientific-computing.html>