



From Amdahl's Law to Big Data: *A Story of* Computer Sciences and Technology

Xian-He Sun

Illinois Institute of Technology
sun@iit.edu



The Journey of Supercomputing

- The Background of Parallel Processing
 - Speedup
 - Sources of overhead
- The Laws of Scalable Computing
 - The Amdahl's law
 - The Gustafson's law
 - The Sun-Ni's law
- Impacts and Discussions



Performance Evaluation

(Improving performance is the goal)

- Performance Measurement
 - Metric, Parameter
- Performance Prediction
 - Model, Application-Resource
- Performance Diagnose/Optimization
 - Post-execution, Algorithm improvement, Architecture improvement, State-of-the-art, Scheduling, Resource management/Scheduling



Performance of Parallel Processing

Models of Speedup

- Speedup

- T_s = time for the best serial algorithm
- T_p = time for parallel algorithm using p processors




$$S_p = \frac{T_s}{T_p}$$

- Simple enough, but also unexpected complex

$$S_p = \frac{\text{Uniprocess or Execution Time}}{\text{Parallel Execution Time}}$$



Example



$$S_p = \frac{100}{25} = 4.0,$$


perfect parallelization

$$S_p = \frac{100}{35} = 2.85,$$


perfect load balancing
but synch cost is 10



Example (cont.)



(d)



(e)

$$S_p = \frac{100}{40} = 2.5,$$

no synch

but load imbalance

$$S_p = \frac{100}{50} = 2.0,$$

load imbalance

and synch cost



What Is “Good” Speedup?

- *Linear* speedup:

$$S_p = p$$

- *Superlinear* speedup


$$S_p > p$$

- *Sub-linear* speedup:

$$S_p < p$$



Speedup





Sources of Parallel Overheads

- Interprocessor communication
- Load imbalance
- Synchronization
- Extra computation




Causes of Superlinear Speedup

- Cache size increased
- Overhead reduced
- Latency hidden
- Randomized algorithms
- Mathematical inefficiency of the serial algorithm
- Higher memory access cost in sequential processing


• X.H. Sun, and J. Zhu, "Performance Considerations of Shared Virtual Memory Machines," *IEEE Trans. on Parallel and Distributed Systems*, Nov. 1995




Degradations of Parallel Processing




Unbalanced Workload



Communication Delay



Overhead Increases with the Ensemble Size




Overheads


- *communication*
- *Load imbalance*
- *Synchronization*
- *Extra computation*




Degradations of Distributed Computing




Unbalanced **Computing Power** and Workload



Shared Computing and Communication Resource




Uncertainty, Heterogeneity, and Overhead Increases with the Ensemble Size





Principals of Architecture Design

- Make common case fast (90/10 Rule)
- Amdahl's Law
 - Law of diminishing returns
- Speedup
 - Achieved performance improvement over original



Gene Amdahl

$$\text{Speedup Overall} = \frac{\text{speed new}}{\text{speed old}} = \frac{\text{execution time old}}{\text{execution time new}}$$

Here performance is measured in **Speed**



Amdahl's Law

Execution time of any code has two portions

Portion I: not affected by enhancement

Portion II: affected by enhancement

$$\text{execution time}_{\text{old}} = \text{execution time}_{p1} + \text{execution time}_{p2}$$

α is % of original code that cannot benefit from enhancement

As $p \rightarrow \text{infinity}$, $\text{execution time}_{\text{new}} \rightarrow \alpha * \text{execution time}_{\text{old}}$

$$\text{execution time}_{\text{new}} = (\alpha) * \text{execution time}_{\text{old}} + (1 - \alpha) * \frac{\text{execution time}_{\text{old}}}{p}$$

Execution time_{p1} *Execution time_{p2}*

p is speedup factor of old/new execution times for portion II



Amdahl's Law for Parallel Processing (1967)

- Let α = fraction of program (algorithm) that is serial and cannot be parallelized. For instance:
 - Loop initialization
 - Reading/writing to a single disk
 - Procedure call overhead
- Parallel run time is given by

$$\text{execution time}_{\text{new}} = (\alpha) * \text{execution time}_{\text{old}} + (1 - \alpha) * \frac{\text{execution time}_{\text{old}}}{p}$$

$$T_p = \left(\alpha + \frac{1 - \alpha}{p} \right) \bullet T_s$$

Gene M Amdahl, "Validity of the single processor approach to achieving large scale computing capabilities," AFIPS spring joint computer conference, 1967



Amdahl's Law

- Amdahl's law gives a limit on speedup in terms of α

$$S_p = \frac{T_s}{T_p} = \frac{T_s}{\alpha T_s + \frac{(1-\alpha)T_s}{p}} = \frac{1}{\alpha + \frac{1-\alpha}{p}}$$

- If we assume that the serial fraction is fixed, then the speedup for infinite processors is limited by $1/\alpha$

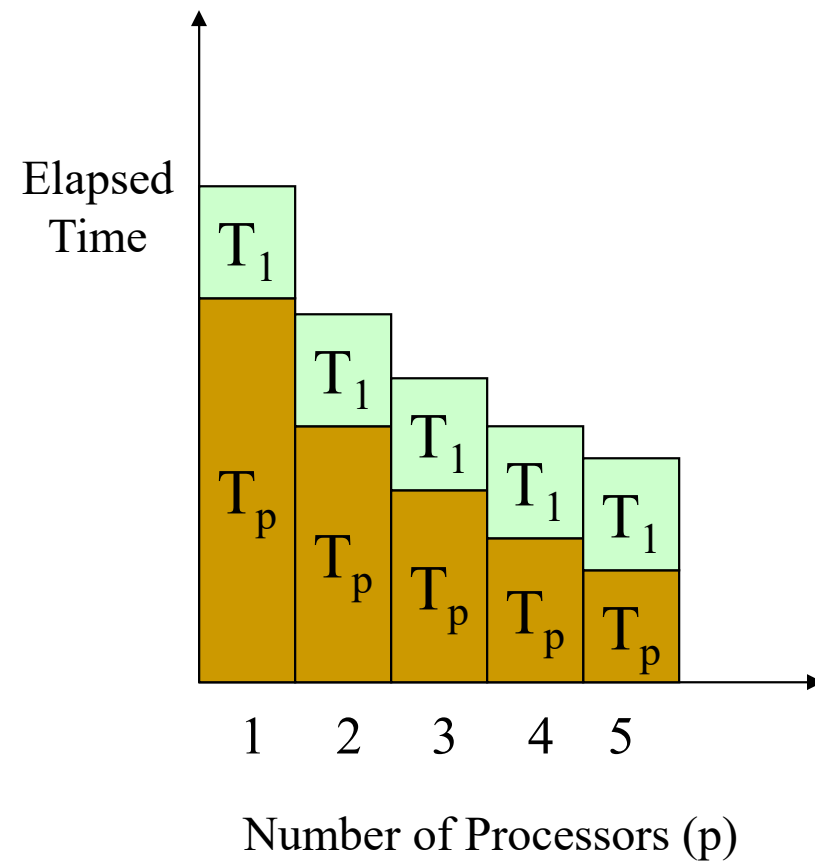
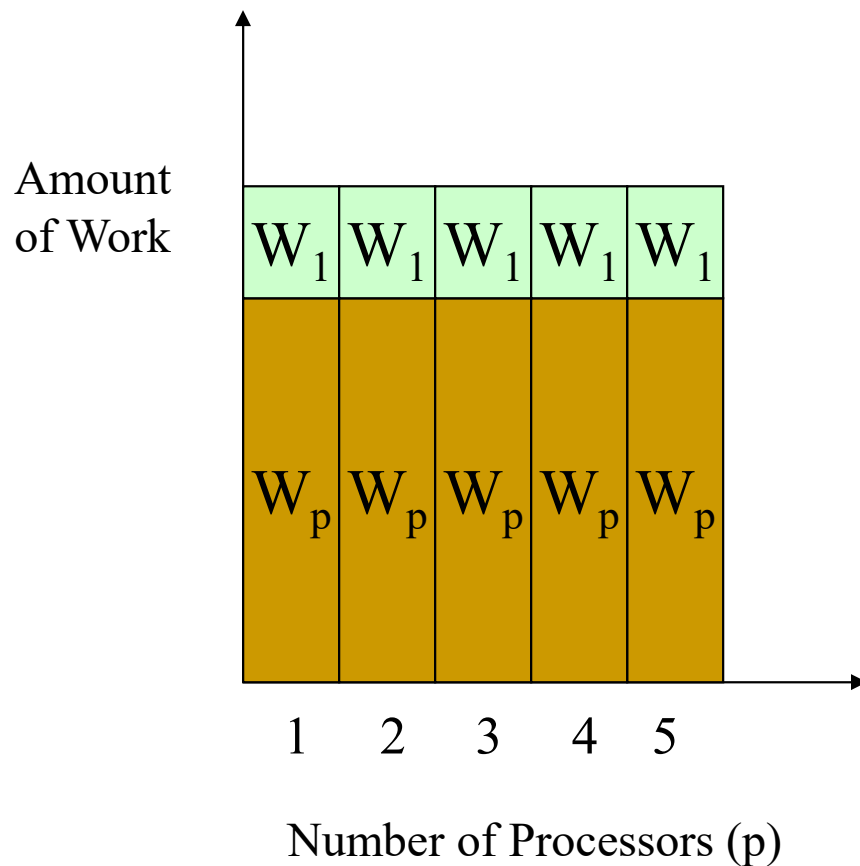
$$\lim_{p \rightarrow \infty} S_p = \frac{1}{\alpha}$$

- For example, **if $\alpha=10\%$** , then the maximum speedup is **10**, even if we use an infinite number of processors



Amdahl Law

- The sequential part becomes the dominate factor quickly







Amdahl's Law

$$\text{execution time}_{\text{new}} = (\alpha) * \text{execution time}_{\text{old}} + (1 - \alpha) * \frac{\text{execution time}_{\text{old}}}{p}$$

Example: alpha = 20%



$$\text{Speedup}_{\text{overall}} = \frac{\text{execution time}_{\text{old}}}{\text{execution time}_{\text{new}}} = \frac{1}{(\alpha) + \frac{1 - \alpha}{p}}$$





Amdahl's Law with Overhead

- To include overhead will be even worse
- The overhead includes parallelism and interaction overheads

$$Speedup_{FS} = \frac{T_1}{\alpha T_1 + \frac{(1-\alpha)T_1}{p} + T_{overhead}} \rightarrow \frac{1}{\alpha + \frac{T_{overhead}}{T_1}} \text{ as } p \rightarrow \infty$$

Amdahl's law: argument against massively parallel systems



Example

- Enhanced mode is used 50% of resulting execution time_{new}
- Portions of code using enhanced mode improve performance by factor 10x
- What is speedup for fast mode?

First find $1-\alpha$ (% of code affected by enhancement)

$$\text{execution time}_{\text{new}} = \underbrace{(.5) * \text{execution time}_{\text{new}}}_{\text{not enhanced}} + \underbrace{(.5) * \text{execution time}_{\text{new}}}_{\text{enhanced}}$$

$$(\alpha) * \text{execution time}_{\text{old}} = \frac{(1 - \alpha) * \text{execution time}_{\text{old}}}{10}$$

$$1 - \alpha = .909091$$

$$\text{Speedup}_{\text{overall}} = \frac{1}{(\alpha) + \frac{1 - \alpha}{n}} = \frac{1}{(1 - .9) + \frac{.9}{10}} = 5.3$$



Example

- Enhancement: Parallel processing with 20 computing nodes
- Portions of code containing computations run 20x faster in parallel processing mode.
- What % of original code must be parallelizable to achieve $\text{speedup}_{\text{overall}} = 2$?

$$\text{Speedup}_{\text{overall}} = \frac{\text{execution time}_{\text{old}}}{\text{execution time}_{\text{new}}} = \frac{1}{(\alpha) + \frac{1-\alpha}{n}}$$

$$2 = \frac{1}{(\alpha) + \frac{1-\alpha}{20}}$$

$$1 - \alpha = .5263$$



History back to 1988



IBM 7030 Stretch



IBM 7950 Harvest



Cray X-MP
Fastest computer 1983-1985



Cray Y-MP

All have up to 8
processors, citing
Amdahl's law,

$$\lim_{p \rightarrow \infty} Speedup_{Amdahl} = \frac{1}{\alpha}$$



Gene Amdahl



Summit: the World Fastest Computer



- 148.6 petaflops (187.66 petaflop theoretical peak)
- 2,282,544 IBM Power 9 core
- 2,090,880 Nvidia Volta GV100 core
- Power efficiency 11.324gigaflop/watt



Bombshell: *Gustafson, etc. Got Speedup of more than 1,000 on Three Applications*

- On a 1024-processor nCUBE parallel computer
- For three applications: wave mechanics, fluid dynamics, and structural analysis.
- Introduced the concept of **Scalable Computing**, *problem size increases with the machine size*

John L. Gustafson, Gary R. Montry, and Robert E. Benner, "Development of Parallel Methods for a 1024-Processor Hypercube," SIAM Journal on Scientific and Statistical Computing, Vol. 9, No.4, 1988 (submitted 3/10/1988, accepted 3/25/1988, appeared April 1988)

John Gustafson, "Reevaluation of Amdahl's Law," Communications of the ACM, Vol. 31, No. 5, May 1988.



Reevaluate Amdahl's Law

- **Amdahl's Law** is designed for technology improvement, but has been widely used to against parallel processing in terms of reducing execution time
- **But:** large computers are not (only) designed for solving existing problem faster, they are designed for solving otherwise unsolvable large problems
- The introduction of **scalable computing**, where *problem size increases with the machine size*



- Fixed-Time Speedup (Gustafson, 88)

- Emphasis on work finished in a fixed time
- Problem size is **scaled** from W to W'
- W' : Work finished within the fixed time with parallel processing



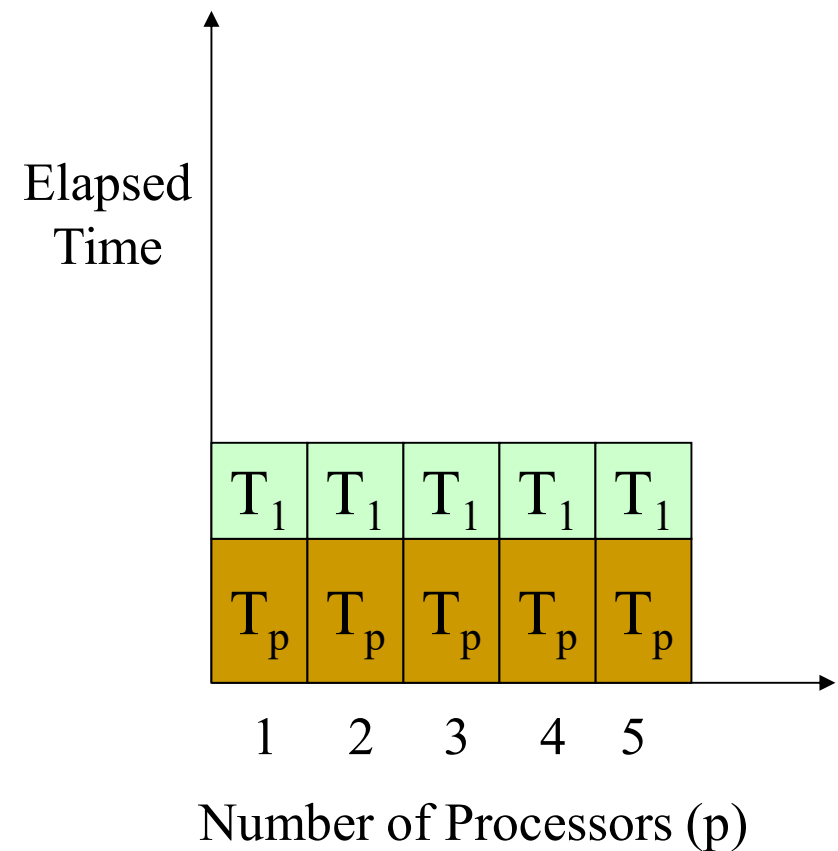

John L. Gustafson

$$\begin{aligned} S'_p &= \frac{\text{Uniprocessor Time of Solving } W'}{\text{Parallel Time of Solving } W'} \\ &= \frac{\text{Uniprocessor Time of Solving } W'}{\text{Uniprocessor Time of Solving } W} \\ &= \frac{W'}{W} \end{aligned}$$



Fixed-Time Speedup (Gustafson)



- Solving a larger application within the time limit





Reexam Amdahl Law (Fixed-Size Speedup)

- It is on time reduction for solving a fixed problem (size)





- Amdahl's law (Fixed-Size Speedup)
 - Emphasis on turnaround time
 - Problem size, W , is fixed


$$S_p = \frac{\text{Uniprocessor Execution Time}}{\text{Parallel Execution Time}}$$

$$S_p = \frac{\text{Uniprocessor Time of Solving } W}{\text{Parallel Time of Solving } W}$$



Gustafson's Law (Without Overhead)

- Under **Gustafson's Law** the parallel processing part is changing with the number of processors, p , and problem size
- Linear speedup



$$\alpha = \frac{t_s}{t_s + t_p}$$

$$Speedup_{FT} = \frac{Work(p)}{Work(1)} = \frac{\alpha W + (1 - \alpha)pW}{W} = \alpha + (1 - \alpha)p$$

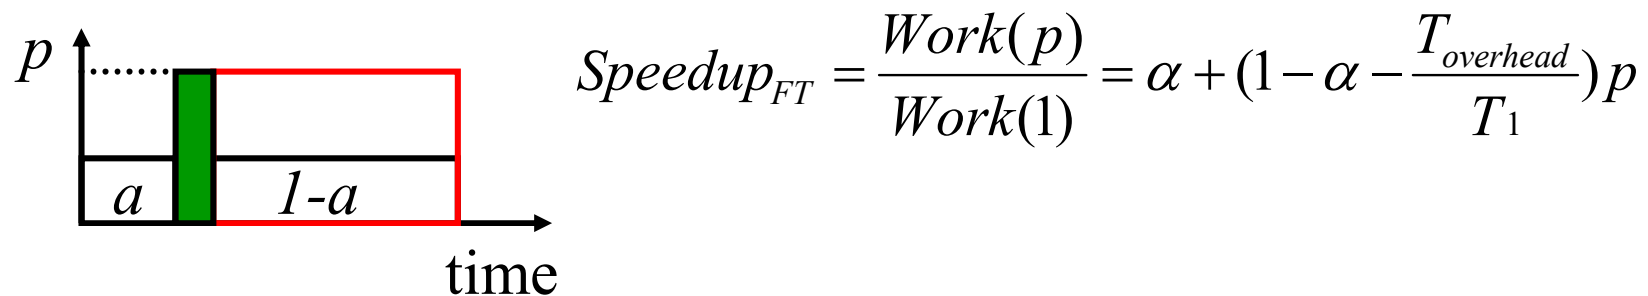
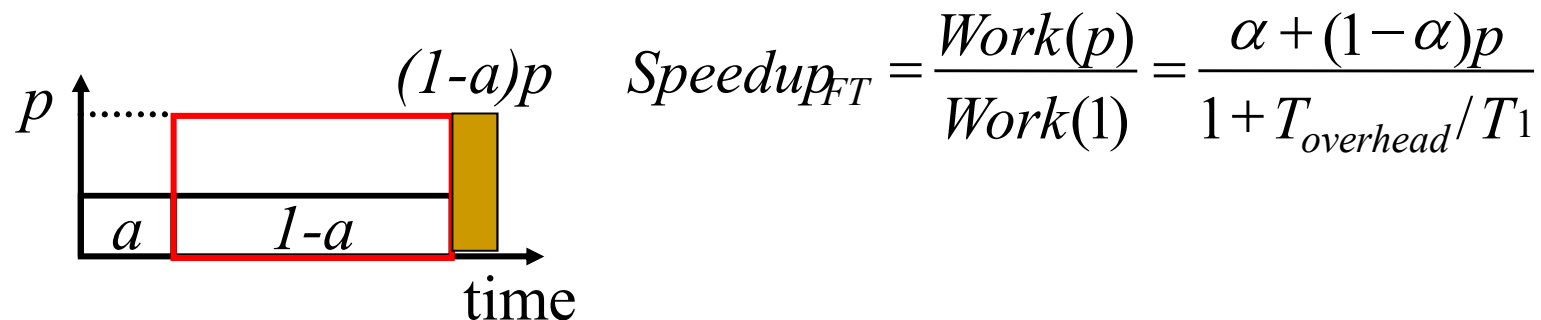
If $\alpha=0.1$

$$Speedup_{FT} = \alpha + (1 - \alpha)p = 0.1 + 0.9p$$



Gustafson's Law (With Overhead)

- In practice, the **overhead** can increase with P, and **limit** the scalability





But: *Gustafson's Applications are not Scalable*

- Most applications cannot get more than 1,000 speedup on a 1024-processor nCUBE parallel computer

Parallel Processing overhead

- Even the three applications are not **Scalable** (increase *problem size further does not help*)

Why?



Memory Constrained Scaling:

Sun and Ni's Law

- **Scaling is limited by memory space** (disk will increase overhead significantly), e.g. fixed memory capacity/usage per processor
 - (ex) N-body problem
- Problem size is scaled from W to W^* , W^* is the work executed under memory limitation
- The relation between memory & computing requirement is determined by the underlying algorithm/program
- **Memory-scaling function**

$$W^* = G(p * M)$$

X.H. Sun, and L. Ni , "Scalable Problems and Memory-Bounded Speedup," *Journal of Parallel and Distributed Computing*, Vol. 19, pp.27-37, Sept. 1993 (**SC90**).




Xian-He Sun

Sun & Ni's Law

存储受限理论



Lionel M. Ni



$$Speedup_{MB} = \frac{Work(p) / Time(p)}{Work(1) / Time(1)} = \frac{\alpha + (1 - \alpha)G(p)}{\alpha + (1 - \alpha)G(p) / p}$$

Assuming $\alpha = 0.1$, the problem needs $2n^3$ computation and $3n^2$ memory
Then $G(p) = G(p) = p^{\frac{3}{2}}$, and

$$Speedup_{MB} = \left(0.1 + 0.9 \times p^{\frac{3}{2}}\right) / \left(0.1 + (0.9 \times p^{\frac{3}{2}}) / p\right)$$



Memory-Bounded Speedup 存储受限理论

(Sun & Ni, 90)

- Emphasis on work finished under current physical limitation
 - Problem size is scaled from W to W^*
 - W^* : Work executed under memory limitation with parallel processing

$$S_p^* = \frac{\text{Uniprocessor Time of Solving } W^*}{\text{Parallel Time of Solving } W^*}$$





Memory-Bounded Speedup (Sun & Ni)

- In practice, memory-bounded performs better than fixed-time but both hard to achieve linear speedup

$$Speedup_{MB} = \frac{Work(p)/Time(p)}{Work(1)/Time(1)} = \frac{\alpha + (1 - \alpha)G(p)}{\alpha + \frac{(1 - \alpha)G(p)}{p} + overhead(p, G(p))}$$

$$Speedup_{FT} = \frac{Work(p)}{Work(1)} = \alpha + (1 - \alpha - \frac{T_{overhead}}{T_1})p$$





Example of calculating Sun & Ni's Law

$$Speedup_{MB} = \frac{Work(p) / Time(p)}{Work(1) / Time(1)} = \frac{\alpha + (1 - \alpha)G(p)}{\alpha + (1 - \alpha)G(p) / p}$$

■ Example:

- Assuming that the problem needs $2n^3$ computation and $3n^2$ memory requirement
- $\alpha = 0.6$
- $p = 8$
- $G(p) = G(8) = 8^{\frac{3}{2}} \approx 22.63$
- $Speedup_{MB} = (0.6 + 0.4 \times 22.63) / (0.6 + (0.4 \times 22.63) / 8) \approx 5.576$



Rethinking of Speedup

- Speedup

$$S_p = \frac{\textit{Uniprocessor ExecutionTime}}{\textit{Parallel ExecutionTime}}$$



- It is only the true speedup if problem size is fixed, but now we have scalable computing
- Generalized speedup

$$S_p = \frac{\text{Parallel Speed}}{\text{Sequential Speed}}$$

X.H. Sun, and J. Gustafson, "Toward A Better Parallel Performance Metric," *Parallel Computing*, Vol. 17, pp.1093-1109, Dec. 1991.



Performance Example

- Performance is measured in **Speed**
- But since they work on the same of work, Speedup equals Time Reduction
- My car (X) travels a distance of 1 in one hour
 - Time between start and completion of event is 1 hour
 - Execution time
- Your car (Y) travels a distance of 1 in two hours
- Intuition: my car is twice as fast as your car



Performance Example

- My car (X) travels a distance of 1 in one hour
 - Time between start and completion of event is 1 hour
 - Execution time
- Your car (Y) travels a distance of 1 in two hours
- Intuition: my car is twice as fast as your car
- Intuition assumes performance = speed



Is intuition correct?

$$\frac{\text{execution time of your car (Y)}}{\text{execution time of my car (X)}} = n$$

Implies X performs n times better than Y

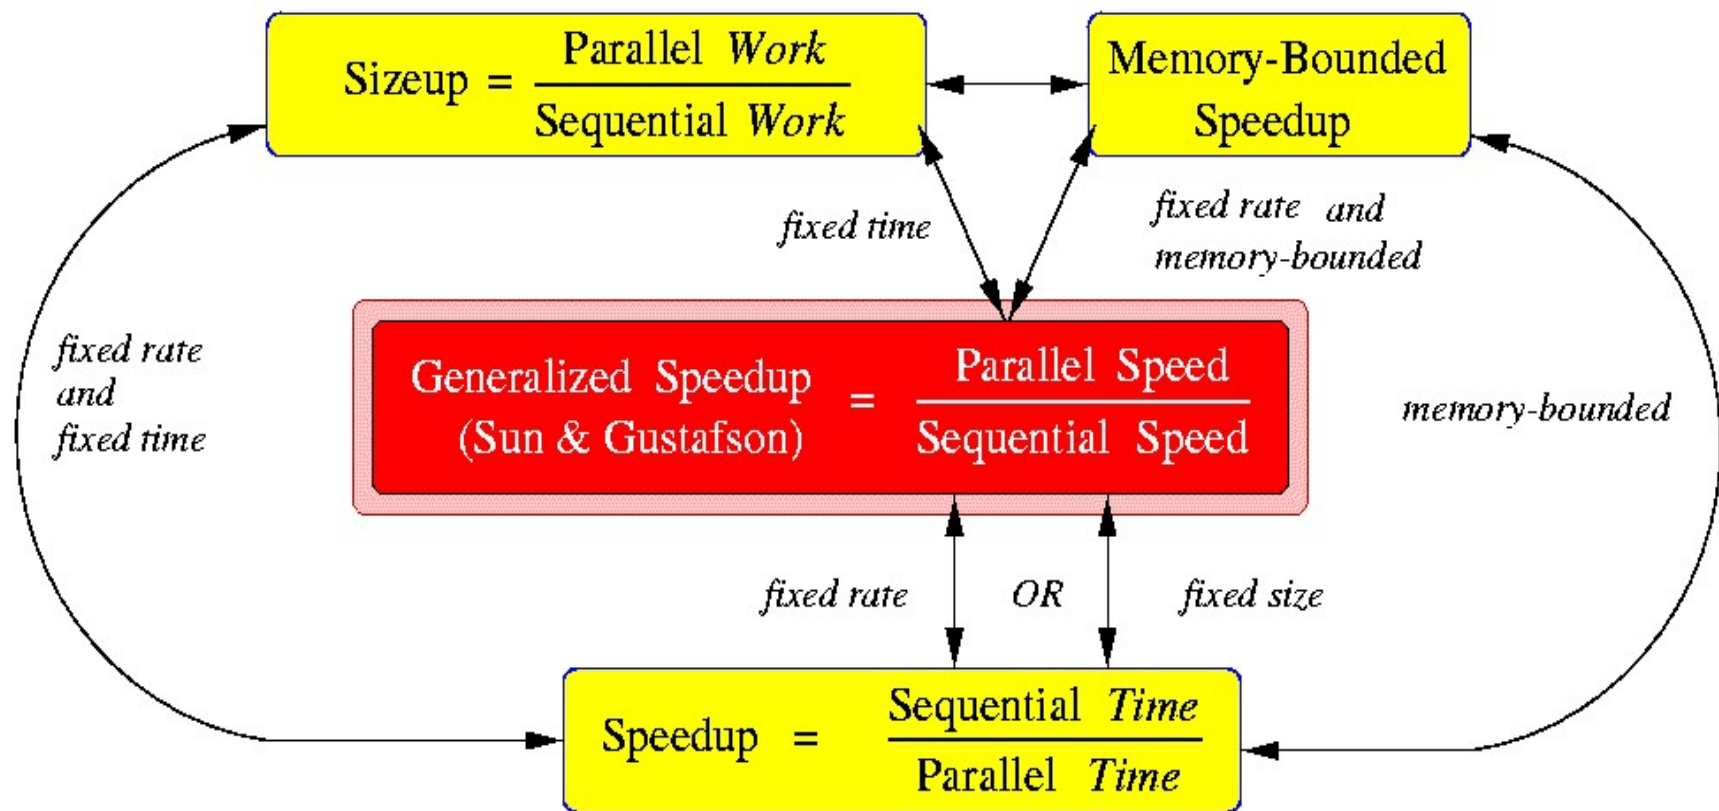
$$\text{performance} = \text{speed} = \frac{1}{\text{execution time}}$$

$$n = \frac{\text{execution time of Y}}{\text{execution time of X}} = \frac{1/\text{speed of Y}}{1/\text{speed of X}} = \frac{\text{speed of X}}{\text{speed of Y}} = \frac{\text{performance of X}}{\text{performance of Y}}$$

Speed is one measure of performance, throughput is another



Models of Speedup






The Three Laws


■ Tacit assumption of Amdahl's law

- Problem size is **fixed**
- Speedup emphasizes on **time reduction**



■ Gustafson's Law, 1988


- Fixed-time speedup model



$$\begin{aligned} \text{Speedup}_{\text{fixed-time}} &= \frac{\text{Sequential Time of Solving Scaled Workload}}{\text{Parallel Time of Solving Scaled Workload}} \\ &= \alpha + (1 - \alpha)p \end{aligned}$$

■ Sun and Ni's law, 1990

- Memory-bounded speedup model



$$\begin{aligned} \text{Speedup}_{\text{memory-bound}} &= \frac{\text{Sequential Time of Solving Scaled Workload}}{\text{Parallel Time of Solving Scaled Workload}} \\ &= \frac{\alpha + (1 - \alpha)G(p)}{\alpha + (1 - \alpha)G(p)/p} \end{aligned}$$

X.-H. Sun, and L. Ni, "Another View of Parallel Speedup," *Proc. of IEEE Supercomputing'90*, NY, NY, Nov.12--Nov.16, 1990.



The Three Laws: and their impact

- **Amdahl's law** (1967) shows the inherent limitation of parallel processing
- **Gustafson's law** (scalable computing, 1988) shows there is no inherent limitation for scalable parallel computing, except engineering issues
- **Sun-Ni's law** (memory-bounded, 1990) shows memory (data) is the constraint of scalable computing (**the** engineering issue)
- The **Memory-Wall Problem** (1994) shows memory-bound is a general performance issue for computing, not just for parallel computing

*I can improve
Amdahl's law*

*I have a
huge
memory*

William Wulf, Sally Mckee, "Hitting the memory wall: implications of the obvious," ACM SIGARCH Computer Architecture News Homepage archive, Vol. 23 Issue 1, March 1995



Impact of Scalable Computing

