

# Class Schedule

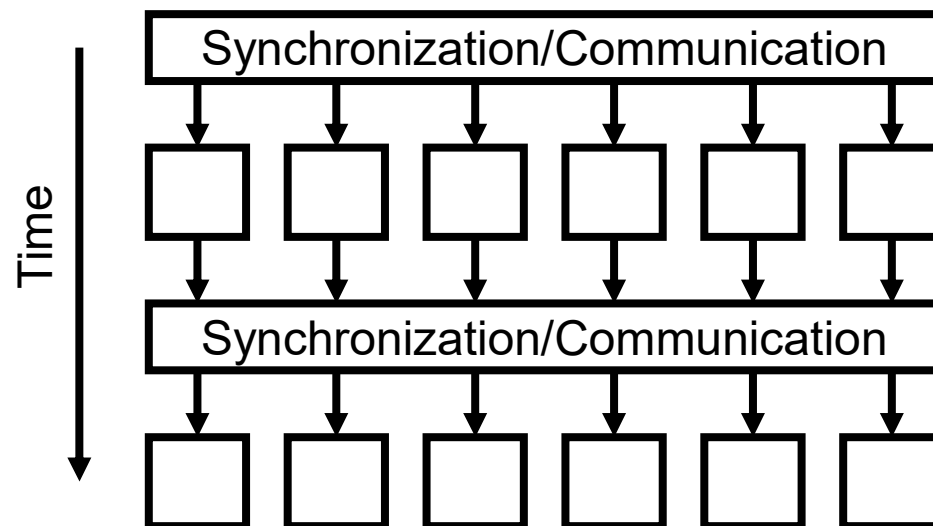
- Review Nov. 7, 2018
- Nov. 12, Regular class, Exam Monday, Nov. 14, 2018
- The week of Nov. 18, no class, individual prepare term project
- Term project presentation: Monday Nov. 25 during and after class hour., Tuesday, Nov. 26 (back up)
- Term report due: Nov. 26

# Application Structure

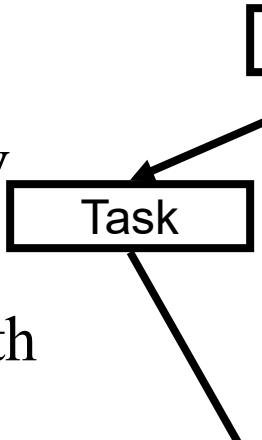
- Frequently used patterns for parallel applications:
  - o Single Program Multiple Data – SPMD (Domain Decomposition)
  - o Embarrassingly Parallel
  - o Master / Slave
  - o Work Pool
  - o Divide and Conquer
  - o Pipeline
  - o Competition

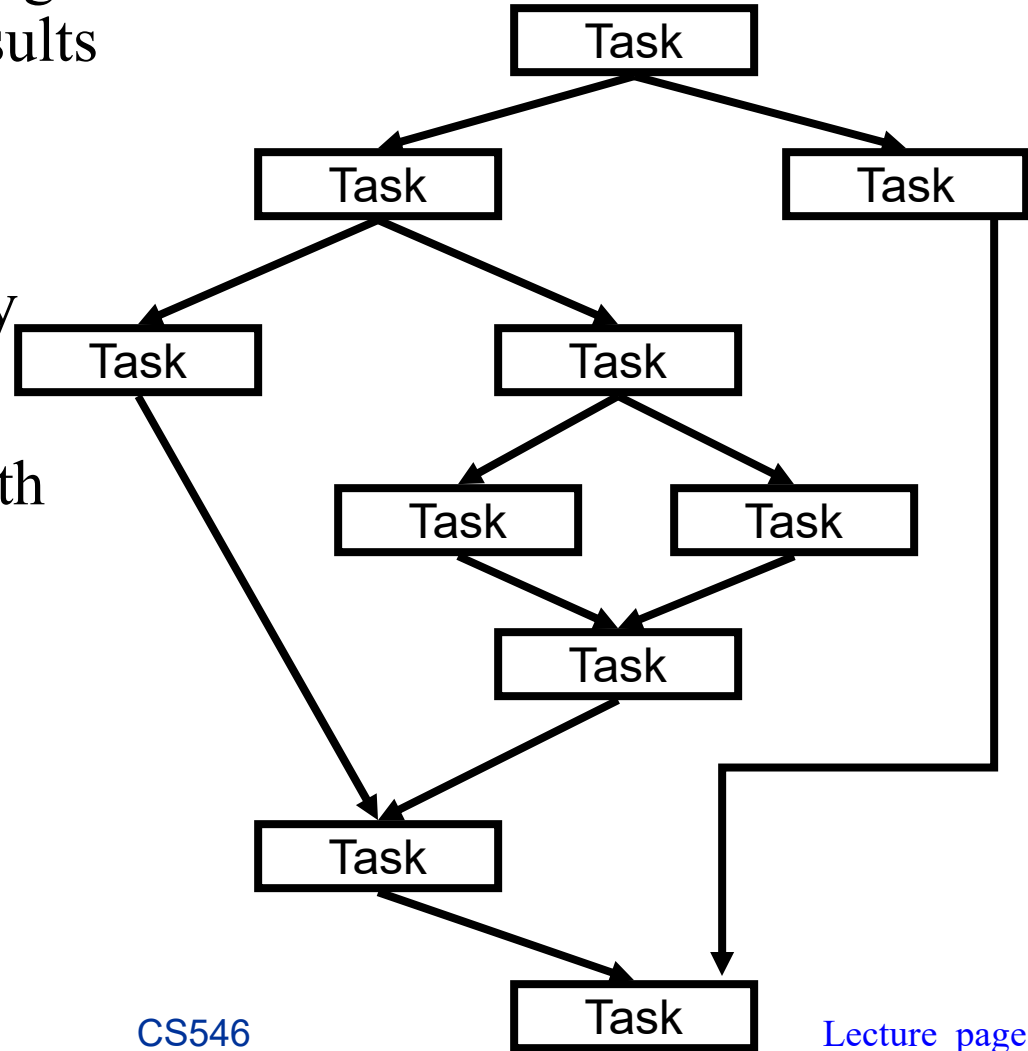
# Structure: Single Program Multiple Data

- Single program is executed in a replicated fashion.
- Processes or threads execute same operations on different data.
- Loosely-synchronous: Sequence of phases of computation and communication/synchronization.

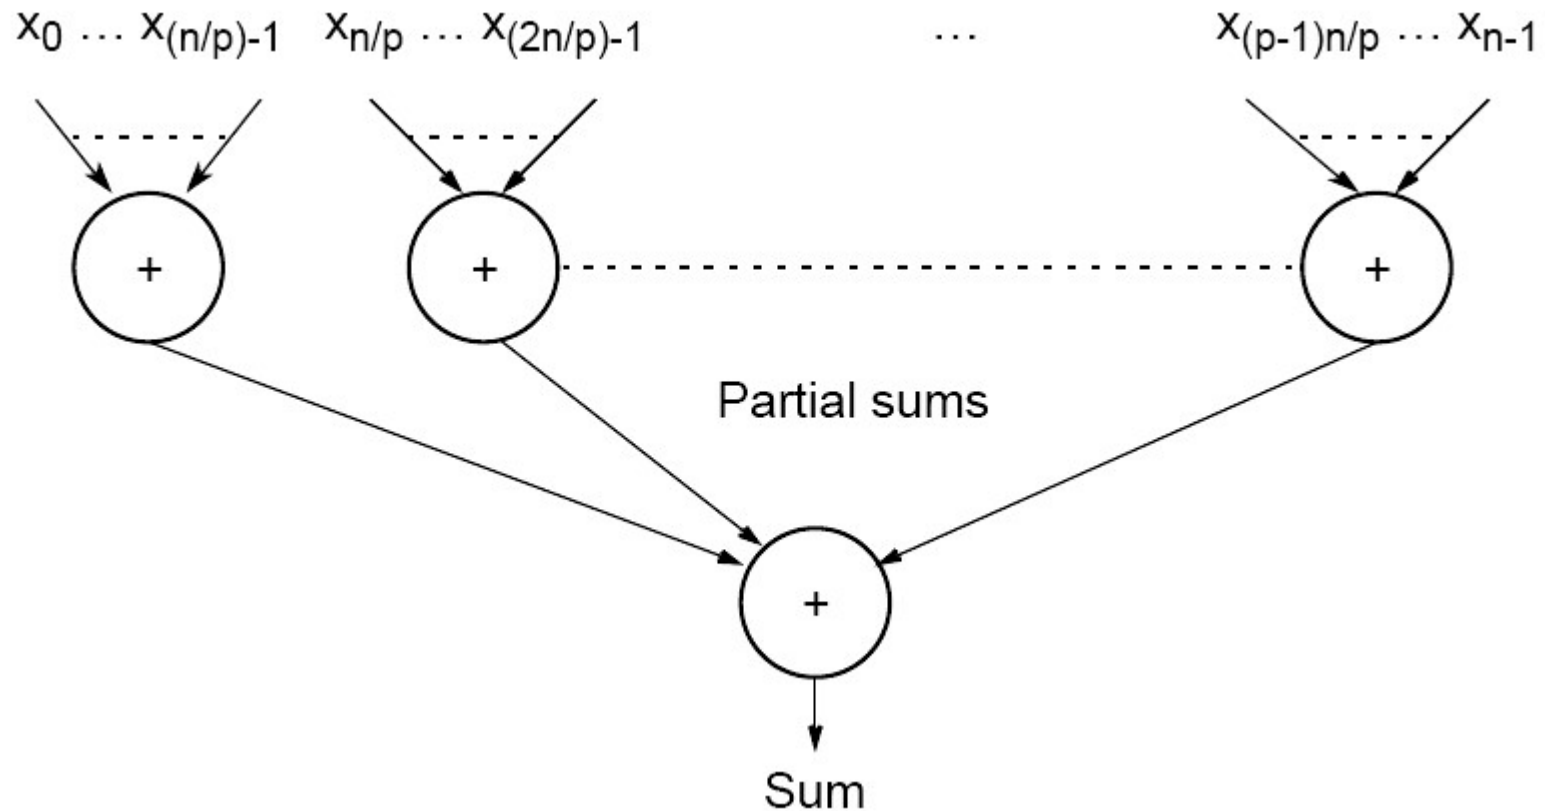


# Structure: Divide and Conquer

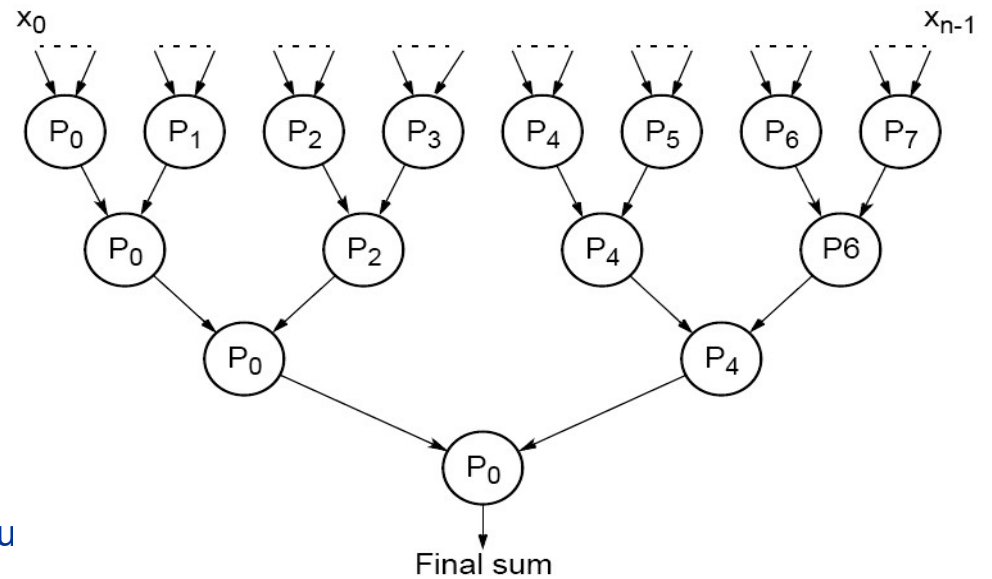
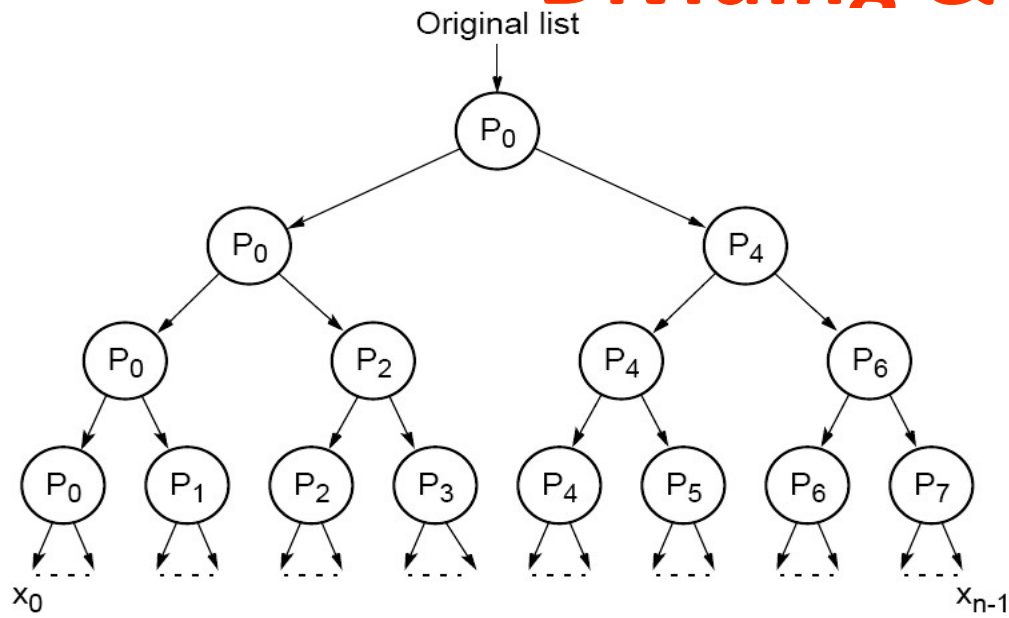
- Recursive partitioning of tasks and collection of results
  - Problems:
    - o load balancing
    - o least granularity
  - Used on systems with background load
- 
- ```
graph TD; In[ ] --> Task[Task]; Task --> Out[ ]
```



# Adding Numbers



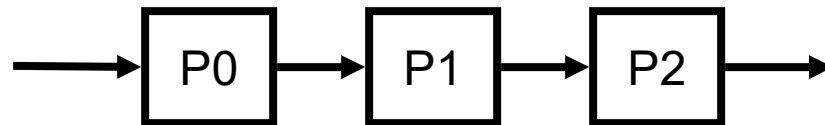
# Dividing & Conquer



# Structure: Pipeline

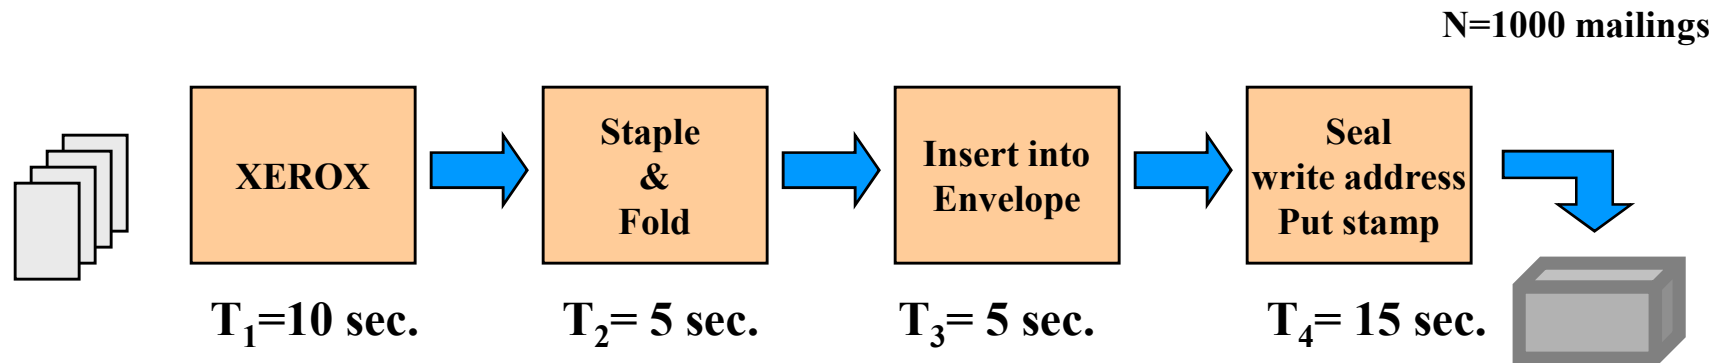
## ➤ Examples

- o Different functions are applied to data: functional decomposition
- o Parallel execution of functions for different data.
- o Signal and image processing
- o Groundwater flow, flow of pollutants, visualization
- o Almost no example of high parallelism



# Pipelining: Example

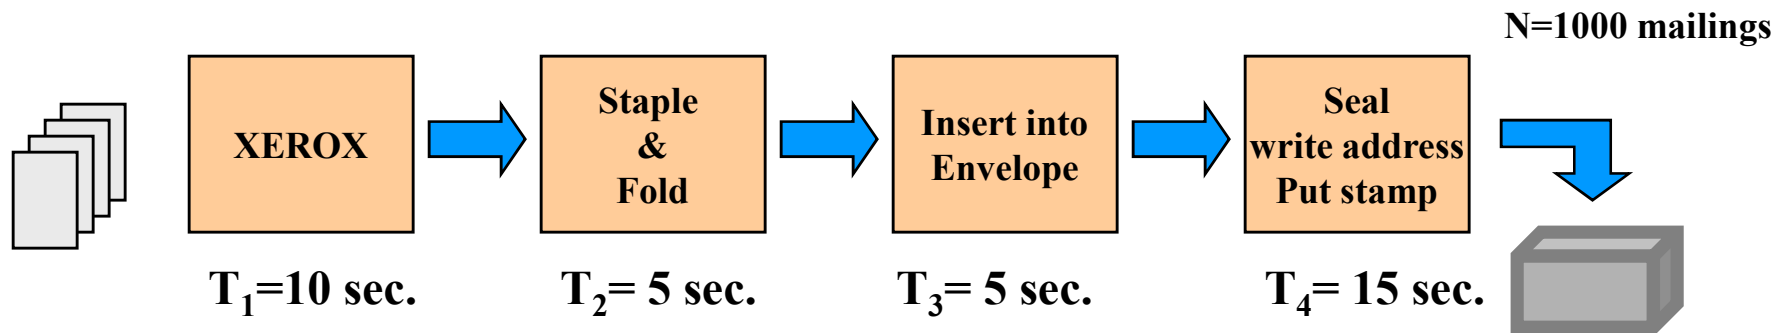
We would like to prepare and mail 1000 envelopes each containing a document of 4 pages to members of an association.



- At what intervals, do we see a new envelope prepared for mailing?  
 $\text{Max}(T_1, T_2, \dots, T_k) = T_{\max} = 15 \text{ sec.}$
- What is the total time to get N envelopes prepared?  
 $\text{Time} = \text{Cold\_Start\_Time} + T_{\max} * (N-1) \cong N * T_{\max}$
- What is the total time we would have spent if pipelining is not used?  
 $N * \sum_i T_i$



## Pipelining: Example (contd.)



How much **speedup** do we get?

$$\text{Speedup} = T_{\text{seq}}/T_{\text{pipe}} = [N * \sum_i T_i] / N * T_{\text{max}} = \sum_i T_i / T_{\text{max}}$$

$$\text{Speedup} = 35/15$$

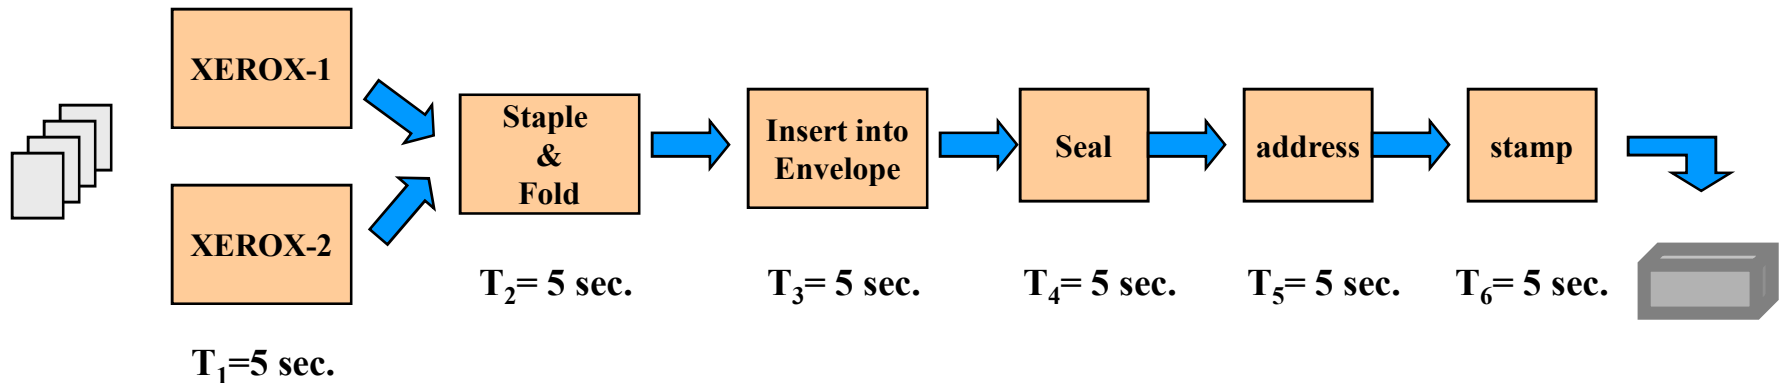
If you can not do much about the completion time for one task (i.e.  $\sum_i T_i$ ); what can you do to **maximize** the **speedup**?

(i) Create as many stations (stages) as possible  
and

(ii) Try to balance the load at each station, i.e.  $T_1 = T_2 = \dots = T_k$

# PIPELINING: EXAMPLE (CONTD.)

One possible configuration to maximize speedup:



- At what intervals, do we see a new envelope prepared for mailing?

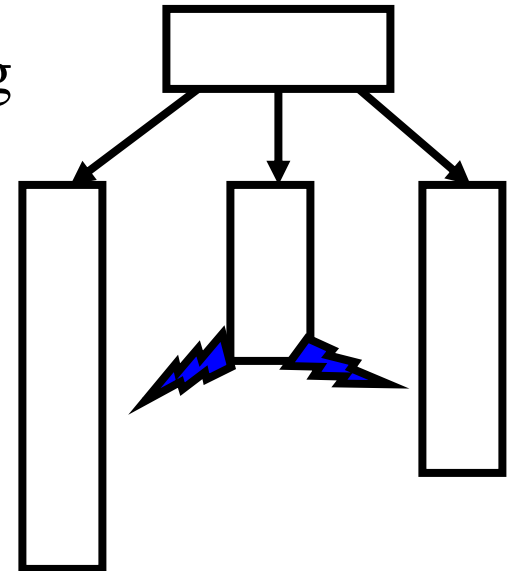
$$\text{Max}(T_1, T_2, \dots, T_k) = T_{\max} = 5 \text{ sec.}$$

- What is the speedup now?

$$\text{Speedup} = 30/5 = 6 = \text{number of stages in the pipeline !}$$

# Structure: Competition

- Evaluation of multiple solution strategies in parallel.
- It might be unknown which strategy is successful or which one is the fastest.
- With  $k$  processors,  $k$  strategies can be tested. If one of the additional strategies - not tested in the sequential program - is very fast, the speedup can be more than  $k$  (Superlinear speedup)
- Random search, speculative computing

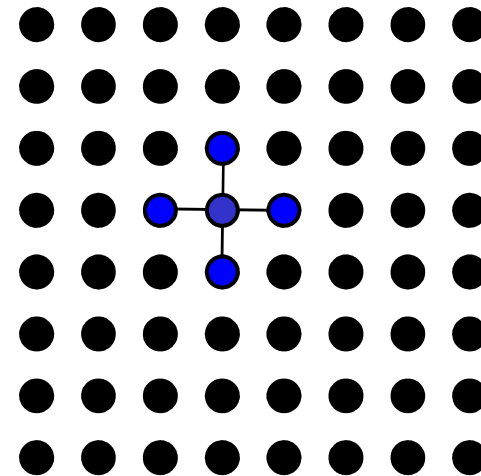


# Application Structure: Application

- Performance Analysis
  - o Pipeline
  - o Based on the parameters to make decision
- System Support
  - o MapReduce
- Optimization

# Example: Equation Solver Kernel

- Solver kernel for a simple partial differential equation.
- Finite difference method
- Grid  $(n+2) \times (n+2)$
- Fixed boundaries
- Interior points are recomputed
  - o Mean value of five points in stencil
  - o In place computation
  - o Gauss-Seidel method
    - New values of upper and right point
    - Old values of lower and left point
  - o Termination if difference between old and new value is below threshold for all points



$$A[i, j] = 0.2 * (A[i, j] + A[i-1, j] + A[i, j-1] + A[i, j+1] + A[i+1, j])$$

# Sequential Code

```
int n;                                /*size of matrix: (n + 2-by-n + 2)*/
float: **A, diff = 0;

main ()
begin
read(n);                             /*read input parameter: matrix size*/
A = malloc (a 2-d array of size n + 2 by n + 2 doubles);
initialize(A);                       /*initialize the matrix A somehow*/
Solve (A);                           /*call the routine to solve equation*/

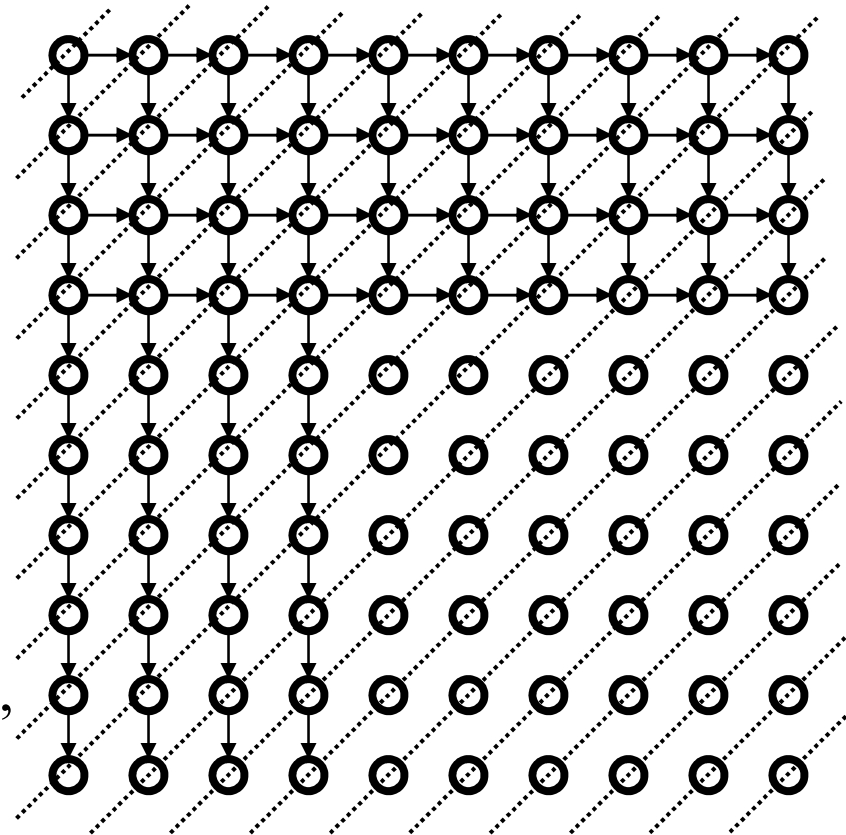
end main
```

# Routine SOLVE

```
procedure Solve (A)          /*solve the equation system'*/
float **A;                   /*A is an (n + 2)-by-(n + 2) array*/
begin
int i, j, done = 0;
float diff = 0, temp;
while (!done) do             /*outermost loop over sweeps*/
    diff = 0;                 /*initialize maximum difference to 0*/
    for i=1 to n do           /*sweep over nonborder points of grid*/
        for j=1 to n do
            temp = A[i,j];     /*save old value of element*/
            A[i,j] = 0.2 * (A[i,j] + A[i,j-1] + A[i-1,j] +
                           A[i,j+1] + A[i+1,j]); /*compute average*/
            diff += abs(A[i,j] - temp);
        end for
    end for
    if (diff/(n*n) < TOL) then done = 1;
end while
end procedure
```

# Dependences in Gauss-Seidel

- Dependences prohibit row or column wise parallelization
- Point-wise synchronization
- Parallel execution along anti-diagonals
  - o Proportional to  $n$
  - o Frequent synchronization, once per anti-diagonal
  - o Load imbalance for short anti-diagonals





# Relaxing Ordering Constraints

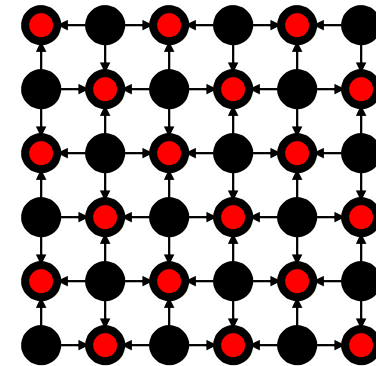
## ➤ Jacobi iterations

- o Full sweep with old values
- o Much slower convergence → more iterations
- o  $N^2$  parallel tasks in each iteration

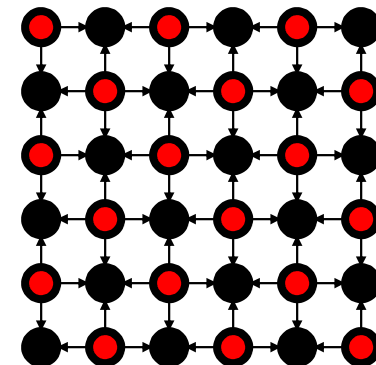
## ➤ Red-Black iterations:

- o Checkerboard like coloring scheme
- o Two phases
  - Computation of red points with old black values
  - Computation of black points with new red values
- o Faster convergence than Jacobi but more iterations than Gauss-Seidel
- o Each phase with  $n^2/2$  parallel tasks

Red phase



Black phase



# Message Passing Programming (1/7)

- Processes have private address spaces
- Values are communicated via send/receive operations
- Writing local code alike to thread programming

# Message Passing Programming (2/7)

```
Int pid, n, nprocs;          /*process id and number of processes*/
Float **myA;

Main()
Begin
    nprocs=get_proc_num();    /*number of started processes*/
    pid=get_my_rank();        /*pid of the executing process*/
    if (pid==0) read(n);      /*master only*/
    broadcast(0,n,sizeof(int),SIZE); /*communicate to all others*/
    solve();                  /*start computation*/
End
```

# Message Passing Programming (3/7)

```
procedure solve()
begin
  int I,j, pid, n1=n/nprocs, done=0;
  float temp, tempdiff, mydiff=0;
  myA=malloc(n/nprocs+2 by n+2);          /*allocate my rows+ghost rows*/
  initialize (myA);                       /*intialize my rows*/

  while (!done) do
    mydiff=0;
    if (pid!=0) then
      send(&myA[1,0],n*sizeof(float),pid-1,ROW)
    if (pid!=nprocs-1) then
      send(&myA[n1,0],n*sizeof(float),pid+1,ROW)
    if (pid!=0) then
      receive(&myA[0,0],n*sizeof(float),pid-1,ROW)
    if (pid!=nprocs-1) then
      receive(myA[n1+1,0],n*sizeof(float),pid+1,ROW)
```

# Message Passing Programming (4/7)

```
for i=1 to n1 do          /*sweep over nonborder points of grid*/
  for j=1 to n do
    temp = myA[i,j];      /*save old value of element*/
    myA[i,j] = 0.2 * (myA[i,j] + myA[i,j-1] + myA[i-1,j] +
                     myA[i,j+1] + myA[i+1,j]); /*compute average*/
    mydiff += abs(myA[i,j] - temp);
  end for
end for
```

# Message Passing Programming (5/7)

```
if (pid !=0) then                                /*send local diffs to P0*/
    send (mydiff,sizeof(float),0,DIFF)
    receive(done,sizeof(int),0,DONE)             /*receive done flag from P0*/
else
    for i=1 to nprocs-1 do
        receive(tempdiff,sizeof(float),*,DIFF) /*receive local diffs*/
        mydiff += tempdiff;                     /*accumulate local diffs*/
    endfor
    if (mydiff/(n*n)<TOL) then done=1 /*check condition*/
    for i=1 to nprocs-1 do
        send(done,sizeof(int),i,done);          /*send done flag to other procs*/
    end for
endif

endwhile
end procedure
```

# Message Passing Programming (6/7)

```
reduce(0,mydiff,sizeof(float),ADD)
if (pid ==0) then
    if (mydiff/(n*n)<TOL) then done=1;
    endif
endif
broadcast(0,done,sizeof(int));
endwhile
end procedure
```

# Message Passing Programming (7/7)

| Syntax                                   | Function                                                                                                                                                  |
|------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------|
| Send(src_addr, size, dest, tag)          | Send size bytes starting at src_addr to the dest process, with tag identifier.                                                                            |
| Receive(buffer_addr, size, src, tag)     | Receive a message with the tag identifier from the src process, and put size bytes of it into buffer starting at buffer_addr.                             |
| Reduce(root_pid, buffer, length, oper)   | Compute global value from local values in buffer of the given length of all processes with operation oper. The global value is delivered to root process. |
| Broadcast(root_pid, buffer, length, tag) | The root process sends the value in buffer to the other processes with the tag identifier.                                                                |



# The Parallel Partition LU (PPT) Algorithm

---

*Step 1.* Allocate  $A_i, d^{(i)}$  and elements  $a_{im}, c_{(i+1)m-1}$  to the  $i$ th node, where  $0 \leq i \leq p-1$ .

*Step 2.* Use the  $LU$  decomposition method to solve

$$A_i[\tilde{x}^{(i)}, v^{(i)}, w^{(i)}] = [d^{(i)}, a_{im}e_0, c_{(i+1)m-1}e_{m-1}]$$

*Step 3.* Send  $\tilde{x}_0^{(i)}, \tilde{x}_{m-1}^{(i)}, v_0^{(i)}, v_{m-1}^{(i)}, w_0^{(i)}, w_{m-1}^{(i)}$  from the  $i$ th node to the other nodes  $0 \leq i \leq p-1$ .

*Step 4.* Use the  $LU$  method to solve  $Zy = h$  on all nodes

*Step 5.* Compute in parallel on  $p$  processors

$$\Delta x^{(i)} = [v^{(i)}, w^{(i)}] \begin{bmatrix} y_{2i-1} \\ y_{2i} \end{bmatrix}$$

$$x^{(i)} = \tilde{x}^{(i)} - \Delta x^{(i)}$$

# The Solving process

1. Solve the subsystems in parallel
2. Solve the reduced system
3. Modification

$$\tilde{\mathbf{A}}^{-1} \mathbf{A} \mathbf{x} = \tilde{\mathbf{A}}^{-1} \mathbf{d}$$

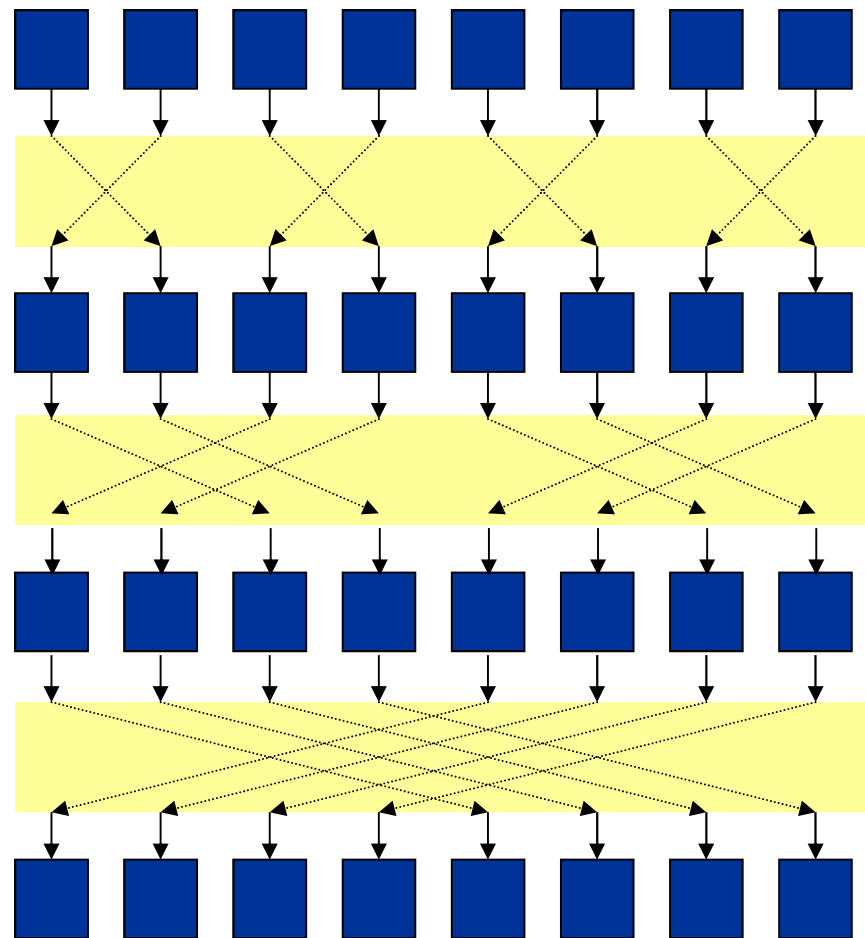
$$\begin{bmatrix} A_0^{-1} & & & \\ & A_1^{-1} & & \\ & & \ddots & \\ & & & A_{p-1}^{-1} \end{bmatrix} \begin{bmatrix} A_0 & * & & \\ * & A_1 & & \\ & * & \ddots & \\ & & * & A_{p-1} \end{bmatrix} \begin{pmatrix} x_0 \\ x_1 \\ \vdots \\ x_{n-1} \end{pmatrix} = \begin{bmatrix} A_0^{-1} & & & \\ & A_1^{-1} & & \\ & & \ddots & \\ & & & A_{p-1}^{-1} \end{bmatrix} \begin{pmatrix} d_0 \\ d_1 \\ \vdots \\ d_{n-1} \end{pmatrix}$$

## The Reduced System $(Zy=h)$

$$\begin{bmatrix} \mathbf{I} & & & \\ & \text{red} & & \\ & \text{yellow} & \mathbf{I} & \\ & \text{red} & & \text{red} \\ & & & \text{red} & \mathbf{I} \end{bmatrix} \begin{pmatrix} \cdot \\ \cdot \\ \cdot \\ \cdot \\ \cdot \end{pmatrix} = \begin{pmatrix} \cdot \\ \cdot \\ \cdot \\ \cdot \\ \cdot \end{pmatrix}$$

**Needs global communication**

# All-to-All Total Data Exchange



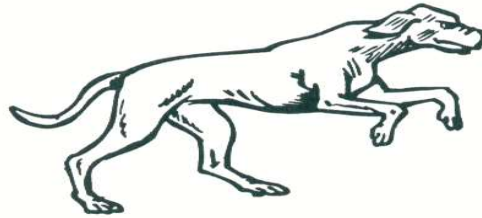
# Summary Parallel Programming

- Steps in the parallelization process
  - o Decomposition, assignment, orchestration, mapping
- Variety of programming models for SM and DM systems
  - o Different parallel languages and APIs
  - o Trend towards standards (OpenMP, MPI, Posix)

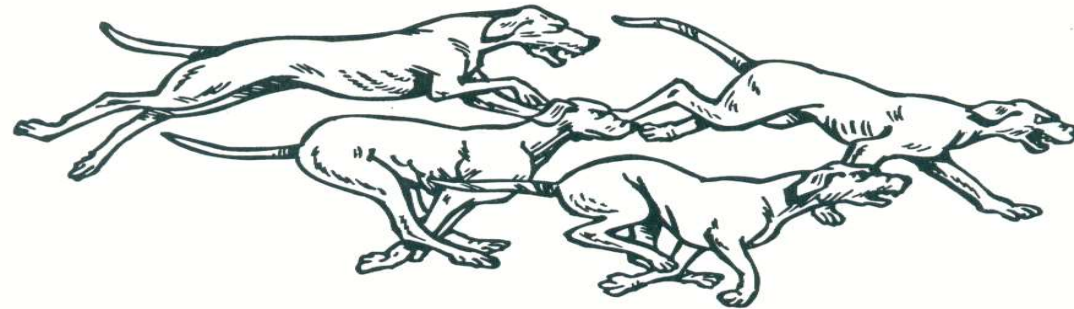
# Summary Parallel Programming Models

- Global vs local programming
  - o Global: directive-based, data parallel
  - o Local: thread-based, remote memory access, message passing
- Data vs functional parallelism
  - o Data parallelism only: data parallel
  - o Both: all the others
- Parallelism
  - o High: remote memory access, message passing, thread-based
  - o Low: directive-based, high level data parallel

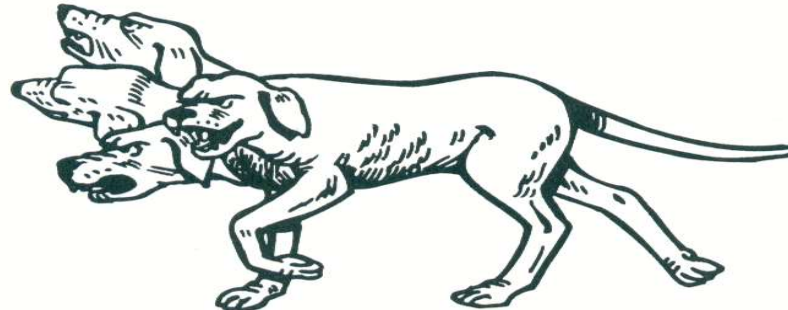
**Figure 4-1 A Dog**



**Figure 4-2 A Pack of Dogs**



**Figure 4-3 A Savage Multiheaded Pooch**





# Concurrent-AMAT: step to optimization

- The traditional AMAT(Average Memory Access Time) :

$$AMAT = HitCycle + MR \times AMP$$

- MR is the miss rate of cache accesses; and AMP is the average miss penalty

- **Concurrent-AMAT (C-AMAT):**

$$C-AMAT = HitCycle/C_H + pMR \times pAMP/C_M = 1/APC$$

- $C_H$  is the hit concurrency;  $C_M$  is the *pure* miss concurrency
- $pMR$  and  $pAMP$  are *pure* miss rate and average *pure* miss penalty
- A pure miss is a miss containing at least one cycle which does not have any hit activity

X.-H. Sun and D. Wang, "Concurrent Average Memory Access Time", in **IEEE Computers**, vol. 47, no. 5, pp. 74-80, May 2014. (IIT Technical Report, IIT/CS-SCS-2012-05)





# Recursive in Memory Hierarchy

- AMAT is recursive
  - $AMAT = HitCycle_1 + MR_1 \times AMP_1$   
Where  $AMP_1 = (HitCycle_2 + MR_2 \times AMP_2)$

- C-AMAT is also recursive

$$C-AMAT_1 = \frac{H_1}{C_{H_1}} + MR_1 \times \kappa_1 \times C-AMAT_2$$

Where

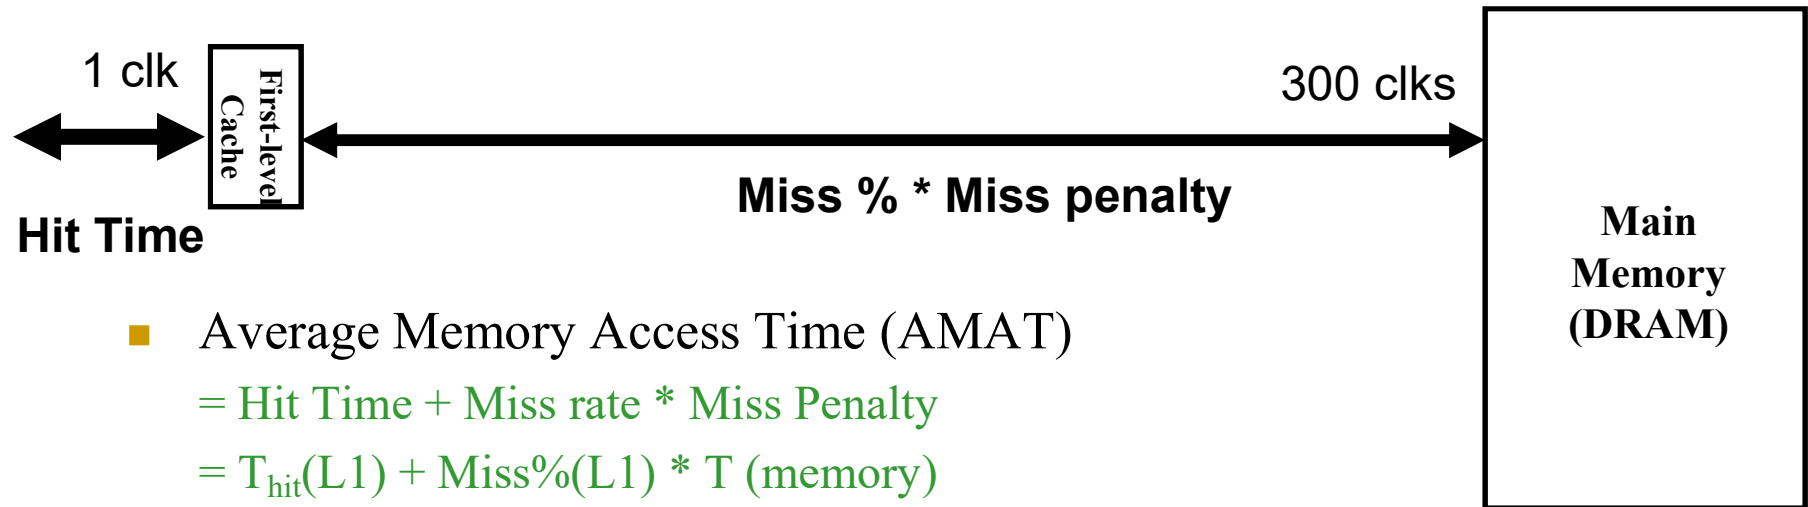
$$C-AMAT_2 = \frac{H_2}{C_{H_2}} + pMR_2 \times \frac{pAMP_2}{C_{M_2}}$$

$$\kappa_1 = \frac{pMR_1}{MR_1} \times \frac{pAMP_1}{AMP_1} \times \frac{C_{m_1}}{C_{M_1}}$$

**With Clear Physical Meaning**



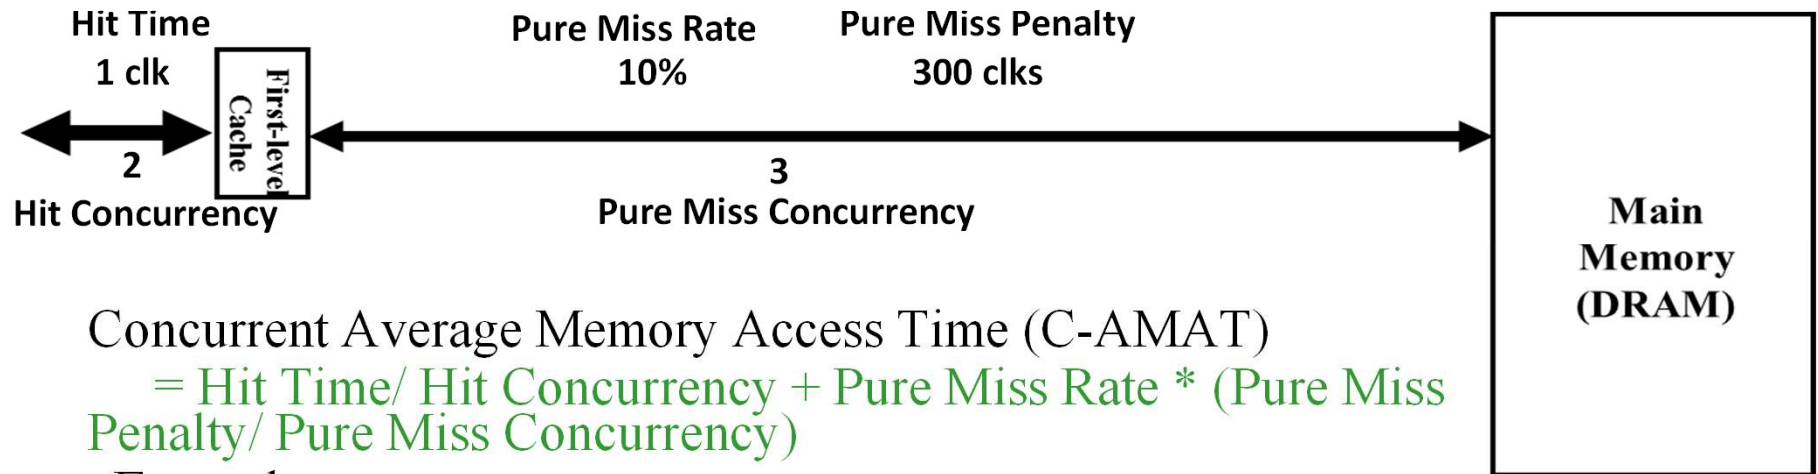
# Memory Hierarchy Performance



- Average Memory Access Time (AMAT)  
= Hit Time + Miss rate \* Miss Penalty  
=  $T_{hit}(L1) + \text{Miss\%}(L1) * T(\text{memory})$
- Example:
  - ❑ Cache Hit = 1 cycle
  - ❑ Miss rate = 10%
  - ❑ Miss penalty = 300 cycles
  - ❑  $AMAT = 1 + 300 * 10\% = 31$  cycles
- To further improve it?



# Example of calculating C-AMAT



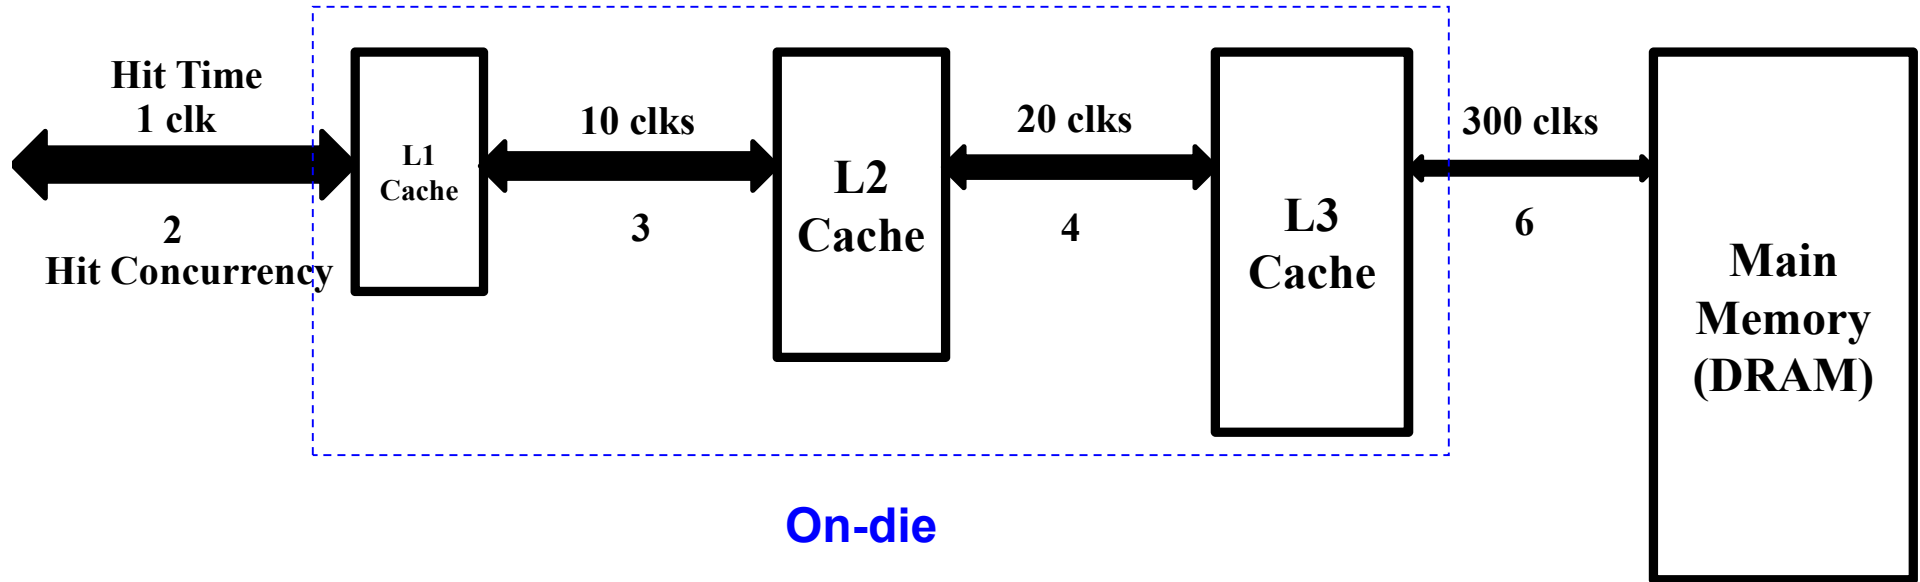
Concurrent Average Memory Access Time (C-AMAT)  
= Hit Time/ Hit Concurrency + Pure Miss Rate \* (Pure Miss Penalty/ Pure Miss Concurrency)

■ Example:

- Hit Time = 1 cycle
- Hit Concurrency = 2
- Pure Miss Rate = 10%
- Pure Miss Penalty = 300 cycles
- Miss Concurrency = 3
- C-AMAT =  $1/2 + 10\% * (300/3) = 0.5 + 10 = 10.5$  cycles



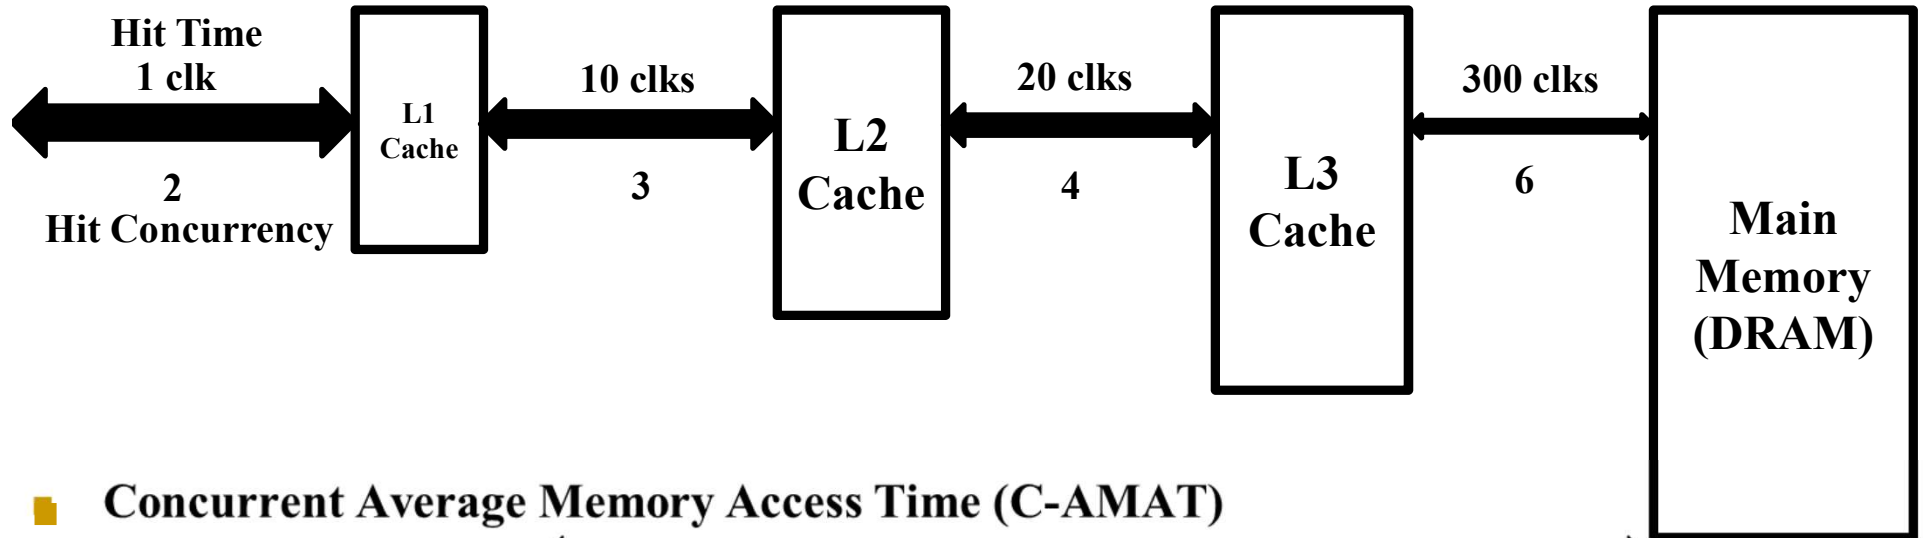
# Data Access Time: AMAT



- Average Memory Access Time (AMAT)  
$$= T_{\text{hit}}(\text{L1}) + \text{Miss}\%(\text{L1}) * (T_{\text{hit}}(\text{L2}) + \text{Miss}\%(\text{L2}) * (T_{\text{hit}}(\text{L3}) + \text{Miss}\%(\text{L3}) * T(\text{memory}) ) )$$
- Example: (Latency as shown above)
  - Miss rate: L1=10%, L2=5%, L3=1% (*Be careful miss rate definition*)
  - AMAT  
 $= 2.115$



# Data Access Time: C-AMAT



## ■ Concurrent Average Memory Access Time (C-AMAT)

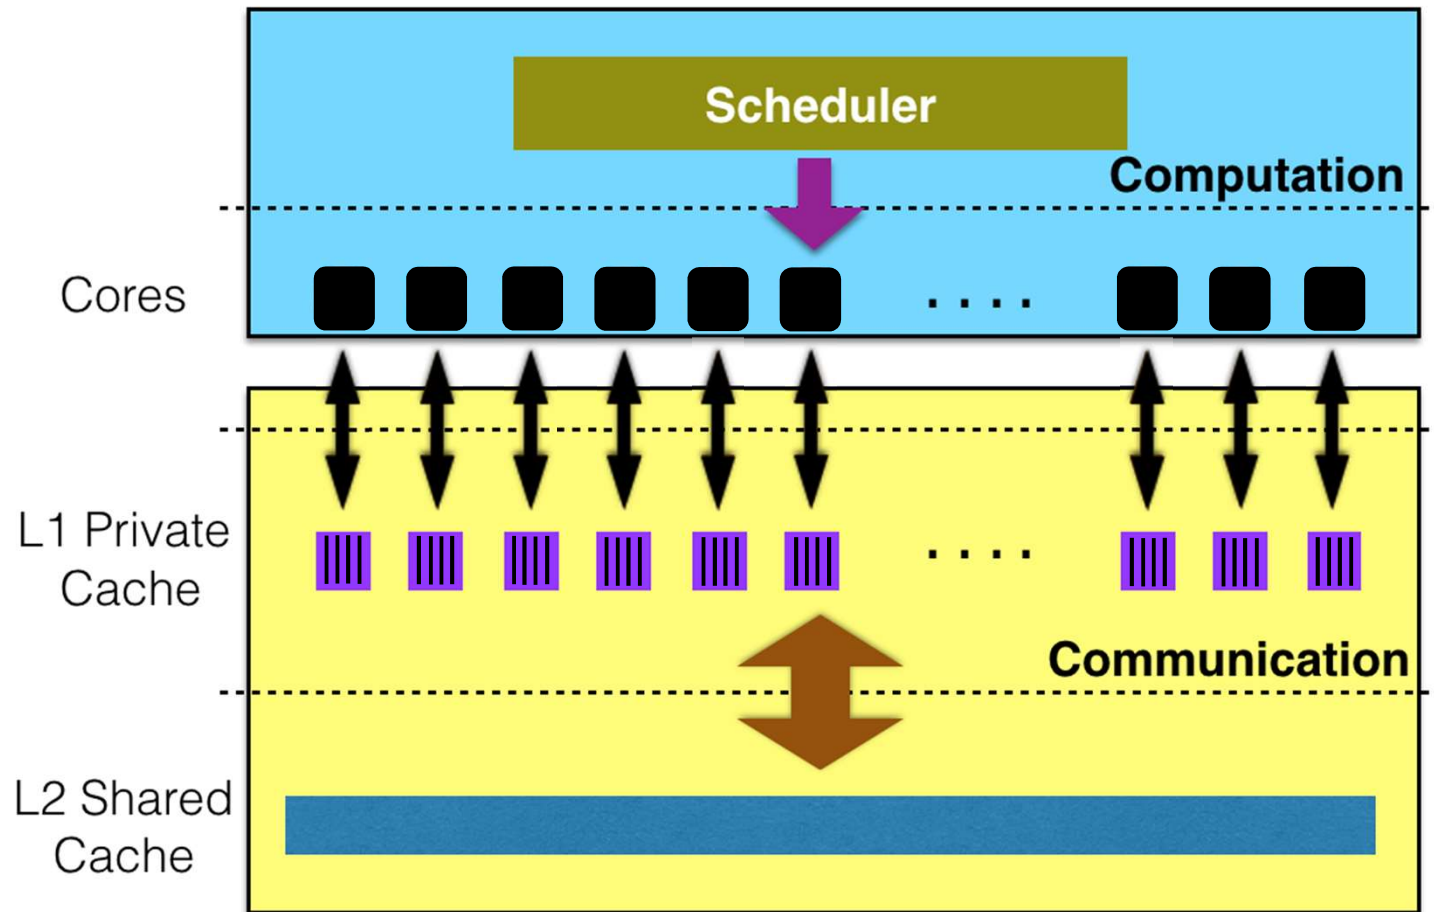
$$= \frac{H_1}{C_{H_1}} + MR_1 \times \kappa_1 \times \left( \frac{H_2}{C_{H_2}} + MR_2 \times \kappa_2 \times \left( \frac{H_3}{C_{H_3}} + MR_3 \times \kappa_3 \times \frac{H_{\text{Mem}}}{C_{H_{\text{Mem}}}} \right) \right)$$

## ■ Example

- Miss Rate: L1=10%, L2=5%, L3=1%      pMR, pAMP, AMP,  $C_M$ ,  $C_m$ : L1=7%, 10, 10, 5, 4
- $\kappa$ : L1=0.56, L2=0.6, L3=0.8      L2=3%, 60, 40, 9, 6
- C-AMAT $\approx$ 0.696      L3=0.8%, 400, 300, 16, 12

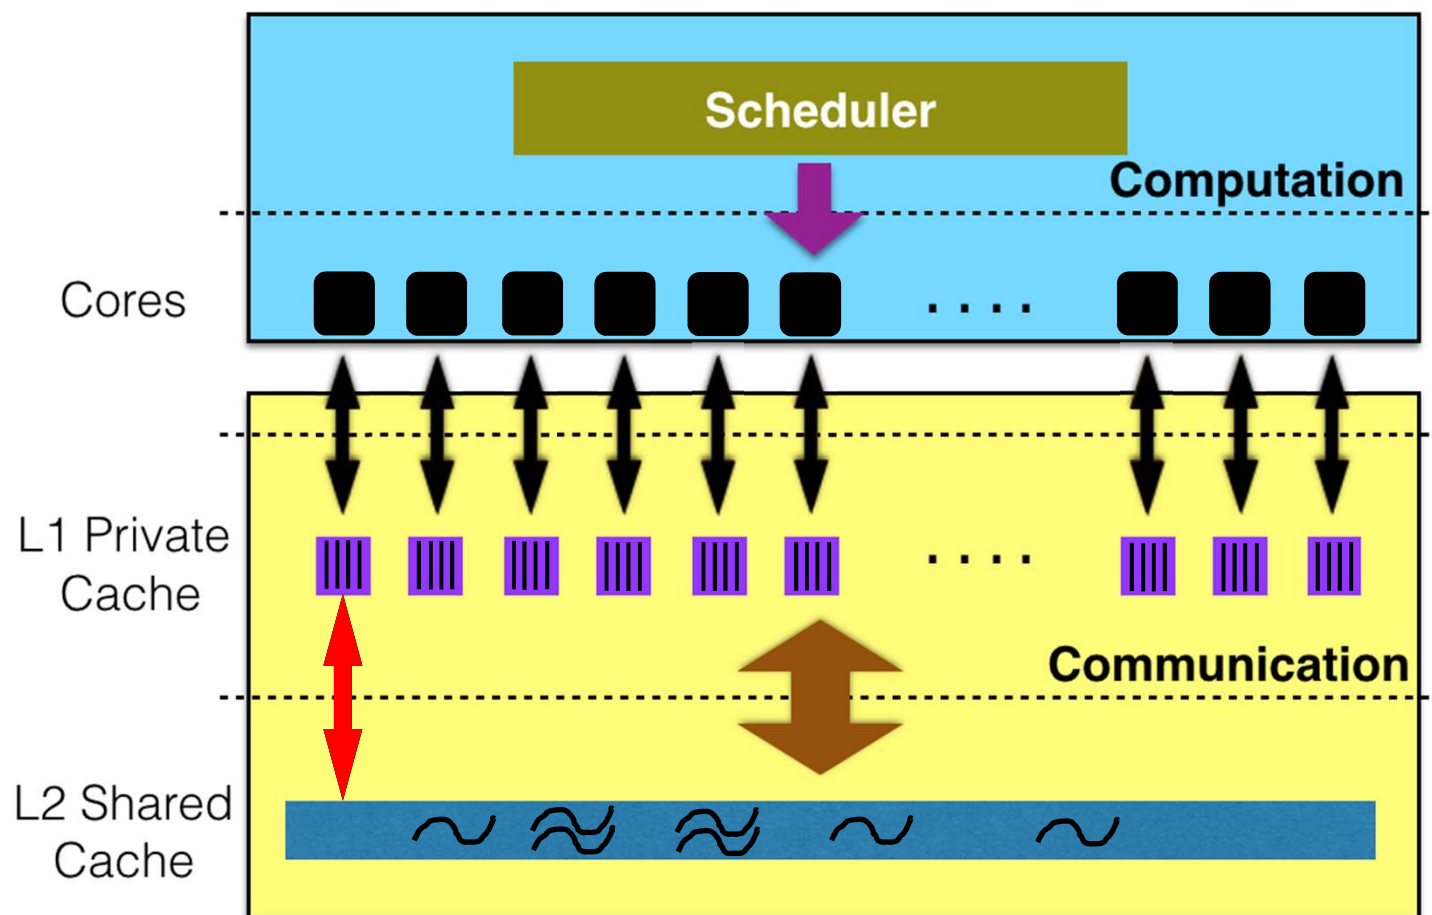


# C-AMAT in Multi-Core Environments





# C-AMAT in Multi-Core Environments





# What Does C-AMAT Say?

- C-AMAT is an extension of AMAT to consider concurrency
  - The same as AMAT, if no concurrency present
- C-AMAT introduces the **Pure Miss** concept:
  - Only pure miss causes performance penalty
- **High locality may hurt performance**
  - High locality may lead to pure miss
- **Two contributions of concurrency**
  - Increase bandwidth
  - Latency hiding (overlapping)





# What Does C-AMAT Say?

- C-AMAT also contains the overlapping factor

$$C-AMAT_1 = \frac{H_1}{C_{H_1}} + MR_1 \times \kappa_1 \times C-AMAT_2$$

- **Balance** locality, concurrency and Overlapping with C-AMAT
- A good explanation for why  
Optimal  $\neq$  Optimal Locality + Optimal Concurrency
- C-AMAT uniquely integrates the **joint impact** of locality concurrency, and overlapping for optimization
- Overlapping in connection locality and concurrency is a new issue of research



# Impact of C-AMAT

- New dimensions for optimization: **concurrency, overlapping and balancing**

$$C-AMAT_1 = \frac{H_1}{C_{H_1}} + MR_1 \times \kappa_1 \times C-AMAT_2$$

- Can apply at **each layer** of a memory hierarchy
- Existing mechanisms are readily to be extended
  - Every AMAT based optimization has a corresponding C-AMAT extension to include concurrency
- Concurrency as bandwidth increaser and **penalty reducer (overlapping)**: Accurate measure the concurrency contribution

$$\kappa_1 = \frac{pMR_1}{MR_1} \times \frac{pAMP_1}{AMP_1} \times \frac{C_{m_1}}{C_{M_1}}$$



# Misunderstanding of Memory Performance

*Optimal* = ~~Optimal~~ Locality

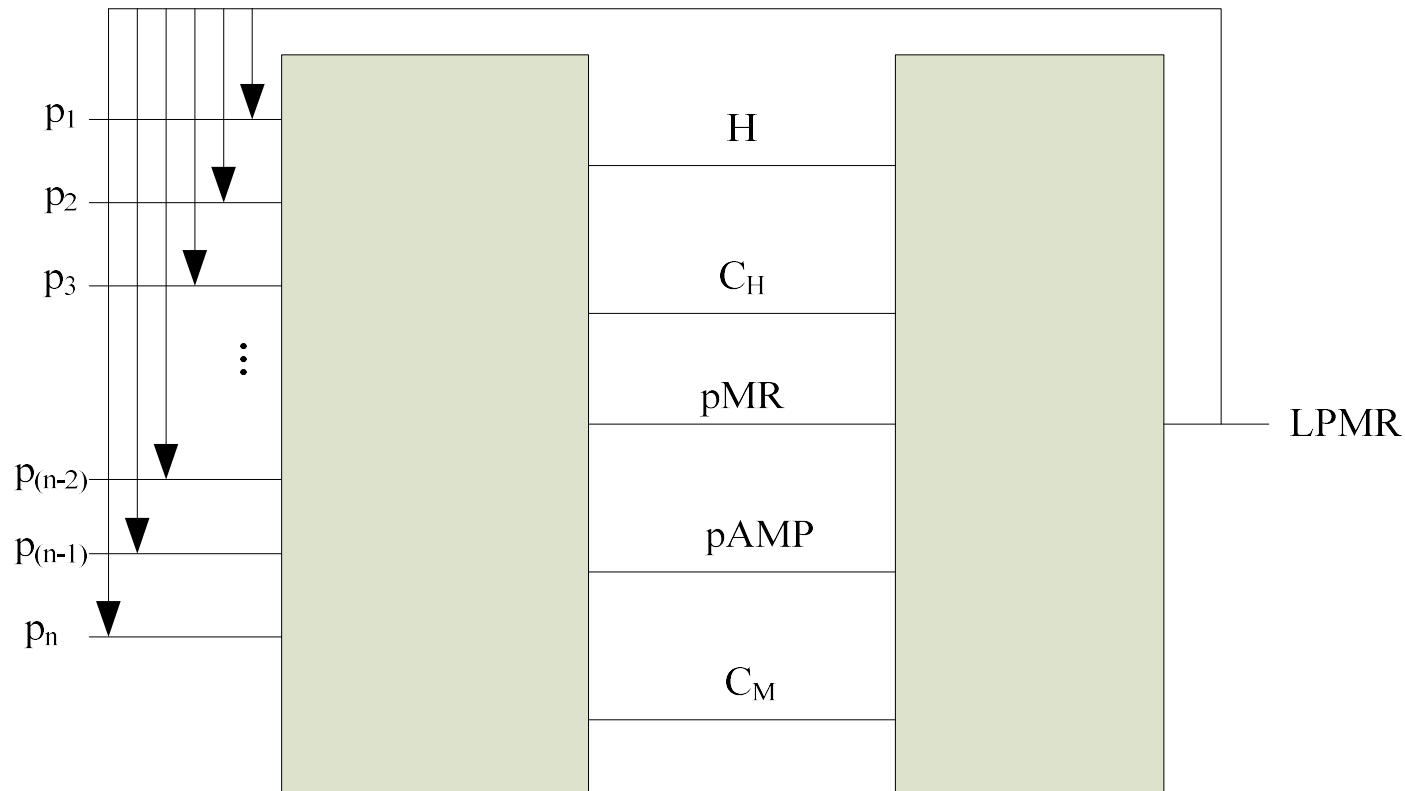


*Optimal* = ~~Optimal~~ Locality + Optimal Concurrence





# Practical: Parameters can be measured in hardware



Feedback-based optimization on scheduling and on reconfigurable architecture



# Challenge in Practice

---

- C-AMAT is general and powerful, but its measurement is environment (hardware) dependent
- Need a good understanding to conduct C-AMAT analysis in Multicore, GPU, GPGPU, FPGA ASIC, etc. environments
- GPU is a special case of C-AMAT matching



# Memory stall time: *the performance we care*

## Traditional AMAT model

$$CPU-time = IC \times (CPI_{exe} + \underbrace{f_{mem} \times AMAT}_{\text{Memory stall time}}) \times Cycle-time$$

## New C-AMAT model

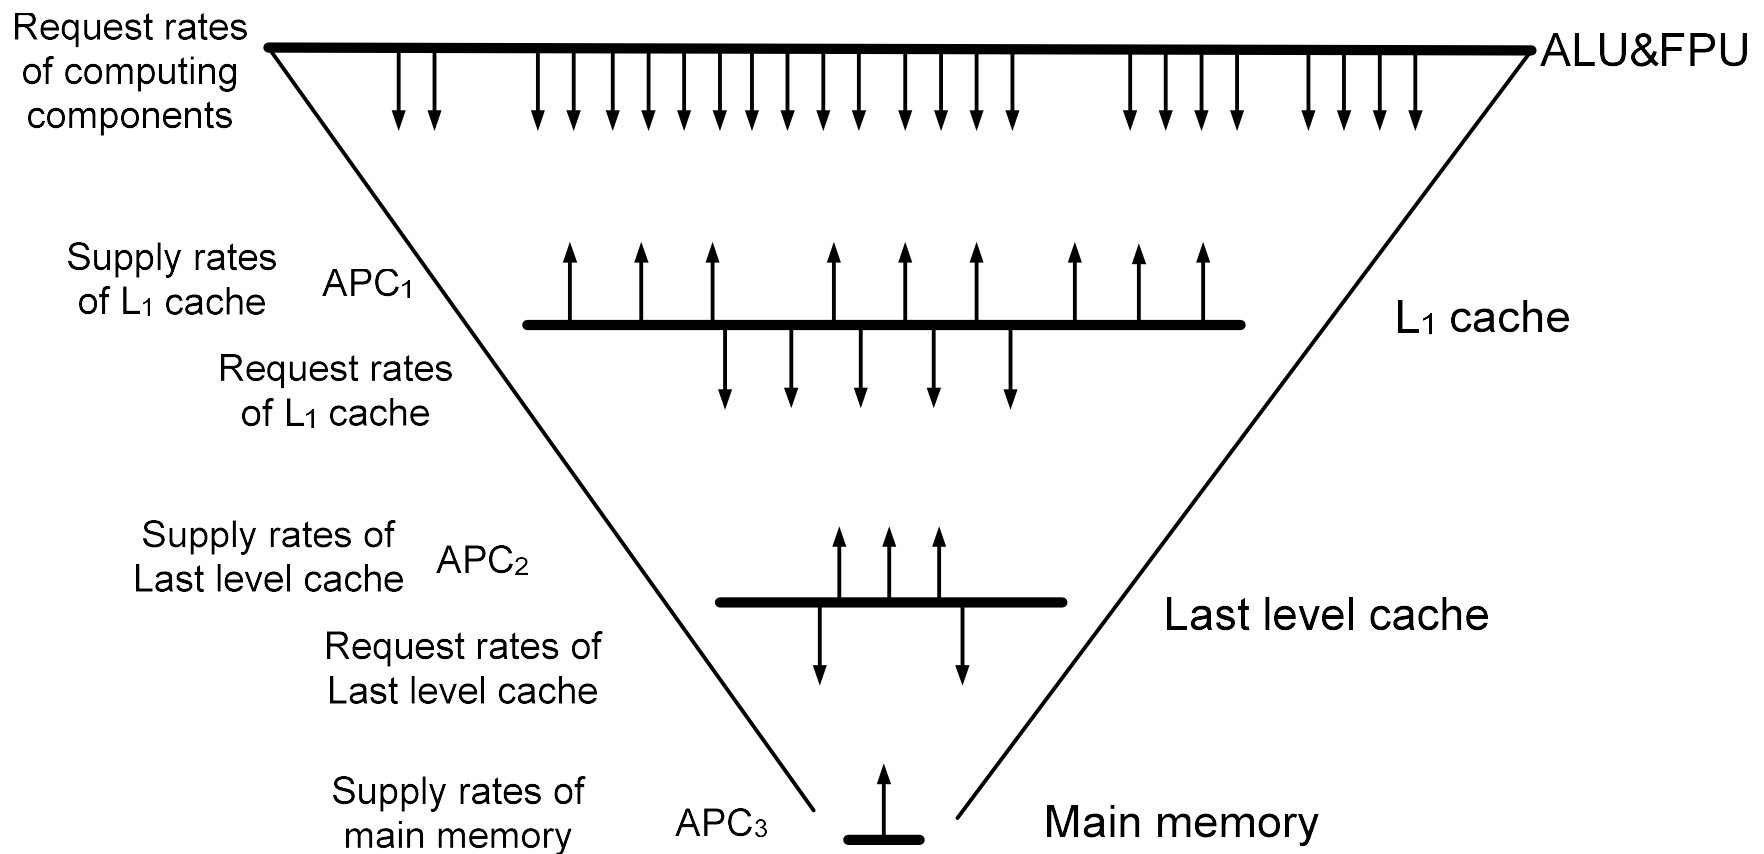
$$CPU-time = IC \times (CPI_{exe} + \underbrace{f_{mem} \times C-AMAT \times (1 - overlapRatio_{c-m})}_{\text{Memory stall time}}) \times cycle-time$$

$$CPU-time = IC \times (CPI_{exe} + \underbrace{f_{mem} \times \frac{pMR \times pAMP}{C_M}}_{\text{Memory stall time}}) \times Cycle-time$$

Only pure miss will cause processor stall, and the penalty is formulated here



# Application: Layered Performance Matching



Yu-Hang Liu, Xian-He Sun, "LPM: Concurrency-driven Layered Performance Matching," in ICPP2015, Beijing, China, Sept. 2015.



# Idea: Match the Request with Supply

$$LPMR(ALU \& FPU, L_1) = \frac{\text{Request rate from ALU \& FPU}}{\text{Supply rate by } L_1 \text{ cache}}$$

$$LPMR(L_1, LLC) = \frac{\text{Request rate from } L_1 \text{ cache}}{\text{Supply rate by LLC}}$$

$$LPMR(LLC, MM) = \frac{\text{Request rate from LLC}}{\text{Supply rate by main memory}}$$

- Match at **each memory layer**
- Adjust the supply performance with concurrency





# Quantify Mismatching: with C-AMAT

$$LPMR_1 = \frac{IPC_{exe} \times f_{mem}}{APC}$$

$$LPMR_2 = \frac{IPC_{exe} \times f_{mem} \times MR_1}{APC_2}$$

$$LPMR_3 = \frac{IPC_{exe} \times f_{mem} \times MR_1 \times MR_2}{APC_3}$$

- C-AMAT measures the request and supply at each layer
- C-AMAT can increase supply with effective concurrency
- Mismatch ratio directly determines memory stall time



# C-AMAT in Action



## New C-AMAT model

$$CPU-time = IC \times (CPI_{exe} + \underbrace{f_{mem} \times \frac{pMR \times pAMP}{C_M}}_{\text{Memory stall time}}) \times Cycle-time$$

Only pure miss will cause processor stall, and the penalty is formulated here

## The Relation of LPMR and Stall time

$$CPU-time = IC \times CPI_{exe} \times \underbrace{(1 + \kappa_1 \times LPMR_2)}_{\text{Memory stall time}} \times Cycle-time$$

Y. Liu and X.-H. Sun, "Reevaluating Data Stall Time with the Consideration of Data Access Concurrency," Journal of Computer Science and Technology (JCST), March, 2015



# The LPM Model

**We get the relation between data stall time and LPMR1**

$$Data - stall - time = CPI_{exe} \times (1 - overlapRatio_{c-m}) \times LP.$$

**We get the relation between data stall time and LPMR2**

$$Data - stall - time = CPI_{exe} \times \kappa_1 \times LPMR_2$$

**Because our final goal is to minimize data stall time, the two equations above. provide the baseline for LPM optimizations**



# The Threshold Requirement

$$\text{overlapRatio}_{c-m} = \frac{H_1/C_{H_1}}{(H_1/C_{H_1} + pAMP_1/C_{M_1})}$$

$$LPMR_1 \leq \frac{\Delta\%}{1 - \text{overlapRatio}_{c-m}}$$

$$LPMR_2 \leq \frac{\Delta\%}{\kappa_1}$$



# A Matching Example (increase performance)

$$LPMR = \frac{IPC_{exe} \times f_{mem}}{APC} \quad (1)$$

- **Assume:**

$$Cycle_{CPU} = 2 \text{ ns}$$

$$Cycle_{mem} = 8 \text{ ns}$$

$$f_{mem} = 20\%$$

$$IPC_{exe} = 2.5$$

$$APC = 1$$

- **Then:**

- In 8 ns, there is 1 memory cycle, and the data supply rate is:

$$APC * N\_Cycle_{mem} = 1$$

- In 8 ns, there is 4 cpu cycle, and the data request rate is:

$$IPC_{exe} * N\_Cycle_{CPU} * f_{mem} = 2$$

- **So:**

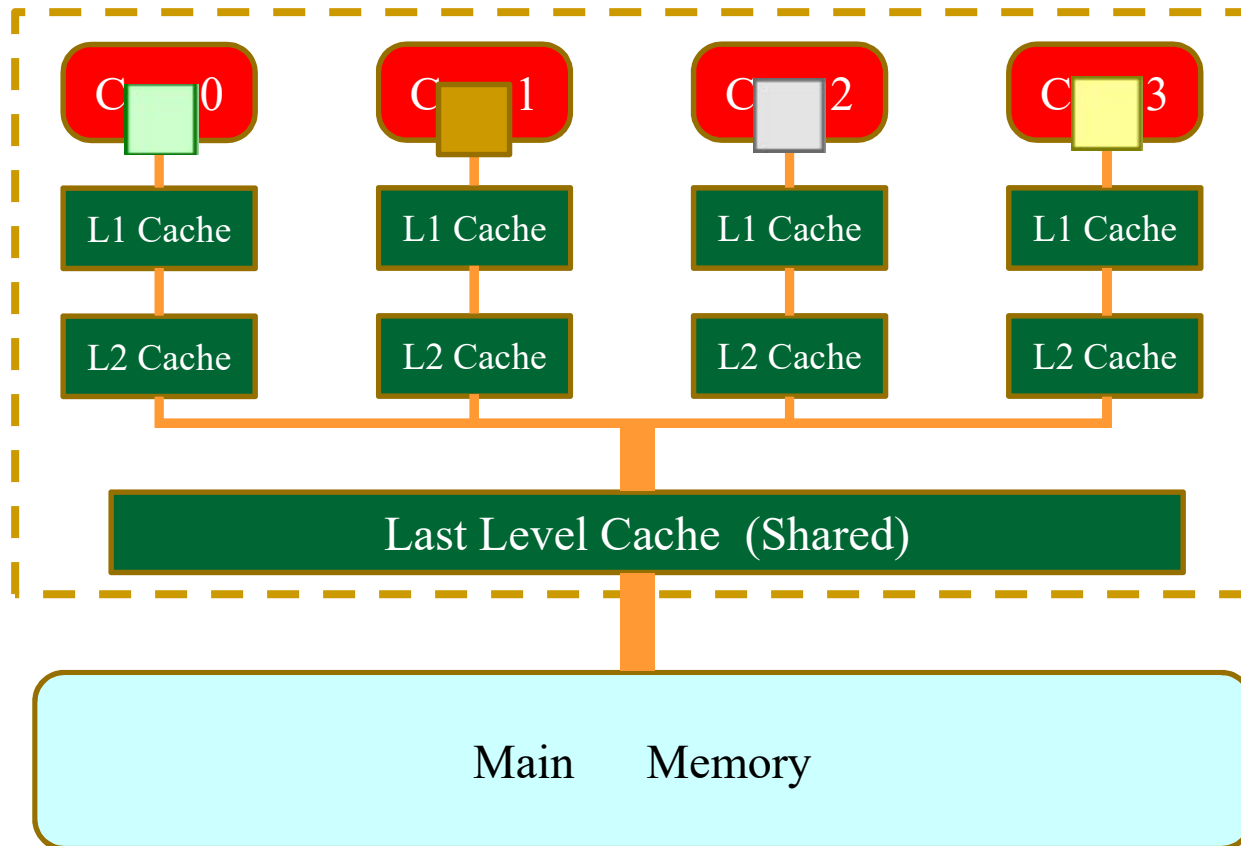
- The data supply rate does not match the data request rate. We can improve memory *concurrency* by adding memory banks or ports to increase data supply ability to adjust the data supply rate. For example, increase *APC* from 1 to 2, so:

$$APC * N\_Cycle_{mem} = 2$$

- And the data supply rate matches with the data request rate.



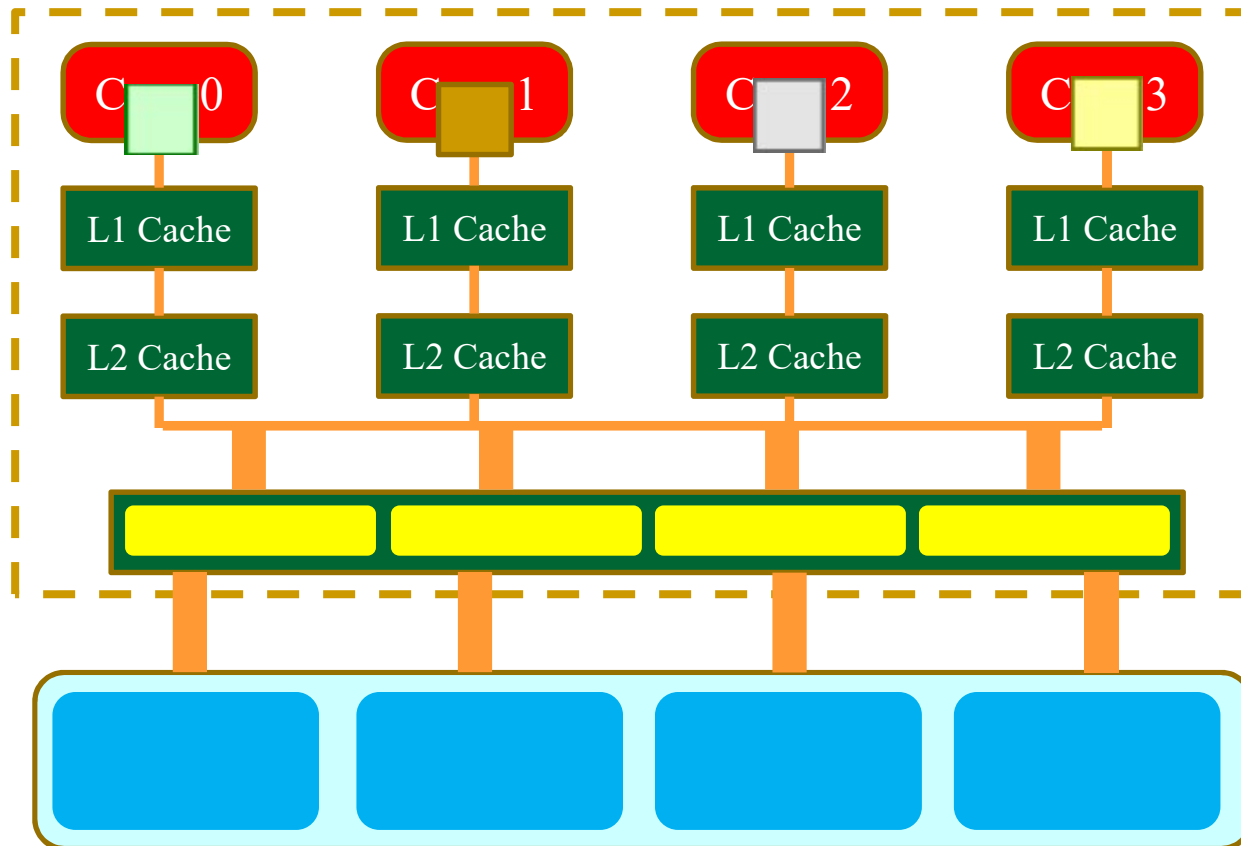
# Memory Access (Concurrency)



**Model and algorithm**



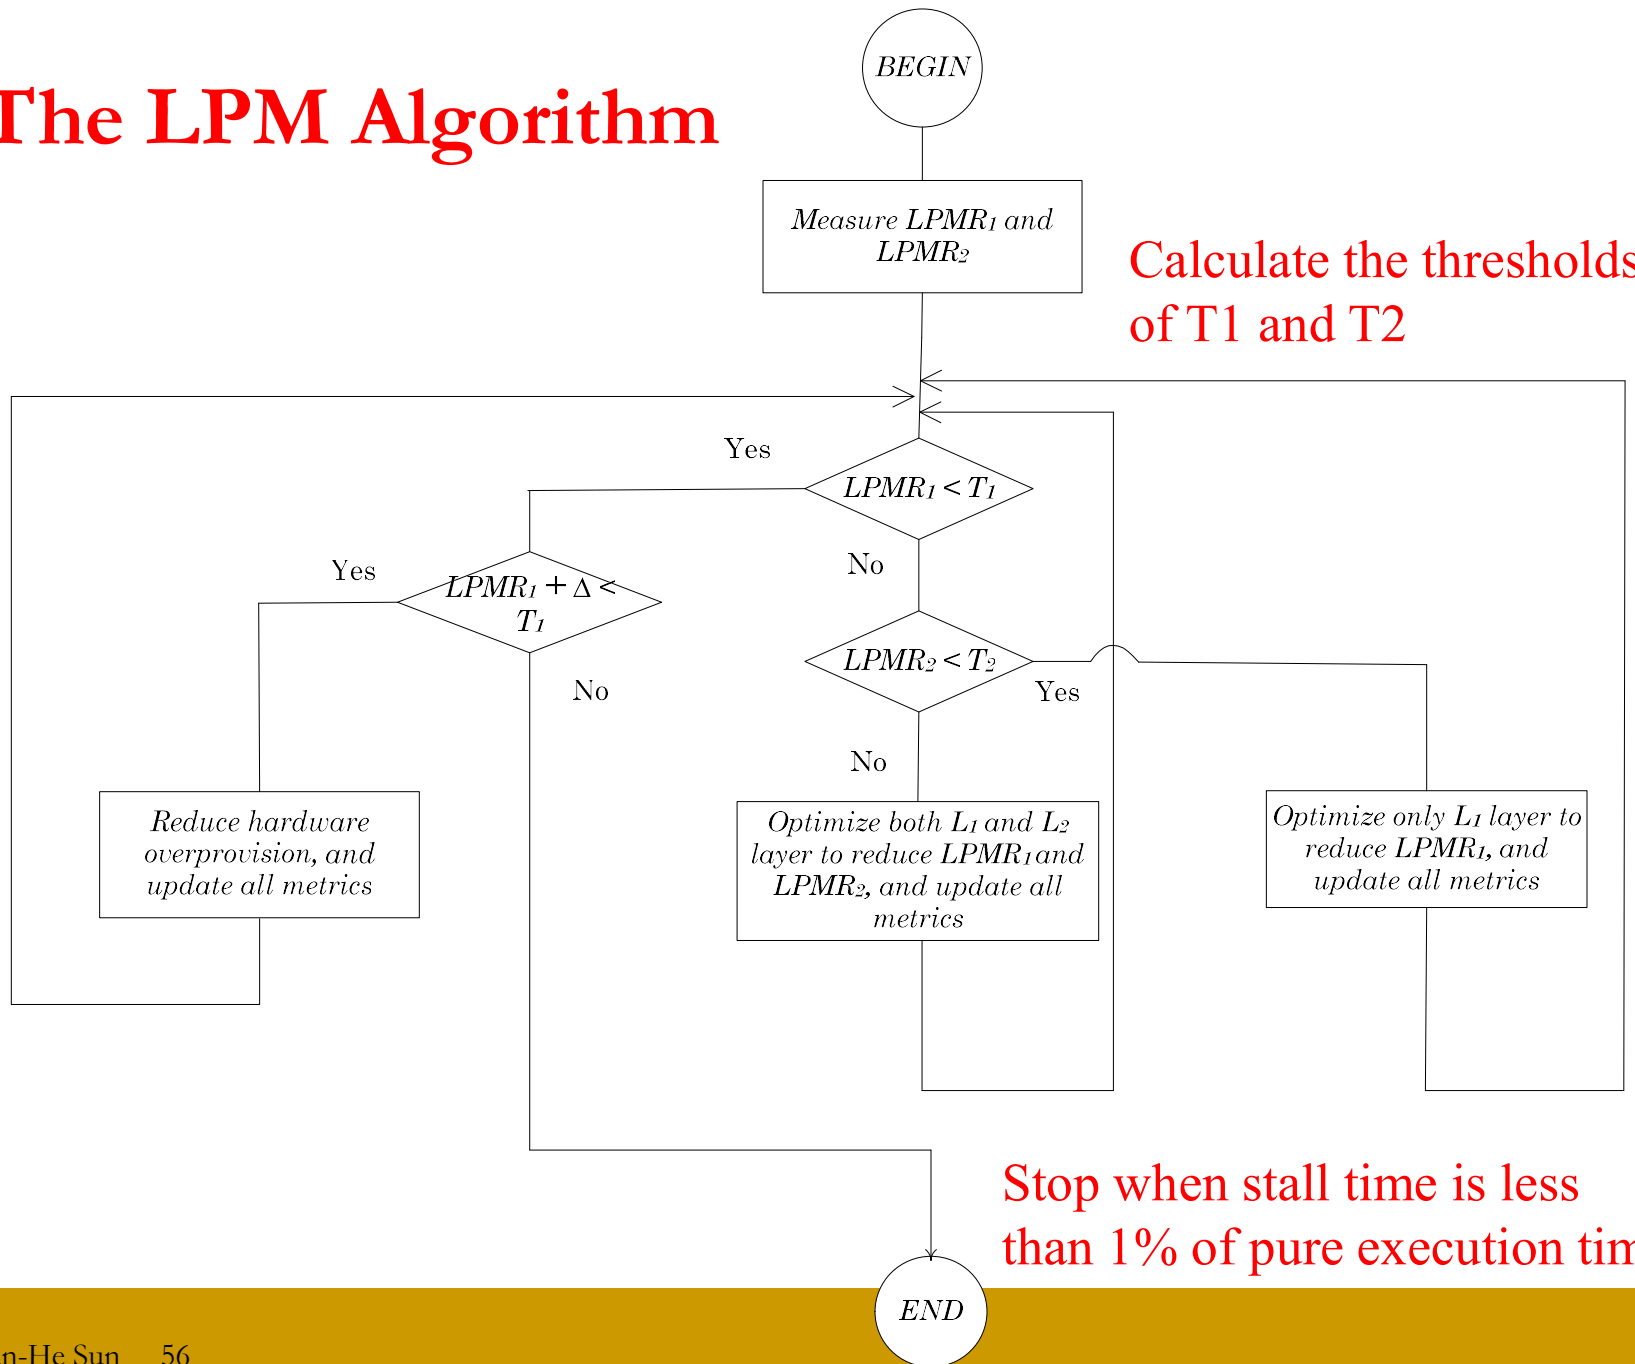
# Memory Access (Concurrency)



**Model and algorithm**



# The LPM Algorithm







## LPMR Reduction Algorithm

```
1: // Initially measure the metrics
2:   For each application or thread, measure the LPMRs in a memory hierarchy
3:   Get the threshold  $T_1$  and  $T_2$  according to Eq. (22) and (23)
4:   Begin Do
5:     // LPM optimization loop
6:     // Case I when both L1 and L2 layer need an optimization
7:     While ( $LPMR_1 > T_1$  and  $LPMR_2 > T_2$ ) Do
8:       Optimizing at L1 layer and L2 layer,
9:       Update all the metrics ( $LPMR_1$ ,  $LPMR_2$ ,  $T_1$  and  $T_2$ )
10:    Until ( $LPMR_1 \leq T_1$  or  $LPMR_2 \leq T_2$ )
11:    // Case II when only L1 layer needs an optimization
12:    While ( $LPMR_1 > T_1$  and  $LPMR_2 \leq T_2$ ) Do
13:      Optimizing at L1 layer
14:      Update all the metrics ( $LPMR_1$ ,  $LPMR_2$ ,  $T_1$  and  $T_2$ )
15:    Until ( $LPMR_1 \leq T_1$ )
16:    // Case III when no layer needs to optimize and overprovision may need to reduce
17:    //  $\delta$  is a positive value
18:    While ( $LPMR_1 + \delta < T_1$ ) Do
19:      Reduce hardware overprovision
20:      Update all the metrics ( $LPMR_1$ ,  $LPMR_2$ ,  $T_1$  and  $T_2$ )
21:    Until ( $LPMR_1 \geq T_1 - \delta$ )
22:    // Case IV when no layer needs to optimize and no overprovision needs to reduce
23:    If ( $T_1 \geq LPMR_1 \geq T_1 - \delta$ )
24:      End the algorithm
25:    Endif
26:  Until (End)
```

Pseudo code



## Comments and Questions:

- Pipelining can be combined with concurrency for best performance: subarray
- Can you derive the general form of LPMR<sub>i</sub>?