



ILLINOIS INSTITUTE OF TECHNOLOGY

File Systems Basics

Anthony Kougkas

akougkas@iit.edu

Outline

- File System Overview
- File System Architecture
- File System Semantics
- File System Operations

File System Overview

- System that permanently stores data
- Usually layered on top of a lower-level physical storage medium
- Divided into logical units called “files”
 - Addressable by a *filename* (“foo.txt”)
 - Usually supports hierarchical nesting (directories)
- A file *path* joins file & directory names into a **relative** or **absolute** address to identify a file (“/home/aaron/foo.txt”)

File (an abstraction)

- A (potentially) large amount of information or data that lives a (potentially) very long time
 - Often *much* larger than the memory of the computer
 - Often *much* longer than any computation
 - Sometimes longer than life of machine itself
- (Usually) organized as a linear array of bytes or blocks
 - Internal structure is imposed by application
 - (Occasionally) blocks may be variable length
- (Often) requiring concurrent access by multiple processes
 - Even by processes on different machines!

Directory – A Special Kind of File

- A tool for users & applications to organize and find files
 - User-friendly names
 - Names that are meaningful over long periods of time
- The data structure for OS to locate files (i.e., containers) on disk

File System Basics

- *File:*
 - Named collection of logically related data
 - Unix file: an uninterpreted sequence of bytes
- *File system:*
 - Provides a logical view of data and storage functions
 - User-friendly interface
 - Provides facility to create, modify, organize, and delete files
 - Provides sharing among users in a controlled manner
 - Provides protection

File Types and Attributes

- *File types:*
 - Regular files
 - Directories
 - Character special files: used for serial I/O
 - Block special files: used to model disks [buffered I/O]
- *File attributes:* varies from OS to OS
 - Name, type, location, size, protection info, password, owner, creator, time and date of creation, last modification, access
- *File operations:*
 - Create, delete, open, close, read, write, append, get/set attributes
- *File access:*
 - Sequential, random

Attributes of Files

- *Name:*
 - Although the name is not always what you think it is!
- *Type:*
 - May be encoded in the name (e.g., *.cpp*, *.txt*)
- *Dates:*
 - Creation, updated, last accessed, etc.
 - (Usually) associated with container
 - Better if associated with content
- *Size:*
 - Length in number of bytes; occasionally rounded up
- *Protection:*
 - Owner, group, etc.
 - Authority to read, update, extend, etc.
- *Locks:*
 - For managing concurrent access
- ...

Definition — *File Metadata*

- Information *about* a file
 - Maintained by the file system
 - Separate from file itself
 - Usually attached or connected to the file
 - E.g., in block # –1
 - Some information visible to user/application
 - Dates, permissions, type, name, etc.
 - Some information primarily for OS
 - Location on disk, locks, cached attributes

Observation – some attributes are not visible to user or program

- E.g., *location*
 - Location is stored in metadata
 - Location can change, even if file does not
 - Location is not visible to user or program

File System

- Components: directory, authorization, file service and system service
 - **Authorization service:** between file and directory services
 - **Directory service:** used to keep track of the location of all resources in the system
 - **File service** provides a transparent way of accessing any file in the system in the same way
 - **System service:** file system's interface to hardware


File Systems

- ***File service:***
 - Specification of what the file system offers to client
 - Actions
 - Client primitives
 - Parameters, application programming interface (API)
 - Does not include how service is implemented
- ***File server:***
 - Process that runs on a machine and implements file service
 - Can have several servers on one machine (UNIX, DOS,...)
 - ideally, clients do not know the distributed nature


Architectures

- How are FS generally organized?
 - Client-server architectures
 - Example: Sun Microsystem's NFS
 - File servers with a standardized view of its local file system; clients can access these files
 - Cluster-based distributed file systems
 - Example: file striping, partitioning the whole file system, GFS
 - Symmetric architectures
 - Fully symmetric organization based on p2p
 - Example: Ivy

Client-Server Architectures




- **Remote access model**
 - Work done at the server
- Stateful server (e.g., databases)
- Pros & cons?



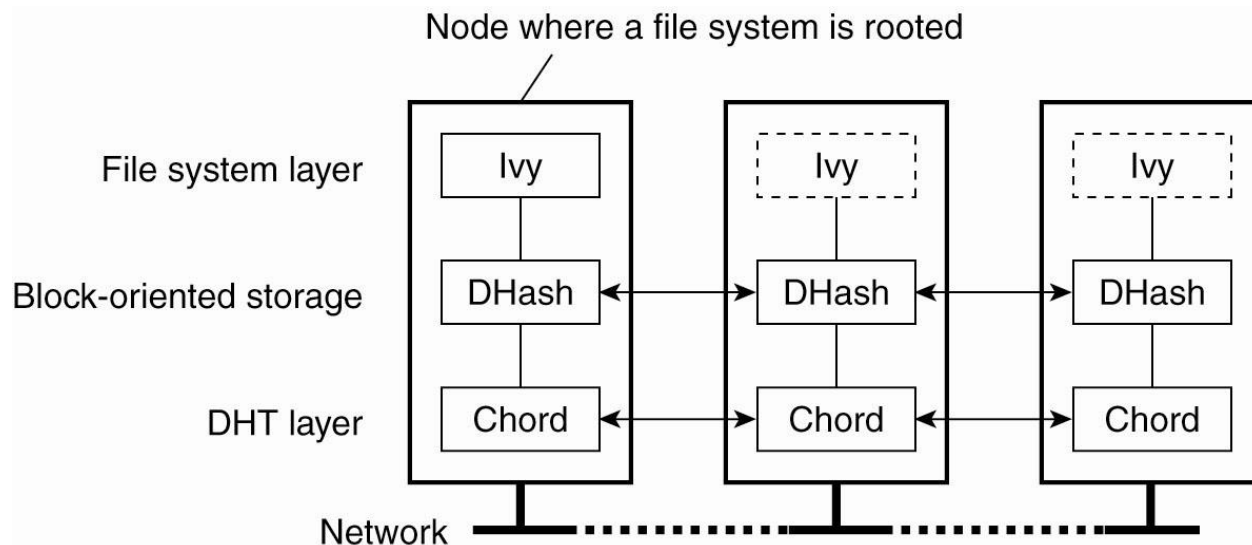
- **Upload/download model**
 - Work done at the client
- Stateless server
- Pros & cons?

Cluster-based DFSs



- Figure (b): When server clusters are used for parallel applications
 - File-striping techniques, a single file is distributed across multiple servers
- Figure (a): For general-purpose applications, when file striping may not be effective
 - Partition the file system as a whole and simply store different files on different servers

Symmetric Architectures



- Commonly used DHT-based systems for distributing data, combined with a key-based lookup mechanism
- Key difference: whether build a file system on top of a distributed storage layer
- Example: Ivy

System Structure

- How should the system be organized?
 - Are clients and server different?
 - Same process implements both functionality
 - Different processes, same machine
 - Different machines (a machine can either be client or server)
 - How are file and directory services organized-same server?
 - Different server processes: cleaner, more flexible, more overhead
 - Same server: just the opposite
 - Caching/no caching
 - server
 - client
 - How are updates handled?
 - File sharing semantics?
 - Server type: stateful vs. stateless

Semantics of File Sharing

- **Unix semantics:** used in centralized systems
 - Read after write returns value written
 - System enforces absolute time ordering on all operations
 - Always returns most recent value
 - Changes immediately visible to all processes
- Issues in distributed systems
 - Single file server (no client caching): easy to implement UNIX semantics
 - Client file caching: improves performance by decreasing demand at the server; updates to the cached file are not seen by other clients
 - Conclusion:?

Semantics of File Sharing

- **Session semantics** (relaxed semantics)
 - Local changes only visible to process that opened file
 - File close => changes made visible to all processes
 - Allows local caching of file at client
- **Problem:**
 - What if two or more clients are caching and modifying a file?

Semantics of File Sharing

- **No file update semantics** (Immutable files):
 - Files are never updated/modified
 - Allowed file operations: CREATE and READ
 - Files are atomically replaced in the directory
 - Problems:
 - what if two clients want to replace a file at the same time?
 - Delete file in use by another process

Semantics of File Sharing

- **Atomic transactions**
 - All file changes are delimited by a Begin and End transaction
 - All files requests within the transaction are carried out in order
 - The complete transaction is either carried out completely or not at all (atomicity)
 - Serializable access
 - Problem: ?

Operations on Files

- *Open, Close*
 - Gain or relinquish access to a file
 - OS returns a *file handle* – an internal data structure letting it cache internal information needed for efficient file access
- *Read, Write, Truncate*
 - *Read*: return a sequence of n bytes from file
 - *Write*: replace n bytes in file, and/or append to end
 - *Truncate*: throw away all but the first n bytes of file
- *Seek, Tell*
 - *Seek*: reposition *file pointer* for subsequent reads and writes
 - *Tell*: get current *file pointer*
- *Create, Delete*:
 - Conjure up a new file; or blow away an existing one

Methods for Accessing Files

- *Sequential* access
- *Random* access
- *Keyed* (or indexed) access