

Review Session

Exam Comments

- 3:00pm-5:00pm Room **SB111** on 11/14/19
 - Closed book, closed notes, no cell phone, no calculator
 - Allow to bring a 4"x6" "cheat card"
 - Anything on the card **MUST HAND WRITING**
- All lectures included on the test
 - Focus on things taught in class (lecture notes, in-class discussions, homework, programming assignment)
- Exam structure:
 - Short answer questions
 - 'big' questions
 - Programming design related questions
 - Performance analysis/evaluation questions
- Good luck!

Work Opportunities

- Research opportunities for graduate students:
 - Always look for self-motivated and hard-working grad students
 - Ph.D. students: CS597 and CS691
 - MS students: CS591 “Research and Thesis for MS Degree”
 - Take CS546 & CS550, check my research projects, send me your CV
- Research opportunities for undergrad students:
 - NSF REU (Research Experiences for Undergraduates) with Prof. Xian-He Sun
 - Various project topics, including development of scheduling simulator, analysis of system logs,
 - UG students: CS491
 - Pay per hour
 - If interested, contact Prof. Sun (sun@iit.edu)

Introduction

- Evolution of Computing:
 - Machine, Grid, Pervasive, Cloud, Edge, and Big Data
- Overview of parallel processing:
 - What is, driven force, aspect
 - Where the performance come from?
 - Multicore, and manycore architecture
 - Challenges of parallelism
 - Memory-wall, and memory hierarchy
 - Principle of locality
 - Shared and distributed memory machines
- Technology advance

Architecture

- Basic components of any architecture:
 - Processors and memory (processing units)
 - Interconnect: static & dynamic (performance factor)
- Logic classification based on:
 - Control mechanism (Flynn's Taxonomy)
 - SISD, SIMD, MISD, MIMD
 - Address space organization
 - Shared Address Space (NUMA)
 - Distributed Address Space
- Cluster
 - Interconnect
 - Parameters, performance

Software

- Important software
 - OS, Compiler, Libraries, Tools, Timing and Profiling, Schedulers, Debugger, Distributed file systems, Parallel virtual file systems
- Software design
 - Style of parallel processing
 - Approach to parallel programming
 - Thread
 - Pthread
 - Synchronization primitives
- Parallel File Systems
 - What is parallel I/O
 - Data Distribution

Data and Task Parallelization

- Data parallel: C program loops where each iteration of a loop is independent and represents a simple statement and is executed on a different processor

```
for (i=0;i<1000;i++)  
    a[i]=b[i]+c[i];
```

- Task parallel: multiple C program loops which cannot be parallelized individually, but the different code blocks are independent and are executed on different processor

```
for (i=0;i<1000;i++) /*block 1 */  
    b[i+1]=b[i]+c[i]  
...  
for (j=0;j<5;j++)  
    a[j+1]=a[j]+d[j];
```

Dependencies

- No dependence (can run in parallel)
S1: $X = K + 3$;
S2: $Y = Z * 5$;
- True dependence (cannot run in parallel)
S1: $X = 3$;
S2: $Y = X * 5$;
- Anti dependence (cannot run in parallel)
S1: $Y = X * 4$;
S2: $X = 3$;
- Output dependence (cannot run in parallel)
S1: $X = Y * 4$;
S2: $X = 3$;
- Loop-carried dependence

Methodology of Parallel Algorithm Design

- Programming model
 - Shared memory, Message passing, Data Parallel, Hybrid
- Four stages of parallel algorithm design:
 - Partitioning (decomposition, assignment): deal with machine independent issues and affect concurrency and scalability
 - Orchestration
 - Mapping
 - Orchestration & mapping: deal with machine dependent issues and affect locality and other performance issues
- Performance goals
- Application structures
 - SPMD, Embarrassingly Parallel, Master/Slave, Work Pool, Divide and Conquer, Pipeline, Competition
- Example

What have covered in MPI lecture

- What is MPI?
- How to write a simple program in MPI
 - MPI Communicators and communication
- Running your application with MPICH
- Slightly more advanced topics:
 - Non-blocking communication in MPI
 - Group (collective) communication in MPI
- Homeworks

Algorithm Design: Tridiagonal solvers

- Parallel Algorithms
 - The Partition Method
 - The PPT Algorithm
 - The PDD Algorithm
 - The LU Pipelining Algorithm
 - The PTH Method and PPD Algorithm Partitioning : deal with machine independent issues and affect concurrency and scalability
- Communication patterns
- Algorithm analysis
- Program assignment

Algorithm Design: Sorting

- Algorithm design
 - Parallelizing sequential alg., design new parallel algorithm, borrowing know alg.
- Existing techniques
 - Balanced tree, Doubling, Partition, Divide-and-Conquer, Pipelining,
- Sorting
 - Bubble sort, Mergesort, Quicksort, Bitonic sort
- Algorithm analysis

What have covered in OpenMP lecture

- What is OpenMP?
 - MPI + X
- OpenMP Basic and API
 - Compiler Directives
 - Runtime Library Routines
 - Environment Variables
- OpenMP slight more Advanced
 - Accelerator
 - Library/optimization
 - Examples
- Homework

Performance Evaluation

- Performance metrics

- Speedup $S_p = \frac{T_s}{T_p}$

- Efficiency $E_p = \frac{S_p}{p}$

- Scalability

- **Amdahl's law** gives a limit on speedup in terms of α (fraction of program (algorithm) that is serial and cannot be parallelized)

$$T_p = \alpha T_s + \frac{(1 - \alpha)T_s}{p}$$

$$S_p = \frac{T_s}{\alpha T_s + \frac{(1 - \alpha)T_s}{p}} = \frac{1}{\alpha + \frac{1 - \alpha}{p}}$$

Performance Evaluation

- Model of speedup
 - Fixed-size
 - Fixed-time
 - Memory-bounded
 - Generalized
- Characteristics of parallel code
 - granularity
 - load balance
 - locality
 - communication and synchronization
- Scalable computing
 - Scalability
 - Examples

Performance Evaluation 2

- Scalable computing
 - Scalability
 - Examples
- Isospeed scalability
 - Definition
 - Examples

Memory Performance Optimization

- Multicore
- Operation of Memory Hierarchy
 - Metrics
 - Cache, localities, cache miss (4Cs), associative, policies
 - Memory hierarchy performance and optimization
- Application Level Optimizations
 - Loop transformations
 - Array padding
 - Cache blocking, etc.
 - Replacement policies
- Examples

Memory Performance Optimization 2

Metrics and their calculation

- AMAT
- APC
- C-AMAT
 - Motivation, definition, recursive
 - Calculation, relation with other metrics
 - optimization
- LPM and Sluice Gate Theory
- Examples
- Revisit

What have covered in performance tool lecture

- What and why Performance tuning?
- Tools
 - Mesurément and profile
 - Pitfalls
- Linux
- Assignment

Lecture – File system basics

- What is the main role of file systems?
- How are data represented?
- How can we search for our data?
- How can we share data between users and applications?
- How are file systems generally organized?
 - Client-server
 - Cluster
 - Symmetric

Lecture – PFS

- What are the main design goals of PFS? Why do we need them?
- What kind of workloads require a PFS?
- Typical access patterns?
 - Shared file
 - File-per-process
- Data distribution in PFS
- What is POSIX I/O?
- How is concurrent access achieved efficiently?

Lecture – HDFS

- What is MapReduce and how does it work?
- What are the core concepts of Hadoop?
- Which workloads are suitable for Hadoop?
- Learn a few examples:
 - Word count in class
 - Can you design your own solution for sorting integers using MapReduce?

Lecture – HDFS (cont.)

- Basic design principles of HDFS?
- Highlight the HDFS architecture
- Data distribution in HDFS
 - Block placement
- Fault tolerance
 - Heart beats
- Data replication
 - Replica placement
 - Replica selection

Lecture – HDFS (cont.)

- Data pipelining
 - Describe all three ways to achieve it
- Namenode
 - Performance
 - Failures
- HDFS optimizations
 - Space reclamation
 - Communication protocols
 - Metadata

Distributed Memory Architecture

- Overview of distributed memory parallel architectures
- Interconnection networks
 - Connection topologies, 2-D mesh, hypercube
- Message passing issues
 - Routing, store-and-forward
 - Group communication

Message Passing Performance & Shared Environments

- The Log(P) model
- The memory Log(P) model
- The probability model for non-dedicated environment
- Why do we need them and how to use them
- The matrix multiple performance analysis

Highlight

- General concepts
 - Architecture, dependence, etc.
- Algorithms
 - Parallelism and analysis
- Performance
 - Speedup, scalability
- Memory performance
 - Hierarchy and C-AMAT
- I/O and file systems
 - PFS and HDFS
- Communication
- Programming

Exam Comments

- 3:00pm-5:00pm Room **SB111** on 11/14/19
 - Closed book, closed notes, no cell phone, no calculator
 - Allow to bring a 4"x6" "cheat card"
 - Anything on the card **MUST HAND WRITING**
- All lectures included on the test
 - Focus on things taught in class (lecture notes, in-class discussions, homework, programming assignment)
- Exam structure:
 - Short answer questions
 - 'big' questions
 - Programming design related questions
 - Performance analysis/evaluation questions
- Good luck!